

BACHELOR OF SCIENCE IN COMPUTER SCIENCE
YEAR 4 SEMESTER 2
CCS 418: ADVANCED DATABASE SYSTEMS
NOTES 1

Object-Oriented Databases

Need for Complex *Data Types*

- **Traditional** database *applications* in data processing had conceptually *simple data types*
 - Relatively *few data types*,
 - *first normal form* holds
- **Complex data types** have grown more important in *recent years*
 - E.g. *Addresses* can be viewed as a
 - Single string, or
 - *Separate attributes* for each part, or
 - *Composite attributes* (which are not in first normal form)
 - E.g. it is often *convenient* to store *multivalued attributes as-is*,
 - *without* creating a *separate relation* to store the values in 1FN
- **Applications:**
 - computer-aided design, computer-aided software engineering
 - multimedia and image databases, and document/hypertext databases.

Object-Oriented *Data Model*

- Loosely speaking, an **object** corresponds to an *entity* in the *E-R model*.
- The *object-oriented paradigm* is based on *encapsulating code* and *data* related to an object *into single unit*.
- The *object-oriented data model* is a *logical data model* (like the *E-R model*).
- **Adaptation** of the *object-oriented programming paradigm*:
 - (e.g., Smalltalk, C++, Java)
 - to *database systems*.

Object Structure

- An object has associated with it:
 - A set of **variables** that contain the *data* for the object. The value of each variable is itself an object.
 - A set of **messages** to which the object responds; each message may have zero, one, or more *parameters*.

- A set of **methods**, each of which is a *body of code* to implement a message; a method returns a value as the *response* to the message
- The physical representation of data is visible only to the implementer of the object
- **Messages** and responses *provide* the only *external interface* to an object.
- The term message does *not necessarily* imply physical *message passing*.
 - **Messages** can be *implemented* as **procedure invocations**.

Messages and Methods

- **Methods** are programs written in *general-purpose language* with the following features:
 - only variables in the object itself may be referenced directly
 - **data** in other objects are *referenced* only by *sending messages*.
- **Methods** can be *read-only* or *update methods*
 - Read-only methods do not change the value of the object
- *Strictly speaking*, **every** attribute of an entity must be represented by:
 - a variable and
 - two methods, one to read and the other to update the attribute
 - e.g., the attribute *address* is represented by a variable *address* and two **messages** *get-address* and *set-address*.
 - For convenience, many object-oriented data models permit direct access to variables of other objects.

Object Classes

- *Similar objects* are grouped into:
 - a **class**; each such object is called
 - an **instance** of its class
- *All objects* in a **class** have the **same**:
 - Variables, with the same types
 - message interface
 - methods

They may differ in the values assigned to variables

- *Example*: Group objects for **people** into a **person class**
- *Classes* are *analogous to* **entity sets** in the **E-R** model

Class Definition Example

```
class employee {
    /*Variables */
    string name;
```

```


        string    address;
        date     start-date;
        int      salary;
/* Messages */
        int      annual-salary();
        string   get-name();
        string   get-address();
        int      set-address(string new-address);
        int      employment-length();
};

```

- Methods to read and set the other variables are also needed with strict encapsulation
- Methods are defined separately
 - E.g. **int** *employment-length*()
 - { **return** *today*() – *start-date*;}
 - **int** *set-address*(**string** *new-address*)
 - { *address* = *new-address*;}

Inheritance

- E.g., class of bank customers is similar to class of bank employees, although there are differences
 - both share some variables and messages, e.g., *name* and *address*.
 - But there are variables and messages specific to each class e.g., *salary* for employees and *credit-rating* for customers.
- Every employee is a person; thus *employee* is a specialization of *person*
- Similarly, *customer* is a specialization of *person*.
- Create classes *person*, *employee* and *customer*
 - variables/messages applicable to all persons associated with class *person*.
 - variables/messages specific to employees associated with class *employee*; similarly for *customer*
- Place classes into a specialization/IS-A hierarchy
 - variables/messages belonging to class *person* are *inherited* by class *employee* as well as *customer*
- Result is a **class hierarchy**



Class Hierarchy Definition

```

class person {
    string name;
    string address;
};

class customer isa person {
    int credit-rating;
};

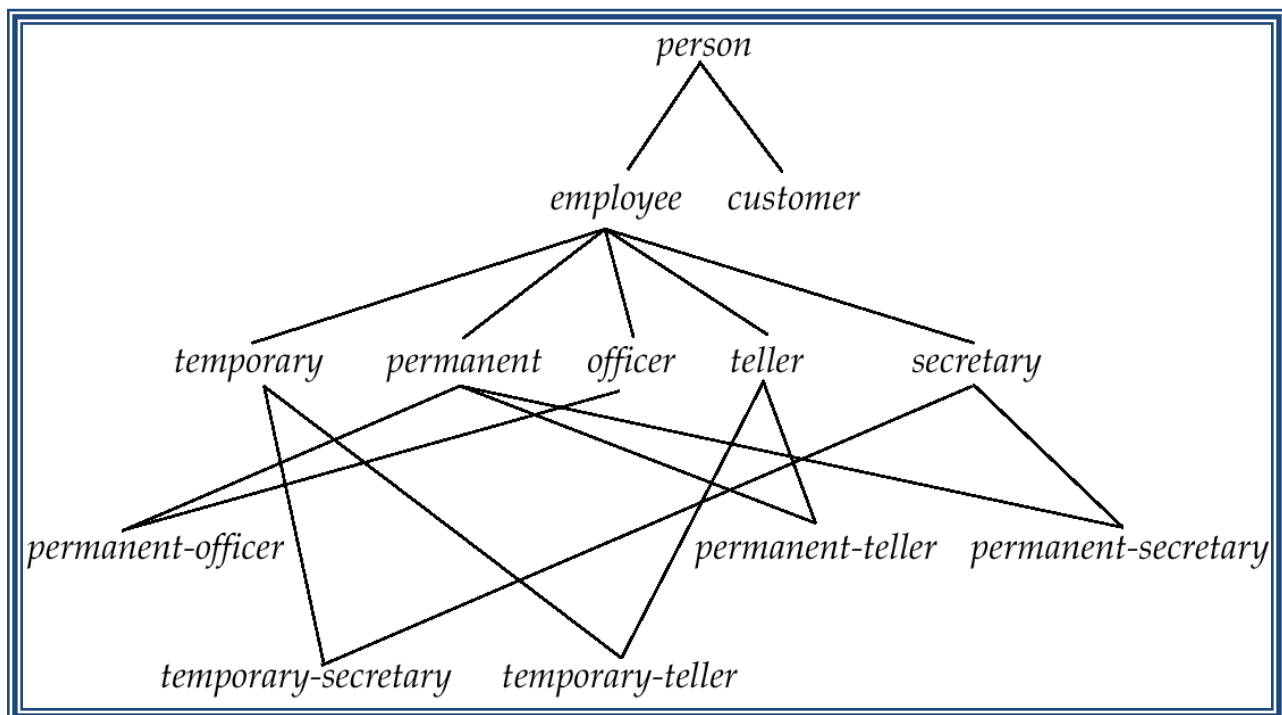
class employee isa person {
    date start-date;
    int salary;
};

class officer isa employee {
    int office-number,
    int expense-account-number,
};
  
```

- Full variable list for objects in the class *officer*:
 - *office-number, expense-account-number*: defined locally
 - *start-date, salary*: inherited from *employee*
 - *name, address*: inherited from *person*
- Methods inherited similar to variables.
- **Substitutability** — any *method* of a class, say *person*, *can be invoked* equally well *with* any object belonging to any *subclass*, such as subclass *officer* of *person*.
- **Class extent**: set of **all objects** in the *class*. Two options:

1. Class extent of *employee* includes **all** *officer*, *teller* and *secretary* objects.
2. Class extent of *employee* includes **only** *employee* objects that are **not** in a subclass such as *officer*, *teller*, or *secretary*
 - This is the **usual choice** in OO systems
 - Can access extents of subclasses to find all objects of subtypes of *employee*

Example of Multiple Inheritance



Class DAG for banking example.

Multiple Inheritance

- With multiple inheritance a class may **have more than one superclass**.
 - The class/subclass relationship is represented by a **directed acyclic graph (DAG)**
 - Particularly useful when objects can be classified in more than one way, which are independent of each other
 - E.g. temporary/permanent is independent of Officer/secretary/teller
 - Create a subclass for each combination of subclasses
 - Need not create subclasses for combinations that are not possible in the database being modeled

- A class **inherits** *variables* and *methods* from all its *superclasses*
- There is potential for **ambiguity** when a variable/message N with the *same name* is inherited from two superclasses A and B
 - No problem if the variable/message is defined in a shared superclass
 - Otherwise, do one of the following
 - flag as an error,
 - rename variables (A.N and B.N)
 - choose one.

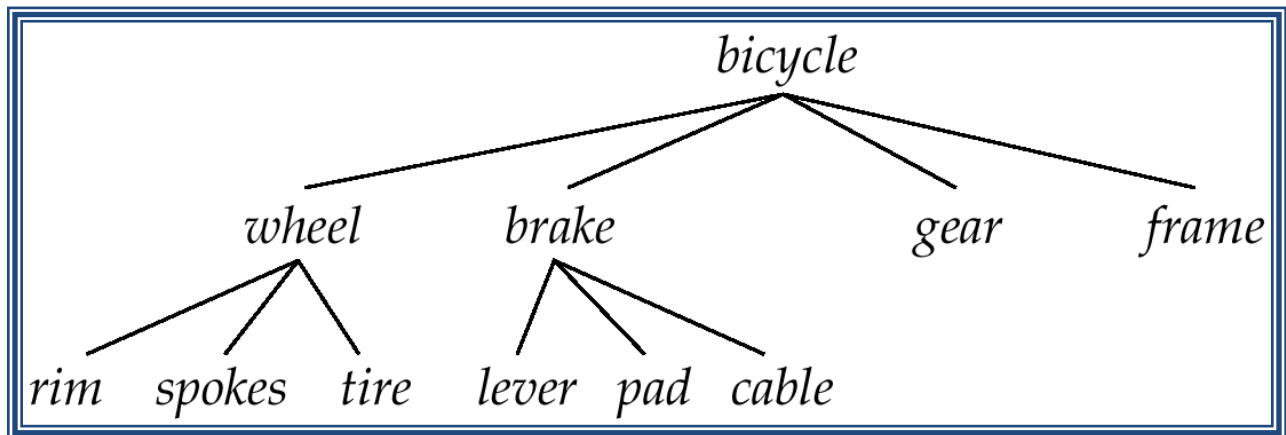
Object Identity

- An object retains its identity even if some or all of the values of variables or definitions of methods change over time.
- Object identity is a stronger notion of identity than in programming languages or data models not based on object orientation.
 - Value – data value; e.g. primary key value used in relational systems.
 - Name – supplied by user; used for variables in procedures.
 - Built-in – identity built into data model or programming language.
 - no user-supplied identifier is required.
 - Is the form of identity used in object-oriented systems.

Object Identifiers

- **Object identifiers** used to uniquely identify objects
 - Object identifiers are unique:
 - no two objects have the same identifier
 - each object has only one object identifier
 - E.g., the *spouse* field of a *person* object may be an identifier of another *person* object.
 - can be stored as a field of an object, to refer to another object.
 - Can be:
 - **system generated** (created by database) or
 - **external** (such as social-security number)
 - *System generated* identifiers:
 - Are *easier* to use, but **cannot** be used *across* database systems
 - May be **redundant** if unique identifier already exists

Object Containment



- Each component in a design may contain other components
- Can be modeled as containment of objects.
- Objects containing other objects are called **composite** objects.
- Multiple levels of containment create a **containment hierarchy**
 - links interpreted as **is-part-of**, not **is-a**.
- Allows data to be viewed at different granularities by different users.

Object-Oriented Languages

- *Object-oriented concepts* can be used in *different ways*
 - Object-orientation can be used as a *design tool*, and be encoded into,
 - for example, a *relational* database
 - analogous to *modeling* data with *E-R* diagram and
 - then *converting* to a set of *relations*
- The *concepts* of object orientation can be *incorporated* into a *programming language* that is used to manipulate the database.
 - **Object-relational systems** – add complex types and object-orientation to relational language.
 - **Persistent programming languages** – extend object-oriented programming language to deal with databases by adding concepts such as persistence and collections.

Persistent Programming Languages

- *Persistent Programming languages* allow objects to be created and stored in a database, and used directly from a programming language
 - allow data to be manipulated directly from the programming language

- **No need** to go through **SQL**.
- **No need** for explicit format (type) changes
 - format changes are carried out *transparently* by system
 - Without a persistent programming language, format changes becomes a burden on the programmer
 - More code to be written
 - More chance of bugs
- **allow objects** to be *manipulated* in-memory
 - **no need** to explicitly **load** from or **store** to the database
 - Saved code, and saved overhead of loading/storing large amounts of data
- **Drawbacks** of *persistent programming* languages
 - Due to *power* of most programming *languages*,
 - it is easy to make *programming errors* that **damage** the *database*.
 - **Complexity** of languages makes
 - automatic high-level *optimization* more **difficult**.
 - **Do not support declarative querying** as well as relational databases

Persistence of Objects

- Approaches to make transient objects persistent include establishing
 - Persistence by Class – declare all objects of a class to be persistent; simple but inflexible.
 - Persistence by Creation – extend the syntax for creating objects to specify that that an object is persistent.
 - Persistence by Marking – an object that is to persist beyond program execution is marked as persistent before program termination.
 - Persistence by Reachability - declare (root) persistent objects; objects are persistent if they are referred to (directly or indirectly) from a root object.
 - Easier for programmer, but more overhead for database system
 - Similar to garbage collection used e.g. in Java, which also performs reachability tests

Object Identity and Pointers

- A persistent object is assigned a persistent object identifier.
- Degrees of permanence of identity:
 - Intraprocedure – identity persists only during the executions of a single procedure
 - Intraprogram – identity persists only during execution of a single program or query.

- Interprogram – identity persists from one program execution to another, but may change if the storage organization is changed
- Persistent – identity persists throughout program executions and structural reorganizations of data; required for object-oriented systems.
- In O-O languages such as C++, an object identifier is actually an in-memory pointer.
- Persistent pointer – persists beyond program execution
 - can be thought of as a pointer into the database
 - E.g. specify file identifier and offset into the file
 - Problems due to database reorganization have to be dealt with by keeping forwarding pointers

Storage and Access of Persistent Objects

How to find objects in the database:

- *Name* objects (as you would name files)
 - Cannot scale to *large number* of objects.
 - Typically given only to class extents and other collections of objects, but not objects.
- Expose *object identifiers* or *persistent pointers* to the objects
 - Can be stored externally.
 - *All objects have object identifiers.*
- *Store collections of objects*, and allow programs to iterate over the collections to find required objects
 - *Model* collections of objects as **collection types**
 - **Class extent** - the collection of all objects belonging to the class;
 - usually maintained for *all classes* that can have *persistent objects*.

Persistent C++ Systems

- C++ language allows support for persistence to be added without changing the language
 - Declare a class called `Persistent_Object` with attributes and methods to support persistence
 - **Overloading** – ability to redefine standard function names and operators (i.e., +, -, the pointer dereference operator `->`) when applied to new types
 - **Template classes** help to build a type-safe type system supporting collections and persistent types.
- Providing persistence without extending the C++ language is
 - relatively easy to implement
 - but more difficult to use

- Persistent C++ systems that add features to the C++ language have been built, as also systems that avoid changing the language

ODMG C++ Object Definition Language

- The *Object Database Management Group* is
 - an **industry consortium** aimed at
 - *standardizing* object-oriented databases
 - in particular *persistent programming* languages
 - Includes standards for C++, Smalltalk and Java
 - ODMG-2.0 and 3.0 (which is 2.0 plus extensions to Java)
 - The description based on ODMG-2.0
- **ODMG** C++ standard **avoids** *changes* to the C++ language
 - *provides* functionality **via** template classes and *class libraries*

ODMG Types

- Template class `d_Ref<class>` used to specify references (persistent pointers)
- Template class `d_Set<class>` used to define sets of objects.
 - Methods include `insert_element(e)` and `delete_element(e)`
- Other collection classes such as `d_Bag` (set with duplicates allowed), `d_List` and `d_Varray` (variable length array) also provided.
- `d_` version of many standard types provided, e.g. `d_Long` and `d_string`
 - Interpretation of these types is platform independent
 - Dynamically allocated data (e.g. for `d_string`) allocated in the database, not in main memory

ODMG C++ ODL: Example

```
class Branch : public d_Object {
    ....
}

class Person : public d_Object {
public:
    d_String  name;    // should not use String!
    d_String  address;
};

class Account : public d_Object {
private:
    d_Long    balance;
```

```

public:
    d_Long    number;
    d_Set <d_Ref<Customer>> owners;
    int    find_balance();
    int    update_balance(int delta);
};

class Customer : public Person {
public:
    d_Date    member_from;
    d_Long    customer_id;
    d_Ref<Branch> home_branch;
    d_Set <d_Ref<Account>> accounts; };

```

Implementing Relationships

- Relationships between classes implemented by references
- Special reference types enforces integrity by adding/removing inverse links.
 - Type `d_Rel_Ref<Class, InvRef>` is a reference to `Class`, where attribute `InvRef` of `Class` is the inverse reference.
 - Similarly, `d_Rel_Set<Class, InvRef>` is used for a set of references
- Assignment method (`=`) of class `d_Rel_Ref` is overloaded
 - Uses type definition to automatically find and update the inverse link
 - Frees programmer from task of updating inverse links
 - Eliminates possibility of inconsistent links
- Similarly, `insert_element()` and `delete_element()` methods of `d_Rel_Set` use type definition to find and update the inverse link automatically
- E.g.

```

extern const char _owners[ ], _accounts[ ];

class Account : public d.Object {
    ....
    d_Rel_Set <Customer, _accounts> owners;
}

// .. Since strings can't be used in templates ...

```

```
const char _owners = "owners";
const char _accounts = "accounts";
```

ODMG C++ Object Manipulation Language

- Uses persistent versions of C++ operators such as new(db)


```
d_Ref<Account> account = new(bank_db, "Account") Account;
```

 - new allocates the object in the specified database, rather than in memory.
 - The second argument ("Account") gives typename used in the database.
- Dereference operator -> when applied on a d_Ref<Account> reference loads the referenced object in memory (if not already present) before continuing with usual C++ dereference.
- Constructor for a class – a special method to initialize objects when they are created; called automatically on new call.
- Class extents maintained automatically on object creation and deletion
 - Only for classes for which this feature has been specified
 - Specification via user interface, not C++
 - Automatic maintenance of class extents not supported in earlier versions of ODMG

ODMG C++OML: Database and Object Functions

- Class d_Database provides methods to
 - open a database : open(databasename)
 - **give names** to objects: set_object_name(object, name)
 - **look up** objects *by name*: lookup_object(name)
 - **rename** objects: rename_object(oldname, newname)
 - close a database (close());
- Class d_Object is inherited by all persistent classes.
 - provides methods to allocate and delete objects
 - method mark_modified() must be called *before* an object is updated.

Is automatically called when object is created

ODMG C++ OML: Example

```
int create_account_owner(String name, String Address){
    Database bank_db.obj;
    Database * bank_db= & bank_db.obj;
```

```

bank_db =>open("Bank-DB");
d.Transaction Trans;
Trans.begin();

d_Ref<Account> account = new(bank_db) Account;
d_Ref<Customer> cust = new(bank_db) Customer;
cust->name = name;
cust->address = address;
cust->accounts.insert_element(account);
... Code to initialize other fields

Trans.commit();
}

```

- Class extents maintained automatically in the database.
- To access a class extent:


```
d_Extent<Customer> customerExtent(bank_db);
```
- Class d_Extent provides method


```
d_Iterator<T> create_iterator()
```

 to create an iterator on the class extent
- Also provides select(pred) method to return iterator on objects that satisfy selection predicate pred.
- Iterators help step through objects in a collection or class extent.
- Collections (sets, lists etc.) also provide create_iterator() method.

ODMG C++ OML: Example of Iterators

```

int print_customers() {
Database bank_db_obj;
Database * bank_db = &bank_db_obj;
bank_db->open ("Bank-DB");
d_Transaction Trans; Trans.begin ();

d_Extent<Customer> all_customers(bank_db);
d_Iterator<d_Ref<Customer>> iter;

```

```

iter = all_customers->create_iterator();
d_Ref <Customer> p;
    while{iter.next (p))
        print_cust (p); // Function assumed to be defined elsewhere
    Trans.commit();
}

```

ODMG C++ Binding: Other *Features*

- *Declarative query language OQL*, looks like **SQL**
 - Form query as a string, and execute it to get a set of results (actually a bag, since duplicates may be present)

```

d_Set<d_Ref<Account>> result;
d_OQL_Query q1("select a
                from Customer c, c.accounts a
                where c.name='Jones'
                and a.find_balance() > 100");

```

```

d_oql_execute(q1, result);

```

- Provides error handling mechanism based on C++ exceptions, through class `d_Error`
- Provides **API** for accessing the schema of a database.

Making Pointer Persistence Transparent

- **Drawback** of the ODMG C++ *approach*:
 - *Two types* of **pointers**
 - *Programmer* has to **ensure**
 - `mark_modified()` is called,
 - else database can become **corrupted**
- *ObjectStore approach*
 - Uses *exactly* the same pointer type for in-memory and database objects
 - Persistence is transparent applications
 - Except when creating objects
 - Same functions can be used on in-memory and persistent objects since pointer types are the same
 - Implemented by a technique called pointer-swizzling which is described in Chapter 11.
 - No need to call `mark_modified()`, modification detected automatically.

Persistent Java Systems

- ODMG-3.0 defines *extensions* to Java for persistence
 - Java does **not** support **templates**,
 - so language *extensions* are *required*
- *Model* for persistence: *persistence by reachability*
 - *Matches* Java's *garbage collection* model
 - Garbage collection needed on the database also
 - **Only** *one* pointer *type* for transient and persistent pointers
- *Class* is made *persistence capable* by running
 - a *post-processor* on object code generated by the Java compiler
 - *Contrast* with *pre-processor* used in C++
 - Post-processor **adds** `mark_modified()` *automatically*
- Defines collection types DSet, DBag, DList, etc.
- Uses Java iterators, no need for new iterator class

ODMG Java

- Transaction must start accessing database from one of the root object (***looked up*** by name)
 - finds other objects by following pointers from the root objects
- Objects ***referred*** to from a *fetch*ed object are:
 - *allocated* space in memory,
 - but *not* necessarily *fetch*ed
- ***Fetch***ing can be done *lazily*
 - An object with space allocated but *not* yet *fetch*ed is called a *hollow object*
 - When a *hollow object* is accessed, its data is fetched from disk.

Object-Relational Data Models

- Extend the relational data model by including object orientation and constructs to deal with added data types.
- Allow attributes of tuples to have complex types, including non-atomic values such as nested relations.
- Preserve relational foundations, in particular the declarative access to data, while extending modeling power.
- Upward compatibility with existing relational languages.

Complex Data Types

- Motivation:
 - Permit non-atomic domains (atomic \equiv indivisible)
 - Example of non-atomic domain: set of integers, or set of tuples
 - Allows more intuitive modeling for applications with complex data
- Intuitive definition:
 - allow relations whenever we allow atomic (scalar) values — relations within relations
 - Retains mathematical foundation of relational model
 - Violates first normal form.

Example of a Nested Relation

- Example: library information system
- Each book has
 - title,
 - a set of authors,
 - Publisher, and
 - a set of keywords
- Non-1NF relation *books*

Title	Author-set	Publisher	Keyword-set
		(Name, Branch)	
Compilers	(Smith, Jones)	(McGraw-Hill, New York)	{ Parsing, Analysis }
Networks	(Jones, Frick)	(Oxford, London)	{ Internet, Web }

4NF Decomposition of Nested Relation

- Remove awkwardness of *flat-books* by assuming that the following multivalued dependencies hold:
 - $title \twoheadrightarrow author$
 - $title \twoheadrightarrow keyword$
 - $title \twoheadrightarrow pub-name, pub-branch$
- Decompose *flat-doc* into 4NF using the schemas:
 - $(title, author)$
 - $(title, keyword)$
 - $(title, pub-name, pub-branch)$

4NF Decomposition of *flat-books*

Title	Author
Compilers	Smith
Compilers	Jones
Networks	Jones
Networks	Frick
Authors	

Title	Keyword
Compilers	Parsing
Compilers	Analysis
Networks	Internet
Networks	Web
Keywords	

Title	Pub-Name	Pub-Branch
Compilers	McGraw-Hill	New York
Networks	Oxford	London
Books4		

Problems with 4NF Schema

- 4NF design requires users to include joins in their queries.
- 1NF relational view *flat-books* defined by join of 4NF relations:
 - eliminates the need for users to perform joins,
 - but loses the one-to-one correspondence between tuples and documents.
 - And has a large amount of redundancy
- Nested relations representation is much more natural here.

Complex Types and SQL

- Extensions to SQL to support complex types include:
 - Collection and large object types
 - ▶ Nested relations are an example of collection types
 - Structured types
 - ▶ Nested record structures like composite attributes
 - Inheritance
 - Object orientation
 - ▶ Including object identifiers and references
- Our description is mainly based on the SQL:1999 standard
 - Not fully implemented in any database system currently
 - But some features are present in each of the major commercial database systems
 - ▶ Read the manual of your database system to see what it supports

Structured Types and Inheritance in SQL

- Structured types can be declared and used in SQL

```
create type Name as  
  (firstname varchar(20),  
   lastname varchar(20))  
final  
create type Address as  
  (street varchar(20),  
   city varchar(20),  
   zipcode varchar(20))  
not final
```

- Note: **final** and **not final** indicate whether subtypes can be created
- Structured types can be used to create tables with composite attributes

```
create table customer (  
    name Name,  
    address Address,  
    dateOfBirth date)
```

- Dot notation used to reference components: *name.firstname*
- User-defined row types

```
create type CustomerType as (  
    name Name,  
    address Address,  
    dateOfBirth date)  
not final
```

- Can then create a table whose rows are a user-defined type

```
create table customer of CustomerType
```

Methods

- Can add a method declaration with a structured type.
 - **method** *ageOnDate* (*onDate* **date**)

- **returns interval year**

- Method body is given separately.

```
create instance method ageOnDate (onDate date)
    returns interval year
    for CustomerType
begin
    return onDate - self.dateOfBirth;
end
```

- We can now find the age of each customer:

```
select name.lastname, ageOnDate (current_date)
from customer
```

Inheritance

- Suppose that we have the following type definition for people:

```
create type Person
    (name varchar(20),
     address varchar(20))
```

- Using inheritance to define the student and teacher types

```
create type Student
    under Person
    (degree varchar(20),
     department varchar(20))

create type Teacher
    under Person
    (salary integer,
     department varchar(20))
```

- Subtypes can redefine methods by using **overriding method** in place of **method** in the method declaration

Multiple Inheritance

- Some SQL do not support multiple inheritance
- If our type system supports multiple inheritance, we can define a type for teaching assistant as follows:

create type *Teaching Assistant*

under *Student, Teacher*

- To avoid a conflict between the two occurrences of *department* we can rename them

create type *Teaching Assistant*

under

Student **with** (*department as student_dept*),

Teacher **with** (*department as teacher_dept*)

Consistency Requirements for Subtables

- Consistency requirements on subtables and supertables.
 - Each tuple of the supertable (e.g. *people*) can correspond to at most one tuple in each of the subtables (e.g. *students* and *teachers*)
 - Additional constraint in SQL:
- All tuples corresponding to each other (that is, with the same values for inherited attributes) must be derived from one tuple (inserted into one table).
 - ▶ That is, each entity must have a most specific type
 - ▶ We cannot have a tuple in *people* corresponding to a tuple each in *students* and *teachers*

Array and Multiset Types in SQL

- Example of array and multiset declaration:

create type *Publisher* **as**

(*name* **varchar**(20),

branch **varchar**(20))

create type *Book* **as**

(*title* **varchar**(20),

author-array **varchar**(20) **array** [10],

pub-date **date**,

publisher *Publisher*,
keyword-set **varchar**(20) **multiset**)

create table *books* **of** *Book*

- Similar to the nested relation *books*, but with array of authors instead of set

Creation of Collection Values

- Array construction

array [‘Silberschatz’, ‘Korth’, ‘Sudarshan’]

- Multisets

▪ **multisetset** [‘computer’, ‘database’, ‘SQL’]

- To create a tuple of the type defined by the *books* relation: (‘Compilers’,

array[‘Smith’, ‘Jones’],

Publisher (‘McGraw-Hill’, ‘New York’),

multiset

[‘parsing’, ‘analysis’])

- To insert the preceding tuple into the relation *books*

insert into *books*

values

(‘Compilers’, **array**[‘Smith’, ‘Jones’],

Publisher (‘McGraw-Hill’, ‘New York’),

multiset

[‘parsing’, ‘analysis’])

Querying Collection-Valued Attributes

- To find all books that have the word “database” as a keyword,

select *title*

from *books*

where ‘database’ **in** (**unnest**(*keyword-set*))

- We can access individual elements of an array by using indices

▪ E.g.: If we know that a particular book has three authors, we could write:

select *author-array*[1], *author-array*[2], *author-array*[3]

from *books*

where *title* = ‘Database System Concepts’

- To get a relation containing pairs of the form “title, author-name” for each book and each author of the book

```
select B.title, A.author
from books as B, unnest (B.author-array) as A (author )
```

- To retain ordering information we add a **with ordinality** clause

```
select B.title, A.author, A.position
from books as B, unnest (B.author-array) with ordinality as
      A (author, position )
```

Unnesting

- The transformation of a nested relation into a form with fewer (or no) relation-valued attributes is called **unnesting**.

- E.g.

```
select title, A as author, publisher.name as pub_name,
      publisher.branch as pub_branch, K.keyword
from books as B, unnest(B.author_array ) as A (author ),
      unnest (B.keyword_set ) as K (keyword )
```

Nesting

- **Nesting** is the opposite of unnesting, creating a collection-valued attribute
- Nesting can be done in a manner similar to aggregation, but using the function **collect()** in place of an aggregation operation, to create a multiset
- To nest the *flat-books* relation on the attribute *keyword*:

```
select title, author, Publisher (pub_name, pub_branch ) as publisher,
      collect (keyword) as keyword_set
```

```
from flat-books
```

```
groupby title, author, publisher
```

- To nest on both authors and keywords:

```
select title, collect (author ) as author_set,
      Publisher (pub_name, pub_branch) as publisher,
      collect (keyword ) as keyword_set
```

```
from flat-books
```

```
group by title, publisher
```

1NF Version of Nested Relation

1NF version of *books*

Title	Author	Pun-Name	Pub-Branch	Keyword
Compilers	Smith	McGraw-Hill	New York	Parsing
Compilers	Jones	McGraw-Hill	New York	Parsing
Compilers	Smith	McGraw-Hill	New York	Analysis
Compilers	Jones	McGraw-Hill	New York	Analysis
Networks	Jones	Oxford	London	Internet
Networks	Frick	Oxford	London	Internet
Networks	Jones	Oxford	London	Web
Networks	Frick	Oxford	London	Web

Flat-books

- Another approach to creating nested relations is to use subqueries in the **select** clause.

select *title*,

array (**select** *author*

from *authors* **as** *A*

where *A.title* = *B.title*

order by *A.position*) **as** *author_array*,

Publisher (*pub-name*, *pub-branch*) **as** *publisher*,

multiset (**select** *keyword*

from *keywords* **as** *K*

where *K.title* = *B.title*) **as** *keyword_set*

from *books4* **as** *B*

Object-Identity and Reference Types

- Define a type *Department* with a field *name* and a field *head* which is a reference to the type *Person*, with table *people* as scope:
 - **create type** *Department* (
name **varchar** (20),
head **ref** (*Person*) **scope** *people*)
- We can then create a table *departments* as follows
 - **create table** *departments* **of** *Department*

- We can omit the declaration **scope** people from the type declaration and instead make an addition to the **create table** statement:

```
create table departments of Department
    (head with options scope people)
```

Initializing Reference-Typed Values

- To create a tuple with a reference value, we can first create the tuple with a null reference and then set the reference separately:

```
insert into departments
    values ('CS', null)

update departments
    set head = (select p.person_id
                from people as p
                where name = 'John')
    where name = 'CS'
```

User Generated Identifiers

- The type of the object-identifier must be specified as part of the type definition of the referenced table, and
- The table definition must specify that the reference is user generated

```
create type Person
    (name varchar(20)
     address varchar(20))
    ref using varchar(20)
create table people of Person
    ref is person_id user generated
```

- When creating a tuple, we must provide a unique value for the identifier:

```
insert into people (person_id, name, address ) values
('01284567', 'John', '23 Coyote Run')
```

- We can then use the identifier value when inserting a tuple into *departments*
 - Avoids need for a separate query to retrieve the identifier:

insert into *departments*
values(`CS`, `02184567`)

- Can use an existing primary key value as the identifier:

create type *Person*

(*name* **varchar** (20) **primary key**,

address **varchar**(20))

ref from (*name*)

create table *people* **of** *Person*

ref is *person_id* **derived**

- When inserting a tuple for *departments*, we can then use

insert into *departments*

values(`CS`, `John`)

Path Expressions

- Find the names and addresses of the heads of all departments:

select *head* \rightarrow *name*, *head* \rightarrow *address*

from *departments*

- An expression such as “*head* \rightarrow *name*” is called a **path expression**
- Path expressions help avoid explicit joins
 - If department head were not a reference, a join of *departments* with *people* would be required to get at the address
 - Makes expressing the query much easier for the user

Implementing O-R Features

- Similar to how E-R features are mapped onto relation schemas
- Subtable implementation
 - Each table stores primary key and those attributes defined in that table or
 - Each table stores both locally defined and inherited attributes

Persistent Programming Languages

- Languages extended with constructs to handle persistent data
- Programmer can manipulate persistent data directly
 - no need to fetch it into memory and store it back to disk (unlike embedded SQL)

- Persistent objects:
 - by class - explicit declaration of persistence
 - by creation - special syntax to create persistent objects
 - by marking - make objects persistent after creation
 - by reachability - object is persistent if it is declared explicitly to be so or is reachable from a persistent object

Object Identity and Pointers

- Degrees of permanence of object identity
 - Intraprocedure: only during execution of a single procedure
 - Intraprogram: only during execution of a single program or query
 - Interprogram: across program executions, but not if data-storage format on disk changes
 - Persistent: interprogram, plus persistent across data reorganizations
- Persistent versions of C++ and Java have been implemented
 - C++
 - ▶ ODMG C++
 - ▶ ObjectStore
 - Java
 - ▶ Java Database Objects (JDO)

Comparison of O-O and O-R Databases

- **Relational systems**
 - simple data types, powerful query languages, high protection.
- **Persistent-programming-language-based OODBs**
 - complex data types, integration with programming language, high performance.
- **Object-relational systems**
 - complex data types, powerful query languages, high protection.
- Note: Many real systems blur these boundaries
 - E.g. persistent programming language built as a wrapper on a relational database offers first two benefits, but may have poor performance.

ADVANCED SQL

Built-in Data Types in SQL

- **date:** Dates, containing a (4 digit) year, month and date
 - Example: **date** '2005-7-27'
- **time:** Time of day, in hours, minutes and seconds.
 - Example: **time** '09:00:30' **time** '09:00:30.75'
- **timestamp:** date plus time of day
 - Example: **timestamp** '2005-7-27 09:00:30.75'
- **interval:** period of time
 - Example: **interval** '1' day
 - Subtracting a date/time/timestamp value from another gives an interval value
 - Interval values can be added to date/time/timestamp values
- Can extract values of individual fields from date/time/timestamp
 - Example: **extract (year from r.starttime)**
- Can cast string types to date/time/timestamp
 - Example: **cast** <string-valued-expression> **as date**
 - Example: **cast** <string-valued-expression> **as time**

User-Defined Types

- **create type** construct in SQL creates user-defined type
 - **create type Dollars as numeric (12,2) final**
- **create domain** construct in SQL-92 creates user-defined domain types
 - **create domain person_name char(20) not null**
- Types and domains are similar. Domains can have constraints, such as **not null**, specified on them.
- **Domain Constraints**
- **Domain constraints** are the most elementary form of integrity constraint. They test values inserted in the database, and test queries to ensure that the comparisons make sense.
- New domains can be created from existing data types

- Example: **create domain Dollars numeric(12, 2)**
 create domain Pounds numeric(12,2)
- We cannot assign or compare a value of type Dollars to a value of type Pounds.
 - However, we can convert type as below
 (cast r.A as Pounds)
 (Should also multiply by the dollar-to-pound conversion-rate)

Large-Object Types

- Large objects (photos, videos, CAD files, etc.) are stored as a *large object*:
 - **blob**: binary large object -- object is a large collection of uninterpreted binary data (whose interpretation is left to an application outside of the database system)
 - **clob**: character large object -- object is a large collection of character data
 - When a query returns a large object, a pointer is returned rather than the large object itself.

Integrity Constraints

- Integrity constraints guard against accidental damage to the database, by ensuring that authorized changes to the database do not result in a loss of data consistency.
 - A checking account must have a balance greater than \$10,000.00
 - A salary of a bank employee must be at least \$4.00 an hour
 - A customer must have a (non-null) phone number

Constraints on a Single Relation

- not null
- primary key
- unique
- check (*P*), where *P* is a predicate
- Not Null Constraint Declare *branch_name* for *branch* is not null
- *branch_name* char(15) not null
- Declare the domain *Dollars* to be not null
- create domain *Dollars* numeric(12,2) not null
- The Unique Constraint unique (*A*₁, *A*₂, ..., *A*_m)

- The unique specification states that the attributes

$A_1, A_2, \dots A_m$

Form a candidate key.

Candidate keys are permitted to be non null (in contrast to primary keys).

The check clause

- **check** (P), where P is a predicate
 - Example: Declare *branch_name* as the primary key for *branch* and ensure that the values of *assets* are non-negative.
 - **create table** *branch*

```
(branch_name  char(15),
  branch_city   char(30),
  assets        integer,
  primary key (branch_name),
  check (assets >= 0))
```
- The **check** clause in SQL permits domains to be restricted:
 - Use **check** clause to ensure that an *hourly_wage* domain allows only values greater than a specified value.
 - **create domain** *hourly_wage* **numeric**(5,2)


```
constraint value_test check(value >= 4.00)
```
 - The domain has a constraint that ensures that the *hourly_wage* is greater than 4.00
 - The clause **constraint** *value_test* is optional; useful to indicate which constraint an update violated.

Referential Integrity

- Ensures that a value that appears in one relation for a given set of attributes also appears for a certain set of attributes in another relation.
 - Example: If “Perryridge” is a branch name appearing in one of the tuples in the *account* relation, then there exists a tuple in the *branch* relation for branch “Perryridge”.

- Primary and candidate keys and foreign keys can be specified as part of the SQL **create table** statement:
 - The primary key clause lists attributes that comprise the primary key.
 - The unique key clause lists attributes that comprise a candidate key.
 - The foreign key clause lists the attributes that comprise the foreign key and the name of the relation referenced by the foreign key. By default, a foreign key references the primary key attributes of the referenced table.

Referential Integrity in SQL – Example

create table *customer*

(*customer_name* **char**(20),
customer_street **char**(30),
customer_city **char**(30),
primary key (*customer_name*))

create table *branch*

(*branch_name* **char**(15),
branch_city **char**(30),
assets **numeric**(12,2),
primary key (*branch_name*))

create table *account*

(*account_number* **char**(10),
branch_name **char**(15),
balance **integer**,
primary key (*account_number*),
foreign key (*branch_name*) **references** *branch*)

create table *depositor*

(*customer_name* **char**(20),
account_number **char**(10),
primary key (*customer_name*, *account_number*),

foreign key (*account_number*) **references** *account*,
foreign key (*customer_name*) **references** *customer*)

Assertions

- An **assertion** is a predicate expressing a condition that we wish the database always to satisfy.
- An assertion in SQL takes the form
 - **create assertion** <assertion-name> **check** <predicate>
- When an assertion is made, the system tests it for validity, and tests it again on every update that may violate the assertion
 - This testing may introduce a significant amount of overhead; hence assertions should be used with great care.
- Asserting
for all $X, P(X)$
is achieved in a round-about fashion using
not exists X such that not $P(X)$

Assertion Example

- Every loan has at least one borrower who maintains an account with a minimum balance or \$1000.00

create assertion *balance_constraint* **check**

(**not exists** (

select *

from *loan*

where not exists (

select *

from *borrower, depositor, account*

where *loan.loan_number* = *borrower.loan_number*

and *borrower.customer_name* = *depositor.customer_name*

and *depositor.account_number* = *account.account_number*

and *account.balance* >= 1000)))

Assertion Example

- The sum of all loan amounts for each branch must be less than the sum of all account balances at the branch.

create assertion *sum_constraint* **check**

(not exists (select *

from *branch*

where (select sum(amount)

from *loan*

where *loan.branch_name =*

branch.branch_name)

>= (select sum (amount)

from *account*

where *loan.branch_name =*

branch.branch_name)))

Authorization

Forms of authorization on parts of the database:

- **Read** - allows reading, but not modification of data.
- **Insert** - allows insertion of new data, but not modification of existing data.
- **Update** - allows modification, but not deletion of data.
- **Delete** - allows deletion of data.
- Forms of authorization to modify the database schema (covered in Chapter 8):
- **Index** - allows creation and deletion of indices.
- **Resources** - allows creation of new relations.
- **Alteration** - allows addition or deletion of attributes in a relation.
- **Drop** - allows deletion of relations.
- **Authorization Specification in SQL**
- The **grant** statement is used to confer authorization
 - **grant** <privilege list>
 - **on** <relation name or view name> **to** <user list>
- <user list> is:

- a user-id
- **public**, which allows all valid users the privilege granted
- Granting a privilege on a view does not imply granting any privileges on the underlying relations.
- The grantor of the privilege must already hold the privilege on the specified item (or be the database administrator).

Privileges in SQL

- **select**: allows read access to relation, or the ability to query using the view
 - Example: grant users U_1 , U_2 , and U_3 **select** authorization on the *branch* relation:
 - **grant select on branch to U_1 , U_2 , U_3**
- **insert**: the ability to insert tuples
- **update**: the ability to update using the SQL update statement
- **delete**: the ability to delete tuples.
- **all privileges**: used as a short form for all the allowable privileges

Revoking Authorization in SQL

- The **revoke** statement is used to revoke authorization.
- **revoke** <privilege list>
- **on** <relation name or view name> **from** <user list>
- Example:
 - **revoke select on branch from U_1 , U_2 , U_3**
- <privilege-list> may be **all** to revoke all privileges the revokee may hold.
- If <revokee-list> includes **public**, all users lose the privilege except those granted it explicitly.
- If the same privilege was granted twice to the same user by different grantees, the user may retain the privilege after the revocation.
- All privileges that depend on the privilege being revoked are also revoked.

Embedded SQL

- The SQL standard defines embeddings of SQL in a variety of programming languages such as C, Java, and Cobol.
- A language to which SQL queries are embedded is referred to as a **host language**, and the SQL structures permitted in the host language comprise *embedded SQL*.
- The basic form of these languages follows that of the System R embedding of SQL into PL/I.
- **EXEC SQL** statement is used to identify embedded SQL request to the preprocessor
 - EXEC SQL <embedded SQL statement> END_EXEC

Note: this varies by language (for example, the Java embedding uses

SQL { } ;)

Example Query

- From within a host language, find the names and cities of customers with more than the variable amount dollars in some account.
- Specify the query in SQL and declare a *cursor* for it

EXEC SQL

declare *c* **cursor** **for**

select *customer_name, customer_city*

from *depositor, customer, account*

where *depositor.customer_name = customer.customer_name*

and *depositor account_number = account.account_number*

and *account.balance > :amount*

END_EXEC

- The **open** statement causes the query to be evaluated

EXEC SQL **open** *c* END_EXEC
- The **fetch** statement causes the values of one tuple in the query result to be placed on host language variables.

EXEC SQL **fetch** *c* **into** *:cn, :cc* END_EXEC

Repeated calls to **fetch** get successive tuples in the query result

- A variable called SQLSTATE in the SQL communication area (SQLCA) gets set to '02000' to indicate no more data is available

- The **close** statement causes the database system to delete the temporary relation that holds the result of the query.

EXEC SQL **close** *c* END_EXEC

Note: above details vary with language. For example, the Java embedding defines Java iterators to step through result tuples.

Updates Through Cursors

- Can update tuples fetched by cursor by declaring that the cursor is for update

declare *c* **cursor for**

select *

from *account*

where *branch_name* = 'Perryridge'

for update

- To update tuple at the current location of cursor *c*

update *account*

set *balance* = *balance* + 100

where current of *c*

Dynamic SQL

- Allows programs to construct and submit SQL queries at run time.
- Example of the use of dynamic SQL from within a C program.

char * *sqlprog* = "**update** *account*

set *balance* = *balance* * 1.05

where *account_number* = ?"

EXEC SQL **prepare** *dynprog* **from** *:sqlprog*;

char *account* [10] = "A-101";

EXEC SQL **execute** *dynprog* **using** *:account*;

- The dynamic SQL program contains a ?, which is a place holder for a value that is provided when the SQL program is executed.

ODBC and JDBC

- API (application-program interface) for a program to interact with a database server
- Application makes calls to
 - Connect with the database server
 - Send SQL commands to the database server
 - Fetch tuples of result one-by-one into program variables
- ODBC (Open Database Connectivity) works with C, C++, C#, and Visual Basic
- JDBC (Java Database Connectivity) works with Java

ODBC

- Open DataBase Connectivity(ODBC) standard
 - standard for application program to communicate with a database server.
 - application program interface (API) to
 - open a connection with a database,
 - send queries and updates,
 - get back results.
- Applications such as GUI, spreadsheets, etc. can use ODBC
- Each database system supporting ODBC provides a "driver" library that must be linked with the client program.
- When client program makes an ODBC API call, the code in the library communicates with the server to carry out the requested action, and fetch results.
- ODBC program first allocates an SQL environment, then a database connection handle.
- Opens database connection using `SQLConnect()`. Parameters for `SQLConnect`:
 - connection handle,
 - the server to which to connect
 - the user identifier,
 - password
- Must also specify types of arguments:
 - `SQL_NTS` denotes previous argument is a null-terminated string.

ODBC Code

```
int ODBCexample()
{
    RETCODE error;
    HENV  env; /* environment */
    HDBC  conn; /* database connection */
    SQLAllocEnv(&env);
    SQLAllocConnect(env, &conn);
    SQLConnect(conn, "aura.bell-labs.com", SQL_NTS, "avi", SQL_NTS, "avipasswd",
    SQL_NTS);
    { .... Do actual work ... }
    SQLDisconnect(conn);
    SQLFreeConnect(conn);
    SQLFreeEnv(env);
}
```

- Program sends SQL commands to the database by using `SQLExecDirect`
- Result tuples are fetched using `SQLFetch()`
- `SQLBindCol()` binds C language variables to attributes of the query result
 - When a tuple is fetched, its attribute values are automatically stored in corresponding C variables.
 - Arguments to `SQLBindCol()`
 - ODBC stmt variable, attribute position in query result
 - The type conversion from SQL to C.
 - The address of the variable.
 - For variable-length types like character arrays,
 - The maximum length of the variable
 - Location to store actual length when a tuple is fetched.
 - Note: A negative value returned for the length field indicates null value
- Good programming requires checking results of every function call for errors; we have omitted most checks for brevity.

- Main body of program

```
char branchname[80];
float balance;
int lenOut1, lenOut2;
HSTMT stmt;
SQLAllocStmt(conn, &stmt);
char * sqlquery = "select branch_name, sum (balance)
                  from account
                  group by branch_name";
error = SQLExecDirect(stmt, sqlquery, SQL_NTS);
if (error == SQL_SUCCESS) {
    SQLBindCol(stmt, 1, SQL_C_CHAR, branchname, 80, &lenOut1);
    SQLBindCol(stmt, 2, SQL_C_FLOAT, &balance, 0, &lenOut2);
    while (SQLFetch(stmt) >= SQL_SUCCESS) {
        printf (" %s %g\n", branchname, balance);
    }
}
SQLFreeStmt(stmt, SQL_DROP);
```

More ODBC Features

- **Prepared Statement**
 - SQL statement prepared: compiled at the database
 - Can have placeholders: E.g. insert into account values(?,?,?)
 - Repeatedly executed with actual values for the placeholders
- **Metadata features**
 - finding all the relations in the database and
 - finding the names and types of columns of a query result or a relation in the database.
- By default, each SQL statement is treated as a separate transaction that is committed automatically.
 - Can turn off automatic commit on a connection

- `SQLSetConnectOption(conn, SQL_AUTOCOMMIT, 0)}`
- transactions must then be committed or rolled back explicitly by
 - `SQLTransact(conn, SQL_COMMIT)` or
 - `SQLTransact(conn, SQL_ROLLBACK)`

ODBC Conformance Levels

- Conformance levels specify subsets of the functionality defined by the standard.
 - Core
 - Level 1 requires support for metadata querying
 - Level 2 requires ability to send and retrieve arrays of parameter values and more detailed catalog information.
- SQL Call Level Interface (CLI) standard similar to ODBC interface, but with some minor differences.

JDBC

- **JDBC** is a Java API for communicating with database systems supporting SQL
- JDBC supports a variety of features for querying and updating data, and for retrieving query results
- JDBC also supports metadata retrieval, such as querying about relations present in the database and the names and types of relation attributes
- Model for communicating with the database:
 - Open a connection
 - Create a “statement” object
 - Execute queries using the Statement object to send queries and fetch results
 - Exception mechanism to handle errors

JDBC Code

```
public static void JDBCexample(String dbid, String userid, String passwd)
{
    try {
        Class.forName ("oracle.jdbc.driver.OracleDriver");
```

```

Connection conn = DriverManager.getConnection( "jdbc:oracle:thin:@aura.bell-
labs.com:2000:bankdb", userid, passwd);

    Statement stmt = conn.createStatement();

    ... Do Actual Work ....

    stmt.close();
    conn.close();
}
catch (SQLException sqle) {
    System.out.println("SQLException : " + sqle);
}
}

```

- **Update to database**

```

try {
    stmt.executeUpdate( "insert into account values
                        ('A-9732', 'Perryridge', 1200)");
} catch (SQLException sqle) {
    System.out.println("Could not insert tuple. " + sqle);
}

```

- **Execute query and fetch and print results**

```

ResultSet rset = stmt.executeQuery( "select branch_name, avg(balance)
                                     from account
                                     group by branch_name");

while (rset.next()) {
    System.out.println(
        rset.getString("branch_name") + " " + rset.getFloat(2));
}

```

JDBC Code Details

- Getting result fields:

- `rs.getString("branchname")` and `rs.getString(1)` equivalent if `branchname` is the first argument of select result.
- Dealing with Null values


```
int a = rs.getInt("a");
if (rs.isNull()) Systems.out.println("Got null value");
```

Procedural Extensions and Stored Procedures

- SQL provides a module language
 - Permits definition of procedures in SQL, with if-then-else statements, for and while loops, etc.
- Stored Procedures
 - Can store procedures in the database
 - then execute them using the call statement
 - permit external applications to operate on the database without knowing about internal details

Functions and Procedures

SQL supports functions and procedures

- Functions/procedures can be written in SQL itself, or in an external programming language
- Functions are particularly useful with specialized data types such as images and geometric objects
 - ▶ Example: functions to check if polygons overlap, or to compare images for similarity
- Some database systems support **table-valued functions**, which can return a relation as a result
- SQL also supports a rich set of imperative constructs, including
 - Loops, if-then-else, assignment

SQL Functions

- Define a function that, given the name of a customer, returns the count of the number of accounts owned by the customer.

```
create function account_count (customer_name varchar(20))
```

```
returns integer
```

```
begin
```

```
declare a_count integer;
```

```
select count (*) into a_count
```

```
from depositor
```

```
where depositor.customer_name = customer_name
```

```
return a_count;
```

```
end
```

- Find the name and address of each customer that has more than one account.

```
select customer_name, customer_street, customer_city
```

```
from customer
```

```
where account_count (customer_name) > 1
```

Table Functions

- SQL added functions that return a relation as a result
- Example: Return all accounts owned by a given customer

```
create function accounts_of (customer_name char(20))
```

```
returns table ( account_number char(10),
```

```
branch_name char(15),
```

```
balance numeric(12,2))
```

```
return table
```

```
(select account_number, branch_name, balance
```

```
from account
```

```
where exists (
```

```
select *
```

```
from depositor
```

where *depositor.customer_name = accounts_of.customer_name*
and *depositor.account_number = account.account_number*))

- Usage

select *
from table (*accounts_of* ('Smith'))

SQL Procedures

- The *author_count* function could instead be written as procedure:
create procedure *account_count_proc* (**in** *title* **varchar**(20), **out** *a_count* **integer**)

begin

select count(*author*) **into** *a_count*
from *depositor*
where *depositor.customer_name = account_count_proc.customer_name*

end

- Procedures can be invoked either from an SQL procedure or from embedded SQL, using the **call** statement.

declare *a_count* **integer**;
call *account_count_proc*('Smith', *a_count*);

Procedures and functions can be invoked also from dynamic SQL

- SQL allows more than one function/procedure of the same name (called name **overloading**), as long as the number of arguments differ, or at least the types of the arguments differ

Procedural Constructs

- Compound statement: **begin ... end**,
 - May contain multiple SQL statements between **begin** and **end**.
 - Local variables can be declared within a compound statements
- **While** and **repeat** statements:

declare *n* **integer** **default** 0;
while *n* < 10 **do**

```

        set  $n = n + 1$ 
    end while
repeat
set  $n = n - 1$ 
until  $n = 0$ 
end repeat

```

- **For loop**
 - Permits iteration over all results of a query
 - Example: find total of all balances at the Perryridge branch

```

declare  $n$  integer default 0;
for  $r$  as
    select  $balance$  from  $account$ 
    where  $branch\_name = 'Perryridge'$ 
do
    set  $n = n + r.balance$ 
end for

```

- Conditional statements (**if-then-else**)
E.g. To find sum of balances for each of three categories of accounts (with balance <1000, >=1000 and <5000, >= 5000)

```

    if  $r.balance < 1000$ 
    then set  $l = l + r.balance$ 
elseif  $r.balance < 5000$ 
    then set  $m = m + r.balance$ 
else set  $h = h + r.balance$ 
end if

```

- SQL also supports a **case** statement similar to C case statement
- Signaling of exception conditions, and declaring handlers for exceptions

```

declare  $out\_of\_stock$  condition
declare exit handler for  $out\_of\_stock$ 

```

```

begin
...
.. signal out-of-stock
end

```

- The handler here is **exit** -- causes enclosing **begin..end** to be exited
- Other actions possible on exception

External Language Functions/Procedures

- SQL permits the use of functions and procedures written in other languages such as C or C++
- Declaring external language procedures and functions

```

create procedure account_count_proc(in customer_name varchar(20), out count
integer)

```

language C

```

external name ' /usr/avi/bin/account_count_proc'

```

```

create function account_count(customer_name varchar(20))

```

returns integer

language C

```

external name ' /usr/avi/bin/author_count'

```

- Benefits of external language functions/procedures:
 - more efficient for many operations, and more expressive power
- Drawbacks
 - Code to implement function may need to be loaded into database system and executed in the database system's address space
 - risk of accidental corruption of database structures
 - security risk, allowing users access to unauthorized data
 - There are alternatives, which give good security at the cost of potentially worse performance

- Direct execution in the database system's space is used when efficiency is more important than security

Security with External Language Routines

- To deal with security problems
 - Use **sandbox** techniques
 - ▶ that is use a safe language like Java, which cannot be used to access/damage other parts of the database code
 - Or, run external language functions/procedures in a separate process, with no access to the database process' memory
 - ▶ Parameters and results communicated via inter-process communication
- Both have performance overheads
- Many database systems support both above approaches as well as direct executing in database system address space

Recursion in SQL

- SQL permits recursive view definition
- Example: find all employee-manager pairs, where the employee reports to the manager directly or indirectly (that is manager's manager, manager's manager's manager, etc.)

```
with recursive empl (employee_name, manager_name) as (
    select employee_name, manager_name
    from   manager
union
    select manager.employee_name, empl.manager_name
    from   manager, empl
    where manager.manager_name = empl.employee_name)
select *
from   empl
```

This example view, *empl*, is called the *transitive closure* of the *manager* relation

The Power of Recursion

- Recursive views make it possible to write queries, such as transitive closure queries, that cannot be written without recursion or iteration.
 - Intuition: Without recursion, a non-recursive non-iterative program can perform only a fixed number of joins of *manager* with itself
 - ▶ This can give only a fixed number of levels of managers
 - ▶ Given a program we can construct a database with a greater number of levels of managers on which the program will not work
 - The next slide shows a *manager* relation and each step of the iterative process that constructs *empl* from its recursive definition. The final result is called the *fixed point* of the recursive view definition.
- Recursive views are required to be *monotonic*. That is, if we add tuples to *manager* the view contains all of the tuples it contained before, plus possibly more

Example of Fixed-Point Computation

<i>employee_name</i>	<i>manager_name</i>
Alon	Barinsky
Barinsky	Estovar
Corbin	Duarte
Duarte	Jones
Estovar	Jones
Jones	Klinger
Rensal	Klinger

<i>Iteration number</i>	<i>Tuples in empl</i>
0	
1	(Duarte), (Estovar)
2	(Duarte), (Estovar), (Barinsky), (Corbin)
3	(Duarte), (Estovar), (Barinsky), (Corbin), (Alon)
4	(Duarte), (Estovar), (Barinsky), (Corbin), (Alon)

Advanced SQL Features

- Create a table with the same schema as an existing table:
create table *temp_account* **like** *account*

- SQL allows subqueries to occur *anywhere* a value is required provided the subquery returns only one value. This applies to updates as well
- SQL:2003 allows subqueries in the **from** clause to access attributes of other relations in the **from** clause using the **lateral** construct:

```
select customer_name, num_accounts
from customer, lateral (
    select count(*)
    from account
    where account.customer_name = customer.customer_name )
as this_customer (num_accounts)
```

- Merge construct allows batch processing of updates.
- Example: relation *funds_received* (*account_number, amount*) has batch of deposits to be added to the proper account in the *account* relation

```
merge into account as A
using (select *
    from funds_received as F )
on (A.account_number = F.account_number )
when matched then
    update set balance = balance + F.amount
```

DATABASE SYSTEM ARCHITECTURES


- **Centralized** and **Client-Server** Systems
- **Server** System Architectures
- **Parallel** Systems
- **Distributed** Systems
- Network Types

Centralized Systems


- Run on a **single** computer system and
 - do not interact with other computer systems.
- General-purpose computer system:

- one to a few **CPUs** and a number of device controllers that are
 - connected through a common **bus** that provides access to **shared** memory.
- Single-user system (e.g., personal computer or workstation):
 - desk-top unit, single user, usually has only one **CPU** and one or two hard **disks**;
 - the **OS** may support only **one** user.
- Multi-user system:
 - more **disks**, more **memory**, multiple **CPUs**, and a multi-user **OS**.
 - Serve a **large number** of users who are connected to the system via terminals
 - Often called *server* systems.

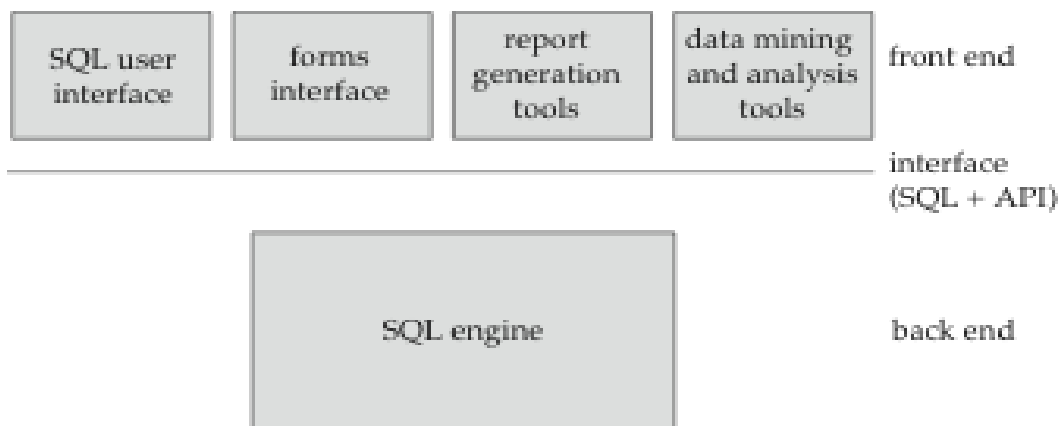
A Centralized Computer System



- **Client-Server Systems**
 - Server systems satisfy requests generated at m client systems,
 - whose general structure is shown below:



- Database *functionality* can be divided into:
 - Back-end: manages access structures, query evaluation and optimization, concurrency control and recovery.
 - Front-end: consists of tools such as *forms*, *report-writers*, and graphical user interface facilities.
- The interface between the front-end and the back-end is:
 - ▶ through SQL or through an Application Program Interface (API).



- **Advantages of replacing mainframes with:**
 - ▶ networks of workstations
 - ▶ or personal computers
 - connected to back-end server machines:
 - ▶ better functionality for the cost
 - ▶ flexibility in
 - locating resources and

- expanding facilities
- better user interfaces
- easier maintenance

Server System Architecture

- **Server** systems can be broadly categorized into **two** kinds:
 - **transaction servers** which are :
 - widely used in **relational** database systems,
 - **data servers**, used in **object-oriented** database systems

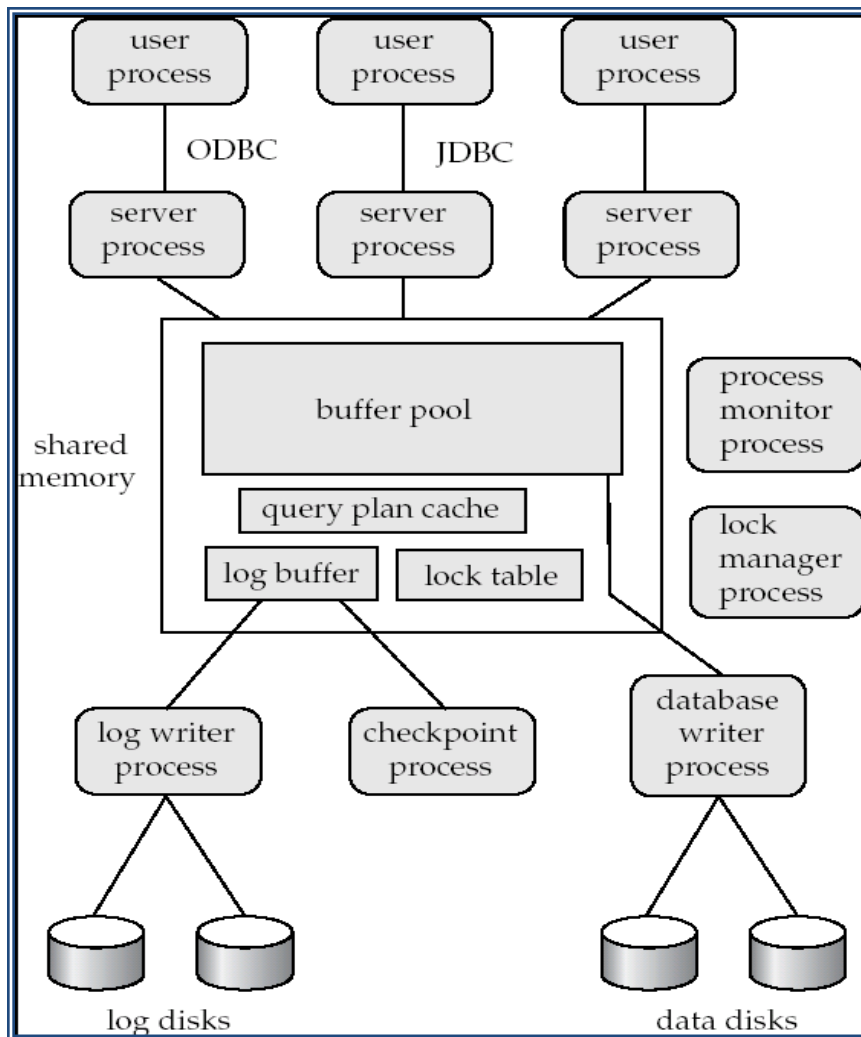
Transaction Servers

- Also called **query server** systems or **SQL server** systems
 - Clients send requests to the server
 - Transactions are executed at the server
 - **Results** are shipped back to the client.
- Requests are specified in SQL, and communicated to the server through a *remote procedure call* (**RPC**) mechanism.
- Transactional RPC **allows** many RPC calls to form a transaction.
- *Open Database Connectivity* (**ODBC**) is a **C** language application program interface standard from Microsoft for connecting to a server,
 - sending **SQL** requests, and receiving **results**.
 - **JDBC** standard is similar to ODBC, for Java

Transaction Server Process Structure

- A typical transaction **server** consists of :
 - multiple processes accessing data in *shared memory*.
- **Server** processes
 - These receive user queries (transactions),
 - execute them and **send** results back
 - Processes may be **multithreaded**,
 - allowing a single process to execute **several** user queries concurrently

- Typically multiple multithreaded server processes
- Lock manager process
 - More on this later
- Database writer process
 - **Output** modified buffer blocks to disks continually
- Log writer process
 - Server processes :
 - ▶ simply **add** log records to log record **buffer**
 - Log writer process:
 - ▶ **outputs** log records to stable storage.
- Checkpoint process
 - Performs periodic **checkpoints**
- Process monitor process
 - **Monitors** other processes, and :
 - ▶ takes **recovery** actions if any of the other processes **fail**
 - ▶ E.g. aborting any transactions being executed by a server process and restarting it



- **Shared memory** contains shared data
 - Buffer pool
 - Lock table
 - Log buffer
 - **Cached** query plans (reused if same query submitted again)
 - **All** database processes can access **shared memory**
- To ensure that **no** two processes are accessing :
 - ▶ the same data structure
 - ▶ at the same time,
 - databases systems implement:
 - ▶ **mutual exclusion** using either:

- Operating system semaphores
 - Atomic instructions such as test-and-set
- To **avoid** overhead of
 - interprocess communication for **lock** request/grant,
 - ▶ each database process operates **directly** on the **lock table**
 - ▶ instead of sending requests to lock manager process
- **Lock manager** process still used for:
 - **deadlock** detection

Data Servers

- Used in **high-speed** LANs, in cases where:
 - The clients are comparable in processing power to the server
 - The tasks to be executed are **compute** intensive.
- Data are shipped to clients where:
 - processing is performed, and
 - then shipped results back to the server.
- This architecture requires **full** back-end functionality at the clients.
- Used in many **object-oriented** database systems
- Issues:
 - Page-Shipping *versus* Item-Shipping
 - Locking
 - Data Caching
 - Lock Caching
- **Page-shipping** versus **item-shipping**
 - Smaller unit of shipping \Rightarrow more messages
 - Worth **prefetching** related items along with requested item
 - Page shipping can be thought of as a form of **prefetching**
- Locking
 - Overhead of **requesting** and **getting locks** from server is high due to message delays
 - Can grant locks on requested and prefetched items;

- ▶ with page shipping, transaction is granted lock on whole page.
- **Locks** on a prefetched item can be **called back** by the **server**,
 - ▶ and returned by **client transaction**:
 - if the prefetched item has **not** been used.
- Locks on the page can be **deescalated** to locks on items in the page :
 - ▶ when there are lock **conflicts**.
 - ▶ Locks on *unused items* can then be **returned** to **server**.
- **Data Caching**
 - **Data** can be **cached** at client even in between transactions
 - But **check** that **data** is up-to-date before it is used (**cache coherency**)
 - Check can be done when **requesting** lock on data item
- **Lock Caching**
 - **Locks** can be **retained** by client system even in between transactions
 - Transactions can **acquire** cached locks locally,
 - ▶ without **contacting server**
 - **Server calls back** locks from clients:
 - ▶ when it **receives conflicting** lock request.
 - ▶ Client **returns** lock once **no** local transaction is using it.
 - Similar to **deescalation**, but **across** transactions

Parallel Systems

- **Parallel** database systems consist of:
 - ▶ multiple **processors** and
 - ▶ multiple **disks**
 - connected by a **fast** interconnection network.
- A **coarse-grain parallel** machine consists of:
 - a small number of ***powerful processors***
- A **massively parallel** or **fine grain parallel** machine utilizes:
 - **thousands** of ***smaller processors***.
- **Two** main performance measures:
 - **Throughput**:

- ▶ the **number of tasks** that can be **completed** in a given **time interval**
- **response time:**
 - ▶ the **amount of time** it takes to **complete** a **single task**
 - from the time it is submitted

Speed-Up and Scale-Up


- **Speedup:**
 - a **fixed-sized** problem executing on a small system is:
 - ▶ given to a system which is **N-times** larger.
 - Measured by:

$$\text{speedup} = \frac{\text{small system elapsed time}}{\text{large system elapsed time}}$$


- Speedup is **linear** if **equation** equals N.
- **Scaleup:**
- **increase** the size of **both** the problem and the system
 - ▶ **N-times** larger system used to perform **N-times** larger job
 - ▶ Measured by:

$$\text{scaleup} = \frac{\text{small system small problem elapsed time}}{\text{big system big problem elapsed time}}$$

- Scale up is **linear** if **equation** equals 1.



Scaleup



Batch and Transaction Scaleup

- **Batch scaleup:**
 - A **single** large job;
 - typical of *most* database queries and **scientific** simulation.
 - Use an N -times **larger computer**
 - on N -times **larger problem**.
- **Transaction scaleup:**
 - **Numerous** small queries submitted by independent users to a shared database;
 - typical transaction processing and timesharing systems.
 - N -times as many users submitting requests (hence, N -times as many requests)
 - to an N -times larger database,
 - on an N -times larger computer.
 - Well-suited to **parallel** execution.

Factors Limiting *Speedup* and *Scaleup*

Speedup and scaleup are often sublinear due to:


- **Startup** costs:
 - Cost of **starting up** multiple processes may **dominate** computation time,
 - if the **degree** of parallelism is **high**.
- **Interference:**
 - Processes accessing shared resources (e.g., system bus, disks, or locks)

- ▶ **compete** with each other,
 - thus **spending time** waiting on other processes,
 - rather than **performing** useful work.
- **Skew:**
 - **Increasing** the degree of parallelism
 - ▶ **increases** the variance in service times of parallelly executing **tasks**.
 - ▶ Overall execution time determined by **slowest** parallel **tasks**.

Interconnection Network Architectures

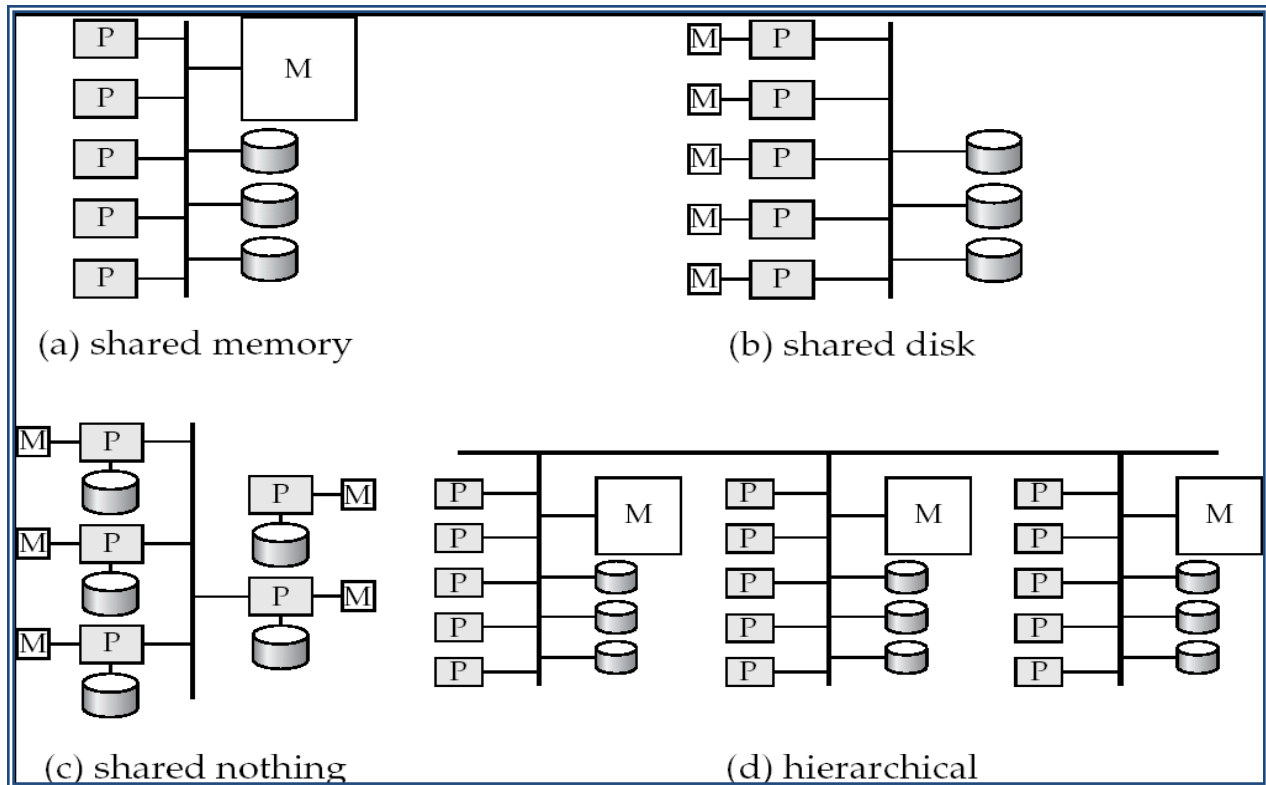
- **Bus.**
 - System components *send data* on and *receive data* from
 - ▶ a **single** communication bus;
 - Does **not** scale **well** with *increasing parallelism*.
- **Mesh.**
 - Components are arranged as **nodes** in a **grid**, and
 - ▶ each component is **connected** to all adjacent components
 - Communication **links grow** with *growing number of components*, and
 - ▶ so scales **better**.
 - But may require $2\sqrt{n}$ **hops** to send message to a node
 - (or \sqrt{n} with wraparound connections at edge of **grid**).
- **Hypercube.**
 - Components are numbered in binary;
 - ▶ components are connected to one another :
 - **if** their binary representations differ in exactly one bit.
 - n components are *connected* to $\log(n)$ other components and
 - ▶ can **reach** each other via at most $\log(n)$ links;
 - ▶ **reduces** communication delays.

Interconnection Architectures



Parallel Database Architectures

- **Shared memory** –
 - processors share a common memory
- **Shared disk** –
 - processors share a common disk
- **Shared nothing** –
 - processors share neither a common memory nor common disk
- **Hierarchical** –
 - **hybrid** of the above architectures



Shared *Memory*

- Processors and disks have access to a common memory,
 - typically via a bus
 - or through an interconnection network.
- Extremely **efficient** communication between processors —
 - data** in shared memory can be **accessed** by any processor
 - without **having to** move it using software.
- Downside – architecture is **not** scalable beyond 32 or 64 processors
 - since the bus or the interconnection network becomes a **bottleneck**
- Widely used for **lower** degrees of parallelism (**4 to 8**).

Shared *Disk*

- All** processors can directly **access all** disks
 - via an interconnection network,
 - but the processors **have private memories**.

- The memory bus is **not** a bottleneck
- Architecture provides a degree of **fault-tolerance** —
 - ▶ if a processor **fails**, the other processors can **take over** its tasks
 - ▶ since the database is resident on disks
 - ▶ that are accessible from all processors.
- Ex: IBM Sysplex and DEC clusters (now part of Compaq)
 - ▶ running Rdb (now Oracle Rdb) were early commercial users
- Downside:
 - **bottleneck** now occurs at:
 - ▶ interconnection to the disk subsystem.
- Shared-disk systems **can scale** to a somewhat **larger number** of processors,
 - but communication between processors is **slower**.

Shared *Nothing*


- Node consists of a processor, memory, and one or more disks.
- Processors at one **node** communicate with another processor at another **node**
 - using an interconnection network.
 - A **node** functions as the **server** for:
 - ▶ the **data** on the **disk** or disks the node **owns**.
- Ex: Teradata, Tandem, Oracle-n CUBE
- **Data** accessed from local disks (and local memory accesses)
 - do not pass through interconnection network,
 - thereby **minimizing** the interference of resource sharing.
- Shared-nothing multiprocessors :
 - can be **scaled up** to **thousands** of processors
 - ▶ without **interference**.
- Main **drawback**: cost of communication and non-local disk access;
 - sending data involves software interaction at both ends.

Hierarchical

- **Combines** characteristics of :
 - shared-memory, shared-disk, and shared-nothing architectures.
- **Top level** is a shared-nothing architecture –
 - **nodes** connected by an interconnection network, and
 - do **not** share disks or memory *with each other*.
- Each node of the system could be:
 - a shared-memory system with a few processors.
- Alternatively, each node could be:
 - a shared-disk system, and each of the systems sharing a set of disks
 - could be a shared-memory system.
- **Reduce the complexity** of programming such systems by:
 - **distributed virtual-memory** architectures
 - Also called **non-uniform** memory architecture (**NUMA**)

Distributed Systems

- **Data spread** over multiple machines
 - (also referred to as **sites** or **nodes**).
- Network interconnects the machines
- **Data shared** by users on multiple machines



Distributed Databases

- Homogeneous distributed databases
 - Same software/schema on all sites, data may be partitioned among sites
 - Goal: provide a view of a single database, hiding details of distribution
- Heterogeneous distributed databases
 - Different software/schema on different sites
 - Goal: integrate existing databases to provide useful functionality
- Differentiate between *local* and *global* transactions
 - A local transaction accesses data in the *single* site at which the transaction was initiated.
 - A global transaction either accesses data in a site different from the one at which the transaction was initiated or accesses data in several different sites.

Trade-offs in Distributed Systems

- Sharing data – users at one site able to access the data residing at some other sites.
- Autonomy – each site is able to retain a degree of control over data stored locally.
- Higher system availability through redundancy — data can be replicated at remote sites, and system can function even if a site fails.
- Disadvantage: added complexity required to ensure proper coordination among sites.
 - Software development cost.
 - Greater potential for bugs.
 - Increased processing overhead.

Implementation Issues for Distributed Databases

- Atomicity needed even for transactions that update data at multiple sites
- The two-phase commit protocol (2PC) is used to **ensure** atomicity
 - Basic idea: each site **executes** transaction until just before commit, and
 - ▶ then **leaves** final decision to a **coordinator**
 - Each site must **follow** decision of coordinator,
 - ▶ even if there is a **failure** while waiting for coordinators decision
- **2PC** is not always appropriate:

- other transaction models based on:
 - persistent messaging, and workflows, are **also used**.
- Distributed concurrency control (and deadlock detection) **required**
- Data items may be **replicated** to **improve** data availability

Network Types

- **Local-area networks (LANs)** –
 - composed of processors that are distributed over **small** geographical areas,
 - such as a **single building**
 - or a **few adjacent buildings**.
- **Wide-area networks (WANs)** –
 - composed of processors distributed over a **large** geographical area.
- WANs with **continuous** connection (e.g. the Internet) are:
 - **needed** for implementing **distributed** database systems
- **Groupware** applications such as *Lotus notes*:
 - can **work** on WANs with **discontinuous** connection:
 - Data is replicated.
 - Updates are propagated to replicas periodically.
 - Copies of data may be updated independently.
 - Non-serializable executions can thus result. Resolution is application dependent.

DISTRIBUTED DATABASES

- **Heterogeneous** and **Homogeneous** Databases
- Distributed **Data** Storage
- Distributed **Transactions**
- Commit **Protocols**
- **Concurrency** Control in Distributed Databases
- **Availability**
- Distributed **Query** Processing

- **Heterogeneous Distributed Databases**
- **Directory Systems**

Distributed Database System

- A **distributed database system** consists of **loosely coupled sites** that **share no physical component**
- **Database** systems that **run** on **each site** are **independent** of each other
- **Transactions** may **access data** at one or **more sites**

Homogeneous Distributed Databases

- In a **homogeneous** distributed database:
 - All sites have **identical software**
 - **Are aware** of **each other** and agree to cooperate in processing user requests.
 - **Each site surrenders** part of **its autonomy** in terms of right to change schemas or software
 - **Appears** to user as a **single system**
- In a **heterogeneous** distributed database
 - **Different sites** may use **different schemas** and software
 - ▶ Difference in schema is a **major problem** for **query processing**
 - ▶ Difference in software is a major problem for **transaction processing**
 - **Sites** may **not** be **aware** of each other and may **provide** only limited facilities for cooperation in transaction processing

Distributed Data Storage

- Assume **relational data model**
- **Replication:**
 - System maintains **multiple copies** of data, stored in different sites, for **faster retrieval** and **fault tolerance**.
- **Fragmentation:**
 - Relation is partitioned into **several fragments** stored in distinct sites
- Replication and fragmentation **can be combined:**

- Relation is partitioned into several fragments: system maintains **several** identical **replicas** of each such **fragment**.

Data Replication

- A relation or fragment of a relation is **replicated** if it is **stored redundantly** in two or more sites.
- **Full replication** of a relation is the case where **the relation** is stored at **all sites**.
- **Fully redundant** databases are those in which **every site** contains a **copy** of the **entire database**.
- **Advantages of Replication**
 - **Availability**: failure of site containing relation r does not result in unavailability of r if replicas exist.
 - **Parallelism**: queries on r may be processed by several nodes in parallel.
 - **Reduced data transfer**: relation r is available locally at each site containing a replica of r .
- **Disadvantages of Replication**
 - **Increased cost of updates**: each replica of relation r must be updated.
 - **Increased complexity** of concurrency control: concurrent updates to distinct replicas may lead to inconsistent data unless **special** concurrency control **mechanisms** are implemented.
 - ▶ One **solution**: **choose** one copy as **primary copy** and apply concurrency control operations on primary copy

Data Fragmentation

- **Division** of relation r into fragments r_1, r_2, \dots, r_n which contain sufficient information to reconstruct relation r .
- **Horizontal fragmentation**: each tuple of r is assigned to one or more fragments
- **Vertical fragmentation**: the schema for relation r is split into several smaller schemas
 - All schemas must contain a common candidate key (or **superkey**) to ensure **lossless join** property.

- A special attribute, the **tuple-id** attribute may be added to each schema to serve as a candidate key.

- **Example** : relation **account** with following schema
- $Account = (branch_name, account_number, balance)$

Horizontal Fragmentation of *account* Relation

<i>branch_name</i>	<i>account_number</i>	<i>balance</i>
Hillside	A-305	500
Hillside	A-226	336
Hillside	A-155	62

$$account_1 = \sigma_{branch_name = \text{"Hillside"}}(account)$$

<i>branch_name</i>	<i>account_number</i>	<i>balance</i>
Valleyview	A-177	205
Valleyview	A-402	10000
Valleyview	A-408	1123
Valleyview	A-639	750

$$account_2 = \sigma_{branch_name = \text{"Valleyview"}}(account)$$

Vertical Fragmentation of *employee_info* Relation

<i>branch_name</i>	<i>customer_name</i>	<i>tuple_id</i>
Hillside	Lowman	1
Hillside	Camp	2
Valleyview	Camp	3
Valleyview	Kahn	4
Hillside	Kahn	5
Valleyview	Kahn	6
Valleyview	Green	7

$$deposit_1 = \Pi_{branch_name, customer_name, tuple_id}(employee_info)$$

<i>account_number</i>	<i>balance</i>	<i>tuple_id</i>
A-305	500	1
A-226	336	2
A-177	205	3
A-402	10000	4
A-155	62	5
A-408	1123	6
A-639	750	7

$$deposit_2 = \Pi_{account_number, balance, tuple_id}(employee_info)$$

Advantages of Fragmentation

- **Horizontal:**
 - **allows parallel processing** on fragments of a relation
 - **allows a relation** to be **split** so that **tuples** are **located** where they are most frequently accessed
- **Vertical:**
 - **allows tuples** to be **split** so that each part of the tuple is stored where it is most frequently accessed
 - **tuple-id** attribute allows **efficient joining** of vertical fragments
 - **allows parallel processing** on a relation
- Vertical and horizontal fragmentation can be **mixed**.
 - Fragments may be successively fragmented to an arbitrary depth.

Data Transparency

- **Data transparency:** Degree to which system **user** may remain **unaware** of the details of how and where the **data** items are **stored** in a distributed system
- Consider transparency **issues** in relation to:
 - Fragmentation transparency
 - Replication transparency
 - Location transparency

Naming of Data Items – Criteria

1. Every data item must have a **system-wide unique name**.
2. It should be possible to **find** the location of **data items efficiently**.
3. It should be possible to **change** the location of **data items transparently**.
4. Each site should be **able** to **create** new **data items autonomously**.

Centralized Scheme - Name Server

- **Structure:**
 - name **server assigns all** names
 - each **site maintains** a record of **local** data items
 - **sites ask name server** to **locate non-local** data items
- **Advantages:**
 - **satisfies** naming criteria 1-3
- **Disadvantages:**
 - does **not** satisfy naming criterion 4
 - name **server** is a potential performance **bottleneck**
 - name **server** is a single point of **failure**

Use of Aliases

- **Alternative** to centralized scheme: each site **prefixes** its own site identifier to any name that it generates **i.e., *site 17.account***.
 - **Fulfills** having a **unique** identifier, and **avoids** problems associated with central control.
 - However, **fails** to achieve network **transparency**.


- **Solution:** Create a set of **aliases** for data items; Store the mapping of aliases to the **real names** at each site.

The **user** can be **unaware** of the physical location of a data item, and is unaffected if the data item is moved from one site to another.

Distributed Transactions

- **Transaction** may **access data** at several sites.
- Each site has a **local transaction manager** responsible for:
 - **Maintaining a log** for recovery purposes
 - **Participating** in coordinating the concurrent execution of the transactions executing at that site.
- Each site has a **transaction coordinator**, which is responsible for:
 - **Starting the execution** of transactions that originate at the site.
 - **Distributing subtransactions** at appropriate sites for execution.
 - **Coordinating the termination** of each transaction that originates at the site, which may result in the transaction being committed at **all** sites or aborted at **all** sites.

Transaction System Architecture



System Failure Modes

- Failures unique to distributed systems:
 - Failure of a site.
 - Loss of messages
 - ▶ Handled by network transmission control protocols such as TCP-IP
 - Failure of a communication link
 - ▶ Handled by network protocols, by routing messages via alternative links
 - Network partition
 - ▶ A network is said to be partitioned when it has been split into two or more subsystems that lack any connection between them
 - Note: a subsystem may consist of a single node

- Network partitioning and site failures are generally indistinguishable.

Commit Protocols

- Commit protocols are used to **ensure atomicity** across sites
 - a transaction which executes at multiple sites **must** either be **committed** at **all** the **sites**, or **aborted** at **all** the **sites**.
 - **not acceptable** to have a transaction committed at one site and aborted at another
- The *two-phase commit* (2PC) protocol is **widely** used
- The *three-phase commit* (3PC) protocol is more **complicated** and more **expensive**, but avoids some drawbacks of two-phase commit protocol. This protocol is **not used** in practice.

Two Phase Commit Protocol (2PC)

- Assumes **fail-stop model** – failed sites simply stop working, and do **not cause any other harm**, such as sending incorrect messages to other sites.
- Execution of the protocol is **initiated** by the **coordinator** after the last step of the transaction has been reached.
- The protocol **involves all** the local **sites** at which the transaction executed
- Let T be a transaction initiated at site S_i , and let the transaction **coordinator** at S_i be C_i

Phase 1: Obtaining a Decision

- Coordinator C_i **asks** all participants to *prepare* to **commit** transaction T_i .
 - C_i **adds** the records $\langle \text{prepare } T \rangle$ to the log and **forces** log to stable storage
 - **sends prepare** T messages to **all** sites at which T executed
- Upon receiving message, **transaction manager** at each site **determines** if it can **commit** the transaction
 - **if not**, **adds** a record $\langle \text{no } T \rangle$ to the log and **sends abort** T message to C_i
 - **if the transaction can be committed**, then:
 - **adds** the record $\langle \text{ready } T \rangle$ to the log
 - **forces all records** for T to stable storage
 - **sends ready** T message to C_i

Phase 2: Recording the Decision

- T can be committed **if** C_i received a **ready** T message from **all** the participating **sites**: otherwise T must be aborted.
- Coordinator C_i **adds** a decision record, $\langle \text{commit } T \rangle$ or $\langle \text{abort } T \rangle$, to the log and **forces** record onto stable storage. Once the record on stable storage it is **irrevocable** (even if failures occur)
- Coordinator C_i **sends** a message to each participant informing it of the decision (commit or abort)
- Participants take appropriate **action** locally.

Handling of Failures - Site Failure

When site S_k recovers, it **examines** its log to determine the **fate** of transactions active at the time of the failure.

- Log contain $\langle \text{commit } T \rangle$ record: site executes **redo** (T)

- Log contains **<abort T >** record: site executes **undo (T)**
- Log contains **<ready T >** record: site **must consult C_i** to determine the **fate** of T .
 - If T committed, **redo (T)**
 - If T aborted, **undo (T)**
- The log contains **no control records** concerning T replies that S_k failed before responding to the **prepare T** message from C_i
 - since the failure of S_k precludes the sending of such a response **C_i must abort T**
 - **S_k must execute undo (T)**

Handling of Failures- Coordinator Failure

- If coordinator C_i **fails** while the commit protocol for T is executing then **participating sites must decide on T 's fate**:
 1. If an active site **contains** a **<commit T >** record in its log, then **T must** be committed.
 2. If an active site **contains** an **<abort T >** record in its log, then **T must** be aborted.
 3. If some active participating site does **not contain** a **<ready T >** record in its log, then the failed coordinator C_i **cannot have decided** to commit T . Can therefore **abort T** .
 4. If none of the above cases holds, then **all** active sites must **have** a **<ready T >** record in their logs, but **no additional control records** (such as **<abort T >** or **<commit T >**). In this case active sites **must** wait for C_i to recover, to find decision.
- **Blocking problem** : active sites may have to wait for failed coordinator C_i to recover.

Handling of Failures - Network Partition

- If the **coordinator** and all its **participants** remain in **one partition**, the failure has no effect on the commit protocol.
- If the **coordinator** and its participants belong to **several partitions**:
 - **Sites** that are **not in** the partition containing the **coordinator** **think** the coordinator has failed, and execute the protocol to deal with failure of the coordinator.
 - ▶ **No harm results**, but sites may still **have to wait** for decision from **coordinator**.
- The **coordinator** (and the sites that are in the **same partition** as the coordinator) **think** that the **sites** in the **other partition have failed**, and follow the usual commit protocol.
 - ▶ Again, **no harm results**

Recovery and Concurrency Control

- **In-doubt transactions** have a **<ready T >**, but neither a **<commit T >**, nor an **<abort T >** log record.
 - The **recovering site must determine** the **commit-abort status** of such transactions by contacting other sites; this **can slow** and potentially **block recovery**.
- **Recovery algorithms** can **note** lock information in the log.

- Instead of **<ready T >**, write out **<ready T, L >** L = **list of locks** held by T when the log is written (read locks can be omitted).
- For every in-doubt transaction T , **all the locks noted** in the **<ready T, L >** log record **are reacquired**.
- After **lock reacquisition**, transaction processing can resume; the commit or rollback of in-doubt transactions is performed **concurrently** with the execution of **new transactions**.

Alternative Models of Transaction Processing

- Notion of a **single transaction** spanning multiple sites **is inappropriate** for many applications
 - E.g. transaction **crossing an organizational boundary**
 - No organization would **like to permit an externally initiated transaction** to **block local transactions** for an indeterminate period
- **Alternative models carry out transactions by sending messages**
 - Code to **handle messages** must be carefully designed to **ensure atomicity** and **durability** properties for updates
 - **Isolation cannot** be guaranteed, in that **intermediate stages** are **visible**, but code **must ensure no inconsistent states** result due to concurrency
 - **Persistent messaging systems** are systems that **provide transactional properties** to messages
 - **Messages** are guaranteed to be **delivered** exactly **once**
 - Will discuss implementation **techniques** later
- Motivating example: **funds transfer** between two banks
 - **Two phase commit** would have the potential to **block updates** on the accounts involved in funds transfer
 - Alternative solution:
 - **Debit** money from source account **and send a message** to other site
 - Site **receives message and** credits destination account
 - **Messaging** has long been used for **distributed transactions** (even before computers were invented!)
- **Atomicity** issue:
 - **Once transaction** sending a message is **committed**, **message must be** guaranteed to **be delivered**
 - Guarantee as long as **destination site** is **up and reachable**, code to **handle undeliverable messages** must also be available
 - n e.g. **credit money back** to source account.
 - If sending **transaction aborts**, **message must not** be sent

Error Conditions with Persistent Messaging

- Code to handle messages has to **take care** of **variety of failure** situations (even **assuming** guaranteed message delivery)
 - E.g. if destination account does **not exist**, failure message **must be sent back** to source site
 - When failure message is received from destination site, **or** destination site itself does **not exist**, **money must be deposited back** in source account
 - **Problem** if source account has been **closed**

- **get humans to take care** of problem
- **User** code executing transaction processing using **2PC** does **not have** to **deal** with such failures
- There are many situations where **extra effort** of **error handling is worth** the benefit of **absence of blocking**
 - E.g. pretty much all **transactions across** organizations

Persistent Messaging and Workflows

- **Workflows** provide a **general model** of transactional processing involving **multiple sites** and possibly **human** processing of certain steps
 - E.g. when a bank receives a **loan** application, it may need to
 - ▶ **Contact** external **credit-checking** agencies
 - ▶ Get **approvals** of one or more **managers** and then **respond** to the **loan** application
 - **Persistent messaging** forms the **underlying** infrastructure for **workflows** in a distributed environment

Concurrency Control

- Modify **concurrency control schemes** for use in **distributed environment**.
- We **assume** that **each site** participates in the execution of a **commit** protocol to **ensure global** transaction **atomicity**.
- We **assume** all **replicas** of any item are updated
 - Will see how to **relax** this in case of site failures later

Single Lock-Manager Approach

- System maintains a **single lock manager** that resides in a **single** chosen **site**, say S_i
- When a transaction needs to **lock** a **data item**, it **sends** a **lock request** to S_i and **lock manager determines** whether the lock can be **granted** immediately
 - If **yes**, lock manager **sends** a **message** to the site which initiated the request
 - If **no**, **request is delayed** until **it can be** granted, at which time a message is sent to the initiating site
- The transaction **can read** the data item from **any** one of the **sites** at which a replica of the data item resides.
- **Writes must** be performed on **all replicas** of a data item
- **Advantages** of scheme:
 - Simple **implementation**
 - Simple **deadlock** handling
- **Disadvantages** of scheme are:
 - **Bottleneck**: lock manager **site** becomes a **bottleneck**
 - **Vulnerability**: system is vulnerable to lock manager **site failure**.

Distributed Lock Manager

- In this approach, functionality of **locking** is implemented by **lock managers** at **each site**
 - Lock managers **control access** to **local data** items
 - ▶ But **special protocols** may be used for **replicas**
- **Advantage**: work is **distributed** and can be made **robust** to failures

- **Disadvantage: deadlock detection** is more complicated
 - Lock managers **cooperate** for deadlock detection
 - More on this later
- Several **variants** of this approach
 - **Primary** copy
 - **Majority** protocol
 - **Biased** protocol
 - **Quorum** consensus

Primary Copy

- **Choose one replica** of data item to be the **primary copy**.
 - **Site containing** the replica is called the **primary site** for that data item
 - **Different data items** can have **different primary sites**
- When a transaction needs to **lock** a data item Q , it requests a lock at the **primary site** of Q .
 - Implicitly gets lock on **all replicas** of the data item
- **Benefit:**
 - **Concurrency control** for replicated data handled **similarly** to unreplicated data - simple implementation.
- **Drawback:**
 - If the **primary site** of Q fails, Q is **inaccessible** even though **other** sites containing a **replica** may be accessible.

Majority Protocol

- **Local lock manager** at each site administers **lock** and **unlock** requests for data items stored at **that site**.
- When a **transaction** wishes to lock an **unreplicated** data item Q residing at site S_i , a message is sent to S_i 's **lock manager**.
 - If Q is **locked** in an **incompatible** mode, then the **request** is **delayed** until it can be granted.
 - **When the lock request can be granted**, the lock manager **sends** a message back to the initiator indicating that the lock request has been granted.
- In case of **replicated data**
 - If Q is replicated at **n sites**, then a lock request **message** must be **sent to more than half** of the **n sites** in which Q is stored.
 - The transaction does **not operate** on Q **until** it has **obtained** a lock on a **majority** of the replicas of Q .
 - When **writing** the data item, transaction performs **writes** on **all replicas**.
- **Benefit**
 - **Can** be used even when some **sites** are **unavailable**
 - n details on how handle **writes** in the presence of **site failure** later
- **Drawback**
 - **Requires $2(n/2 + 1)$ messages** for handling **lock** requests, and **$(n/2 + 1)$ messages** for handling **unlock** requests.
 - **Potential for deadlock** even with **single item** - e.g., each of **3** transactions may have **locks** on **1/3rd of the replicas** of a data.

Biased Protocol


- **Local lock manager** at each site as in majority protocol, however, **requests for shared locks** are handled differently than requests for **exclusive locks**.
- **Shared locks**. When a transaction needs to lock data item Q , it simply requests a lock on Q from the lock manager at **one site** containing a replica of Q .
- **Exclusive locks**. When transaction needs to lock data item Q , it requests a lock on Q from the lock manager at **all sites** containing a replica of Q .
- Advantage - imposes **less overhead** on **read** operations.
- Disadvantage - **additional overhead** on **writes**

Quorum Consensus Protocol

- A **generalization** of both **majority** and **biased** protocols
- **Each site** is assigned a **weight**.
 - Let S be the **total** of **all site weights**
- Choose two values **read quorum** Q_r and **write quorum** Q_w
 - Such that $Q_r + Q_w > S$ and $2 * Q_w > S$
 - **Quorums** can be chosen (and S computed) separately for **each item**
- **Each read** must lock enough replicas that the **sum** of the site weights is $\geq Q_r$
- **Each write** must lock enough replicas that the **sum** of the site weights is $\geq Q_w$
- For now we assume **all replicas** are **written**
 - Extensions to allow some **sites** to be **unavailable** described later

Timestamping

- **Timestamp based** concurrency-control protocols can be used in **distributed** systems
- Each **transaction** must be given a **unique** timestamp
- Main **problem**: how to **generate** a **timestamp** in a distributed fashion:
 - **Each site** generates a **unique local** timestamp using either a logical **counter** or the local **clock**.
 - **Global unique** timestamp is **obtained** by **concatenating** the unique **local** timestamp with the unique **identifier**.



- A site S_i with a **slow clock** will assign smaller timestamps
 - Still logically correct: **serializability not affected**
 - **But**: “**disadvantages**” transactions of the site S_i
- To **fix** this problem:

- Define within each site S_i a **logical clock** (LC_i), which generates the **unique local** timestamp
- Require that S_i **advance** its logical clock whenever a request is received from a transaction T_i with timestamp $\langle x, y \rangle$ and x is greater than the current value of LC_i .
- In this case, site S_i **advances** its logical clock to the value $x + 1$.

Replication with Weak Consistency

- Many **commercial** databases **support** replication of data with **weak** degrees of **consistency** (I.e., **without** a guarantee of **serializability**)
- **E.g.: master-slave replication:** **updates** are performed at a **single “master”** site, and propagated to “**slave**” sites.
 - **Propagation** is not part of the **update transaction**: it is decoupled
 - May be **immediately** after transaction **commits**
 - May be **periodic**
 - Data may **only** be **read** at **slave sites**, not updated
 - **No need** to obtain **locks** at any remote site
 - Particularly **useful** for **distributing information**
 - E.g. from central office to branch-office
 - Also **useful** for running **read-only** queries **offline** from the main database
- Replicas **should** see a **transaction-consistent snapshot** of the database
 - That is, a **state of the database** reflecting **all** effects of **all** transactions up to **some point** in the serialization order, and no effects of any later transactions.
- E.g. **Oracle** provides a **create snapshot** statement to create a snapshot of a relation or a set of relations at a **remote** site:
 - **Snapshot** refresh either by **recomputation** or by **incremental** update
 - **Automatic** refresh (**continuous** or **periodic**) or **manual** refresh

Multimaster and Lazy Replication

- With **multimaster replication** (also called **update-anywhere** replication) **updates** are **permitted at any replica**, and are automatically **propagated** to **all** replicas
 - **Basic model:** in distributed databases, where:
 - **transactions** are **unaware** of the details of replication, and
 - **database** system **propagates** updates **as part of** the same transaction
 - Coupled with 2 phase commit
- Many systems **support lazy propagation** where:
 - **Updates** are **transmitted** after transaction commits
 - **Allows** updates to occur even if some **sites** are **disconnected** from the network,
 - **But** at the cost of **consistency!**

Deadlock Handling

Consider the following **two transactions** and history, with item **X** and transaction T_1 at site 1, and item **Y** and transaction T_2 at site 2:

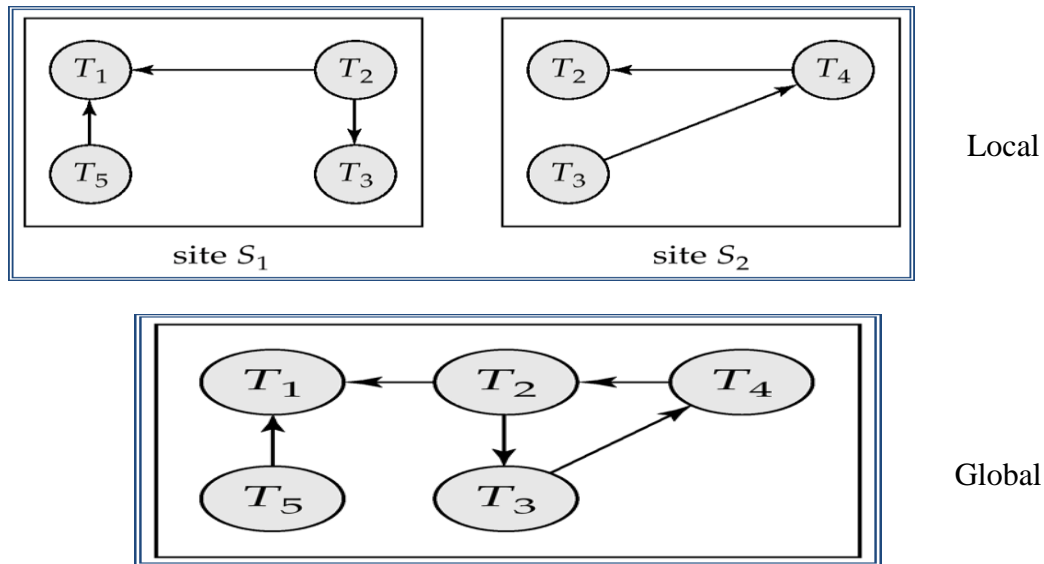
T_1 :	write (X) write (Y)	T_2 :	write (Y) write (X)
	X-lock on X write (X)		X-lock on Y write (Y) wait for X-lock on X

Result: **deadlock** which **cannot** be detected locally at either site

Centralized Approach

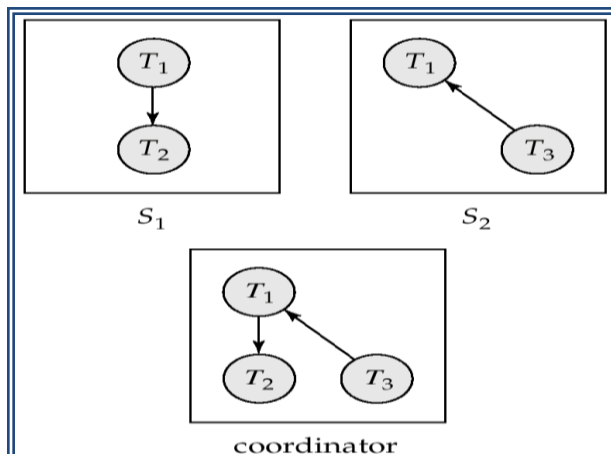
- A **global wait-for graph** is constructed and maintained in a *single site*; the deadlock-detection **coordinator**
 - *Real graph*: Real, but **unknown**, state of the system.
 - *Constructed graph*: **Approximation generated** by the controller during the execution of its algorithm .
- the **global wait-for graph** can be **constructed** when:
 - a **new edge** is **inserted** in or **removed** from one of the **local wait-for graphs**.
 - a number of **changes** have **occurred** in a **local wait-for graph**.
 - the **coordinator** needs to **invoke cycle-detection**.
- If the **coordinator** finds a cycle, it selects a **victim** and **notifies** all sites. The sites **roll back** the **victim** transaction.

Local and Global Wait-For Graphs



Example Wait-For Graph for False Cycles

Initial state:



- Suppose that starting from the state shown in figure,
 1. T_2 releases resources at S_1
 - ▶ resulting in a message **remove** $T_1 \rightarrow T_2$ message from the Transaction Manager at site S_1 to the coordinator)
 2. And then T_2 requests a resource held by T_3 at site S_2
 - ▶ resulting in a message **insert** $T_2 \rightarrow T_3$ from S_2 to the coordinator
- Suppose further that the **insert** message reaches before the **delete** message
 - this can **happen** due to **network delays**

- The **coordinator** would then **find** a **false cycle**

$$T_1 \rightarrow T_2 \rightarrow T_3 \rightarrow T_1$$
- The **false cycle** above **never existed** in reality.
- **False cycles cannot occur** if **two-phase locking** is **used**.

Unnecessary Rollbacks

- **Unnecessary rollbacks** may result when deadlock has indeed occurred and a **victim** has been **picked**, and **meanwhile another** of the **transactions** was **aborted** for reasons unrelated to the deadlock.
- Unnecessary rollbacks **can result from false cycles** in the global wait-for graph;
 - however, **likelihood** of **false cycles** is **low**.

Availability

- **High availability**: time for which system is not fully usable should be extremely low (e.g. **99.99% availability**)
- **Robustness**: ability of system to **function** in spite of **failures** of **components**
- **Failures** are **more** likely in **large** distributed systems
- To **be robust**, a distributed system **must**:
 - **Detect** failures
 - **Reconfigure** the system so computation may continue
 - **Recover/reintegrate** when a site or link is repaired
- **Failure detection**: distinguishing **link failure** from **site failure** is **hard**
 - (partial) **solution**: have **multiple links**, multiple link failure is likely a site failure

Reconfiguration

- **Reconfiguration**:
 - **Abort** all **transactions** that were active at a **failed site**
 - ▶ **Making** them **wait** could **interfere with other** transactions since they may **hold locks** on other sites
 - ▶ However, in case **only some replicas** of a data item failed, it may be **possible** to continue transactions that **had accessed** data at a failed site (more on this later)
 - If **replicated data** items were **at failed site**, update system catalog to **remove** them from the **list** of replicas.
 - ▶ This should be **reversed** when failed site **recovers**, but additional **care** needs to be taken to bring values **up to date**
 - If a **failed site** was a **central server** for some subsystem, an **election** must be **held** to determine the **new server**
 - ▶ E.g. **name server**, concurrency **coordinator**, global **deadlock detector**
- Since **network partition** may **not** be **distinguishable** from **site failure**, the following situations **must** be **avoided**:
 - **Two** or more **central servers** elected in distinct partitions
 - **More** than one **partition** updates a replicated data item
- **Updates** must be able to **continue** even if some **sites** are **down**
- **Solution: majority based approach**
 - Alternative of “read **one** write **all available**” is tantalizing but **causes problems**

Majority-Based Approach

- The **majority protocol** for distributed concurrency control can be modified to **work** even if some **sites** are **unavailable**
 - **Each replica** of each item has a **version number** which is **updated** when the replica is updated, as outlined below:
 - A **lock request** is sent to at least $\frac{1}{2}$ the **sites** at which item replicas are stored and **operation continues only** when a **lock** is **obtained** on a **majority** of the sites
 - **Read** operations **look** at **all replicas** locked, and **read** the value from the replica with **largest** version number
 - ▶ **May write** this **value** and version number **back** to replicas with lower version numbers (**no need** to obtain **locks** on **all replicas** for this task)
 - **Write** operations
 - ▶ **find** highest version number like reads, and set new version number to old highest version + 1
 - ▶ **Writes** are then **performed** on all locked replicas and version number on these replicas is set to new version number
 - **Failures** (network and site) cause **no problems** as long as:
 - ▶ Sites at commit contain a **majority** of replicas of any updated data items
 - ▶ During reads a **majority** of replicas are **available** to find version numbers
 - ▶ Subject to above, 2 phase commit can be used to update replicas
 - Note: **reads** are **guaranteed** to see latest version of data item
 - **Reintegration is trivial: nothing** needs to be done
- **Quorum consensus** algorithm can be similarly extended

Read One Write All (Available?)

- **Biased protocol** is a **special case** of **quorum consensus**
 - Allows **reads** to read **any one replica** but **updates require all replicas** to be **available** at commit time (called **read one write all**)
- Read **one write all available** (**ignoring failed sites**) is **attractive, but incorrect**:
 - If **failed link** may **come back up**, **without** a disconnected **site** ever **being aware** that it was disconnected
 - The **site** then **has old values**, and a **read** from that site **would return** an **incorrect value**
 - If **site** was **aware** of failure **reintegration could** have been **performed**, **but no** way to **guarantee** this
 - With **network partitioning**, sites in **each partition** may **update** same item **concurrently**
 - ▶ **believing sites in other partition** have **all failed**

Site Reintegration

- When **failed site recovers**, it **must catch up** with **all updates** that it **missed** while it was down
 - Problem: **updates** may be **happening** to items whose replica is stored at the site **while the site is recovering**
 - Solution 1: **halt** all **updates** on system while **reintegrating** a **site**

- ▶ **Unacceptable** disruption
- Solution 2: **lock** all **replicas** of all data items **at the site**, update to latest version, then release locks
 - ▶ Other **solutions** with **better concurrency** also available

Comparison with Remote Backup

- **Remote backup systems (hot spare)** are also designed to **provide high availability**
- **Remote backup systems** are **simpler** and have **lower overhead**
 - **All actions** performed at a **single site**, and only **log** records **shipped**
 - **No need** for distributed concurrency control, or 2 phase commit
- Using distributed databases with replicas of data items can provide higher availability by having **multiple (> 2) replicas** and using the majority protocol
 - Also **avoid** failure detection and **switchover** time associated with **remote backup** systems

Coordinator Selection

- Backup **coordinators**
 - **site** which **maintains** enough information locally to **assume** the **role** of coordinator **if the actual coordinator fails**
 - **executes** the same algorithms and **maintains** the same internal state information
 - **executes** state information **as the actual coordinator**, **if the actual coordinator fails**
 - **allows fast recovery** from coordinator failure **but involves overhead** during normal processing.
- **Election algorithms**
 - used to **elect** a **new coordinator** in case of failures
 - Example: **Bully Algorithm** - applicable to systems where **every site** can send a message to every other site.

Bully Algorithm

- If site S_i sends a **request** that is **not answered** by the coordinator within a time interval T , **assumes** that the **coordinator** has **failed**, S_i tries to **elect itself** as the new coordinator:
 - S_i **sends** an election message to **every site** with a higher identification number,
 - S_i then **waits** for any of these **processes** to **answer** within T ,
 - If **no response** within T , **assumes** that **all sites** with number **greater** than i have **failed**, S_i **elects itself** the new coordinator,
 - If **answer** is **received**, S_i **begins** time interval T' , **waiting** to **receive** a message that a **site** with a **higher** identification number has been **elected**.
 - If **no message** is **received** within T' , **assumes** the **site** with a **higher** number has **failed**; S_i **restarts** the algorithm.
- After a failed **site recovers**, it immediately **begins** execution of the same algorithm:
 - If there are **no active sites** with higher numbers, the **recovered site forces** all processes with lower numbers to **let it become** the **coordinator** site,
 - even if there is a **currently active coordinator** with a lower number.

Distributed Query Processing

- For **centralized systems**, the primary criterion for measuring the **cost** of a particular strategy is the **number of disk accesses**.
- In a **distributed system**, other issues must be taken into account:
 - The **cost** of a **data transmission** over the **network**.
 - The **potential gain** in performance from having **several sites** process **parts of the query in parallel**.

Query Transformation

- **Translating algebraic queries on fragments**.
 - It **must** be possible to **construct** relation ***r*** from its fragments
 - Replace relation ***r*** by the **expression to construct** relation ***r*** from its fragments
- Consider the **horizontal** fragmentation of the *account* relation into

$account_1 = \sigma_{branch_name = \text{"Hillside"}}(account)$

$account_2 = \sigma_{branch_name = \text{"Valleyview"}}(account)$

- The **query** $\sigma_{branch_name = \text{"Hillside"}}(account)$ becomes

$\sigma_{branch_name = \text{"Hillside"}}(account_1 \cup account_2)$

which is **optimized into**

$\sigma_{branch_name = \text{"Hillside"}}(account_1) \cup \sigma_{branch_name = \text{"Hillside"}}(account_2)$

- Since *account₁* has **only** tuples **pertaining** to the Hillside branch, we can eliminate the selection operation.
- Apply the definition of *account₂* to obtain
$$\sigma_{branch_name = \text{"Hillside"}}(\sigma_{branch_name = \text{"Valleyview"}}(account))$$
- This **expression** is the **empty** set regardless of the contents of the *account* relation.
- Final strategy is for the Hillside site to return *account₁* as the result of the query.

Simple Join Processing

- Consider the following relational algebra expression in which the **three relations** are **neither** replicated **nor** fragmented
$$account \bowtie depositor \bowtie branch$$
- *account* is stored at site *S₁*
- *depositor* at *S₂*
- *branch* at *S₃*
- For a **query** issued **at site *S₁***, the system needs to produce the **result at site *S₁***

Possible Query Processing Strategies

- **Ship** copies of **all three** relations to **site *S₁*** and choose a strategy for processing the entire locally at site *S₁*.
- **Ship** a copy of the *account* relation to **site *S₂*** and compute $temp_1 = account \bowtie depositor$ at *S₂*. **Ship** *temp₁* to *S₁*.
- Devise similar strategies, exchanging the roles *S₁*, *S₂*, *S₃*
- Must consider following **factors**:
 - **amount of data** being **shipped**
 - **cost of transmitting** a **data block** between sites
 - relative **processing speed** at **each site**

Heterogeneous Distributed Databases

- Many database applications **require data** from a variety of **preexisting databases** located in a heterogeneous collection of hardware and software platforms
- **Data models** may differ (**hierarchical**, relational , etc.)
- **Transaction** commit **protocols** may be **incompatible**
- **Concurrency control** may be based on **different** techniques (locking, **timestamping**, etc.)
- **System-level** details almost certainly are totally **incompatible**.
- A **multidatabase system** is a **software layer** on top of existing database systems, which is designed to manipulate information in heterogeneous databases
 - Creates an **illusion** of logical **database integration** without any physical database integration

Advantages

- **Preservation of investment** in existing
 - hardware
 - system **software**
 - **Applications**
- **Local autonomy** and administrative control
- **Allows** use of **special-purpose DBMSs**
- **Step** towards a **unified homogeneous DBMS**
 - **Full integration** into a homogeneous DBMS faces:
 - ▶ Technical difficulties and **cost of conversion**
 - ▶ Organizational/political difficulties
 - Organizations do not want to **give up control** on their data
 - Local databases wish to retain a **great** deal of **autonomy**

Unified View of Data

- **Agreement** on a common **data model**
 - Typically the **relational** model
- **Agreement** on a common **conceptual schema**
 - **Different names** for **same** relation/attribute
 - **Same** relation/attribute **name** means **different things**
- **Agreement** on a single **representation** of shared data
 - E.g. data types, precision,
 - Character sets
 - ▶ ASCII vs EBCDIC
 - ▶ Sort order variations
- **Agreement** on **units of measure**
- Variations in **names**
 - E.g. Köln vs Cologne, Mumbai vs Bombay

Query Processing

- Several issues in **query** processing in a **heterogeneous** database
- **Schema translation**
 - Write a **wrapper** for each **data source** to translate data to a global schema
 - **Wrappers** must also **translate updates** on global schema to updates on local schema
- **Limited query capabilities:**
 - Some data sources allow only **restricted** forms of **selections**
 - ▶ E.g. web forms, **flat file** data sources
 - **Queries** have to be **broken** up and processed partly at the source and partly at a different site
- Removal of **duplicate information** when sites have overlapping information
 - **Decide which** sites to execute query
- **Global query optimization**

Mediator Systems

- **Mediator** systems are systems that **integrate** multiple heterogeneous data sources by providing an **integrated global view**, and providing **query facilities** on global view
 - Unlike full fledged multidatabase systems, **mediators** generally **do not** bother about **transaction processing**
 - But the terms **mediator** and **multidatabase** are sometimes used interchangeably
 - The term **virtual database** is also used to refer to mediator/multidatabase systems

Directory Systems

- Typical kinds of **directory information**
 - **Employee** information such as **name, id, email, phone**, office addr, ..
 - Even **personal** information to be accessed from multiple places
 - ▶ e.g. Web browser bookmarks
- **White pages**
 - Entries organized by **name** or **identifier**
 - ▶ Meant for forward lookup to find more about an entry
- **Yellow pages**
 - Entries organized by **properties**
 - For reverse lookup to find entries matching specific requirements
- When **directories** are to be **accessed across** an organization
 - Alternative 1: Web interface. Not great for programs
 - **Alternative 2: Specialized directory access protocols**
 - ▶ Coupled with specialized **user interfaces**

Directory Access Protocols

- Most commonly used directory access protocol:
 - **LDAP (Lightweight Directory Access Protocol)**
 - Simplified from earlier **X.500 protocol**
- Question: **Why not** use database protocols like **ODBC/JDBC**?
- Answer:

- **Simplified protocols** for a **limited** type of **data access**, **evolved** parallel to ODBC/JDBC
- Provide a **nice hierarchical naming** mechanism similar to file system directories
 - **Data** can be **partitioned** amongst **multiple servers** for different parts of the hierarchy, yet give a single view to user
 - E.g. different servers for **Bell Labs** Murray Hill and **Bell Labs** Bangalore
- Directories **may use** databases as storage mechanism

LDAP: Lightweight Directory Access Protocol

- LDAP Data Model
- Data Manipulation
- Distributed Directory Trees

LDAP Data Model

- LDAP directories store **entries**
 - Entries are similar to objects
- Each entry must have unique **distinguished name (DN)**
- DN made up of a sequence of **relative distinguished names (RDNs)**
- E.g. of a DN
 - cn=Silberschatz, ou=Bell Labs, o=Lucent, c=USA
 - Standard RDNs (can be specified as part of schema)
 - cn: common name ou: organizational unit
 - o: organization c: country
 - Similar to paths in a file system but written in reverse direction
- Entries can have attributes
 - Attributes are multi-valued by default
 - LDAP has several built-in types
 - Binary, string, time types
 - Tel: telephone number PostalAddress: postal address
- LDAP allows definition of **object classes**
 - Object classes specify attribute names and types
 - Can use inheritance to define object classes
 - Entry can be specified to be of one or more object classes
 - No need to have single most-specific type
- Entries organized into a **directory information tree** according to their DN's
 - Leaf level usually represent specific objects
 - Internal node entries represent objects such as organizational units, organizations or countries
 - Children of a node inherit the DN of the parent, and add on RDNs
 - E.g. internal node with DN c = USA
 - Children nodes have DN starting with c=USA and further RDNs such as o or ou
 - DN of an entry can be generated by traversing path from root
 - Leaf level can be an alias pointing to another entry
 - Entries can thus have more than one DN

- E.g. person in more than one organizational unit

LDAP Data Manipulation

- Unlike SQL, LDAP does not define DDL or DML
- Instead, it defines a network protocol for DDL and DML
 - Users use an API or vendor specific front ends
 - LDAP also defines a file format
 - ▶ LDAP Data Interchange Format (LDIF)
- Querying mechanism is very simple: only selection & projection

LDAP Queries

- LDAP query must specify
 - Base: a node in the DIT from where search is to start
 - A search condition
 - ▶ Boolean combination of conditions on attributes of entries
 - Equality, wild-cards and approximate equality supported
 - A scope
 - ▶ Just the base, the base and its children, or the entire subtree from the base
 - Attributes to be returned
 - Limits on number of results and on resource consumption
 - May also specify whether to automatically dereference aliases
- LDAP URLs are one way of specifying query
- LDAP API is another alternative

LDAP URLs

- First part of URL specifies server and DN of base
 - ldap://aura.research.bell-labs.com/o=Lucent,c=USA
- Optional further parts separated by ? symbol
 - ldap://aura.research.bell-labs.com/o=Lucent,c=USA??sub?cn=Korth
 - Optional parts specify
 1. attributes to return (empty means all)
 2. Scope (sub indicates entire subtree)
 3. Search condition (cn=Korth)

Distributed Directory Trees

- Organizational information may be split into multiple directory information trees
 - Suffix of a DIT gives RDN to be tagged onto to all entries to get an overall DN
 - ▶ E.g. two DITs, one with suffix o = Lucent, c = USA and another with suffix o = Lucent, c = India
 - Organizations often split up DITs based on geographical location or by organizational structure
 - Many LDAP implementations support replication (master-slave or multi-master replication) of DITs (not part of LDAP 3 standard)
- A node in a DIT may be a **referral** to a node in another DIT

- E.g. Ou= Bell Labs may have a separate DIT, and DIT for o=Lucent may have a leaf with ou=Bell Labs containing a referral to the Bell Labs DIT
- Referrals are the key to integrating a distributed collection of directories
- When a server gets a query reaching a referral node, it may either
 - ▶ Forward query to referred DIT and return answer to client, or
 - ▶ Give referral back to client, which transparently sends query to referred DIT (without user intervention)