

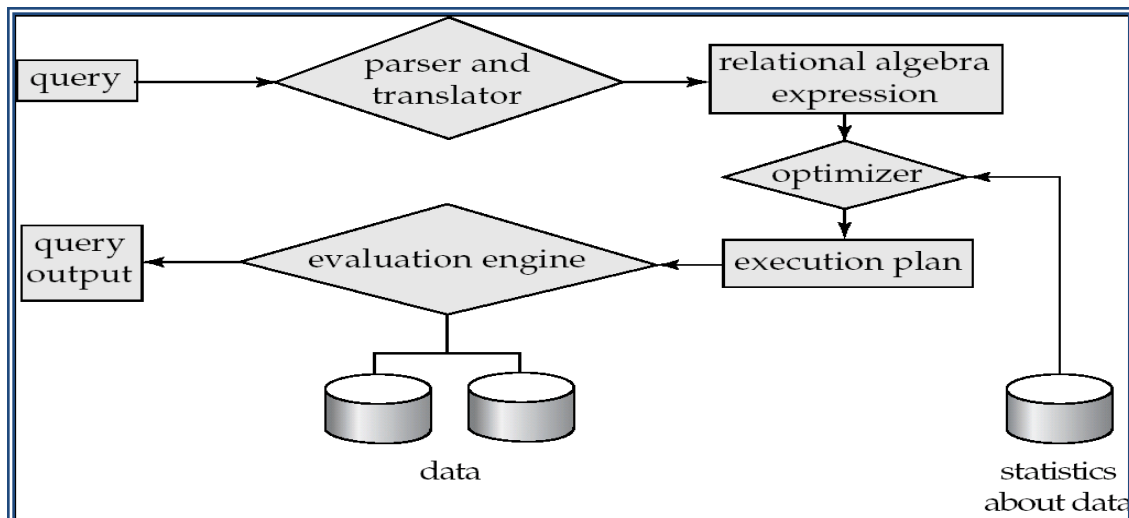
BACHELOR OF SCIENCE IN COMPUTER SCIENCE
YEAR 4 SEMESTER 2
CCS 418: ADVANCED DATABASE SYSTEMS
NOTES 2

QUERY PROCESSING

- Overview
- Measures of *Query Cost*
- *Selection* Operation
- *Sorting*
- *Join* Operation
- *Other* Operations
- Evaluation of *Expressions*

Basic Steps in Query Processing

1. *Parsing* and translation
2. Optimization
3. Evaluation



- **Parsing and translation**
 - *translate* the *query* into its *internal form*.
 - This is then *translated* into *relational algebra*.
 - *Parser* checks *syntax*,
 - verifies relations
- **Evaluation**
 - The query-execution engine takes a *query-evaluation plan*,
 - *executes* that *plan*, and
 - *returns* the *answers* to the query.
- **Query Optimization:** Amongst *all equivalent evaluation plans*:

- **choose** the **one** with **lowest cost**.
- **Cost** is **estimated** using **statistical information** from:
 - ▶ the *database catalog*
 - ▶ e.g. *number of tuples* in each relation, *size of tuples*, etc.
- In *this chapter* we study:
 - **How** to **measure query costs**
 - **Algorithms** for **evaluating relational algebra operations**
 - **How** to **combine algorithms** for individual operations in order to:
 - ▶ *evaluate a complete expression*

Measures of Query Cost

- **Cost** is generally **measured** as:
 - ▶ **total elapsed time** for *answering query*
 - Many **factors** contribute to **time cost**:
 - ▶ *disk accesses*, *CPU*, or even *network communication*
- Typically **disk access** is the *predominant cost*, and
 - ▶ is also *relatively easy to estimate*.
 - **Measured** by taking into account:
 - ▶ *Number of seeks* * *average-seek-cost*
 - ▶ *Number of blocks read* * *average-block-read-cost*
 - ▶ *Number of blocks written* * *average-block-write-cost*
 - ▶ **Cost** to **write** a block is **greater** than cost to **read** a block
 - data is *read back* after being written:
 - » to ensure that the write was *successful*
- As the **cost measures**:
 - for *simplicity* we just use:
 - ▶ the *number of block transfers* from disk and
 - ▶ the *number of seeks*
 - t_R – *time to transfer one block*
 - t_S – *time for one seek*
 - **Cost** for **b** block transfers plus **S** seeks:

$$\text{Cost} = b * t_R + S * t_S$$
- We **ignore CPU costs** for *simplicity*
 - *Real systems* do **take CPU cost** into account
- We do **not include** cost to **writing output** to disk in our cost formulae
- Several **algorithms** can **reduce disk I/O** by:
 - ▶ using *extra buffer space*
 - Amount of **real memory** available to **buffer** depends on:
 - ▶ other concurrent **queries** and **OS processes**,
 - known **only** during *execution*
 - ▶ We often use **worst case estimates**, assuming:
 - only the **minimum** amount of **memory**
 - needed for the operation is **available**
- **Required data** may be **buffer resident** already,
 - ▶ avoiding **disk I/O**
 - But **hard** to take into account for **cost estimation**

Selection Operation

- **File scan algorithms:**
 - ▶ search algorithms that **locate** and **retrieve** records
 - that **fulfill** a selection condition. (No Index use!)
- Algorithm **A1** (**linear search**): **Scan** each **file block** and
 - ▶ test all records to see whether they **satisfy** the selection condition.
 - **Cost** estimate = b_r block transfers + 1 seek
 - ▶ b_r denotes number of **blocks** containing records from **relation r**
 - If selection is on a **key attribute**:
 - ▶ can **stop** on finding record
 - ▶ cost = $(b_r/2)$ block transfers + 1 seek
 - **Linear search** can be applied **regardless of**:
 - ▶ selection **condition** or
 - ▶ **ordering** of records in the file, or
 - ▶ availability of **indices**
- **A2** (**binary search**). Applicable if selection is:
 - ▶ an **equality** comparison :
 - ▶ on the **attribute** on which **file** is **ordered**.
 - Assume that the **blocks** of a relation are **stored contiguously**
 - **Cost** estimate (number of **disk blocks** to be scanned):
 - ▶ cost of **locating** the **first** tuple by a **binary search** on the blocks
 - $\lceil \log_2(b_r) \rceil * (t_T + t_S)$
 - ▶ If there are **multiple records** satisfying selection
 - Add transfer cost of the number of **blocks** containing records that **satisfy** selection condition

Selections Using Indices

- **Index scan algorithms:** search algorithms that **use** an index
 - ▶ So, selection condition must be on search-key of an index.
- **A3** (primary index on **candidate key**, equality).
 - Retrieve a **single** record that satisfies the corresponding equality condition
 - ▶ Cost = $(h_i + 1) * (t_T + t_S)$
- **A4** (primary index on **nonkey**, equality) :
 - ▶ Retrieve **multiple** records.
 - Records will be on consecutive blocks
 - ▶ Let **b** = number of **blocks** containing matching records
 - ▶ Cost = $h_i * (t_T + t_S) + t_S + t_T * b$
- **A5** (equality on search-key of secondary index).
 - Retrieve a **single** record **if** the search-key is a candidate key
 - ▶ Cost = $(h_i + 1) * (t_T + t_S)$
 - Retrieve **multiple** records **if** search-key is **not** a candidate key
 - ▶ each of **n** matching records may be on a different block
 - ▶ Cost = $(h_i + n) * (t_T + t_S)$
 - be very expensive!

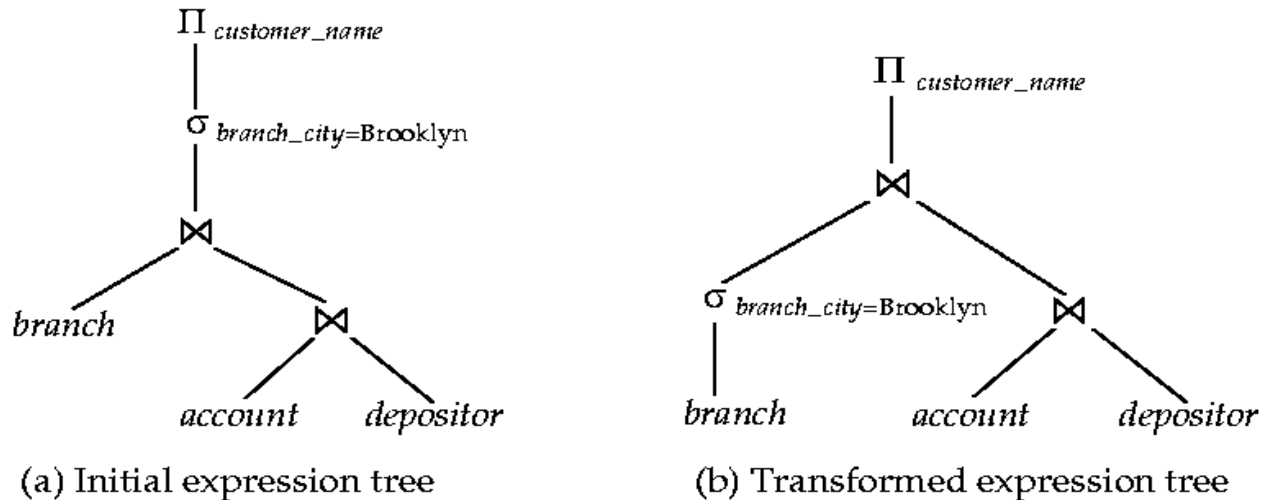
Selections Involving Comparisons

- Can **implement** selections of the form $\sigma_{A \leq v}(r)$ or $\sigma_{A \geq v}(r)$ by using:
 - a **linear file scan** or **binary search**,
 - **or** by using **indices** in the following ways:
- **A6** (primary index, comparison). (Relation is **sorted** on **A**)
 - ▶ For $\sigma_{A \geq v}(r)$ use **index** to find first tuple $\geq v$ and
 - then scan relation sequentially from there
 - ▶ For $\sigma_{A \leq v}(r)$ just scan relation sequentially till first tuple $> v$;
 - do **not** use index
- **A7** (secondary index, comparison).
 - ▶ For $\sigma_{A \geq v}(r)$ use index to find first index entry $\geq v$ and
 - then scan index sequentially from there,
n to find pointers to records.
 - ▶ For $\sigma_{A \leq v}(r)$ just scan **leaf pages** of index finding pointers to records,
 - till first entry $> v$
 - ▶ In either case, retrieve records that are pointed to,
 - requires **an I/O** for each record !
 - Linear file scan may be cheaper !

QUERY OPTIMIZATION

Introduction

- **Alternative ways** of evaluating a given query
 - **Equivalent expressions**
 - **Different algorithms** for each operation
- **Cost difference** between a **good** and a **bad way** of evaluating a query
 - ▶ can be **enormous!**
- **Need to estimate** the **cost** of operations
 - **Statistical** information about relations.
E.g.:
 - ▶ number of tuples,
 - ▶ number of distinct values for an attributes,
 - Statistics **estimation** for **intermediate** results
 - ▶ to **compute cost** of **complex** expressions
- ▶ Relations generated by two **equivalent expressions** :
 - ▶ have the **same** set of attributes and
 - ▶ contain the **same** set of tuples
 - although their tuples/attributes may be **ordered differently**.



- **Generation** of *query-evaluation plans* for an expression
 - ▶ involves *several steps*:
 1. Generating logically **equivalent** expressions using *equivalence rules*.
 2. Annotating resultant **expressions** to get *alternative query plans*
 3. Choosing the **cheapest plan** based on *estimated cost*
- The overall process is called **cost based optimization**.

Transformation of Relational Expressions

- Two relational algebra **expressions** are said to be **equivalent**:
 - if on every legal *database instance* the **two expressions** generate
 - ▶ the **same set of tuples**
 - Note: **order** of tuples *is irrelevant!*
- In SQL, inputs and outputs **are multisets of tuples**:
 - Two **expressions** in the multiset version of the relational algebra are said to be equivalent:
 - ▶ if on every legal *database instance* the **two expressions** generate
 - the **same multiset of tuples**
- An **equivalence rule** says that:
 - If expressions of **two forms** are equivalent:
 - ▶ Can **replace** expression of **first form** by **second**,
 - ▶ or *vice versa*

Enumeration of Equivalent Expressions

- *Query optimizers* use *equivalence rules* to systematically
 - ▶ generate expressions equivalent to the given expression
- Conceptually, generate *all equivalent expressions* by repeatedly executing the following step until no more expressions can be found:
 - for each *expression* found so far,
 - ▶ use all applicable *equivalence rules*
 - ▶ add *newly generated expressions* to the *set of expressions* found so far

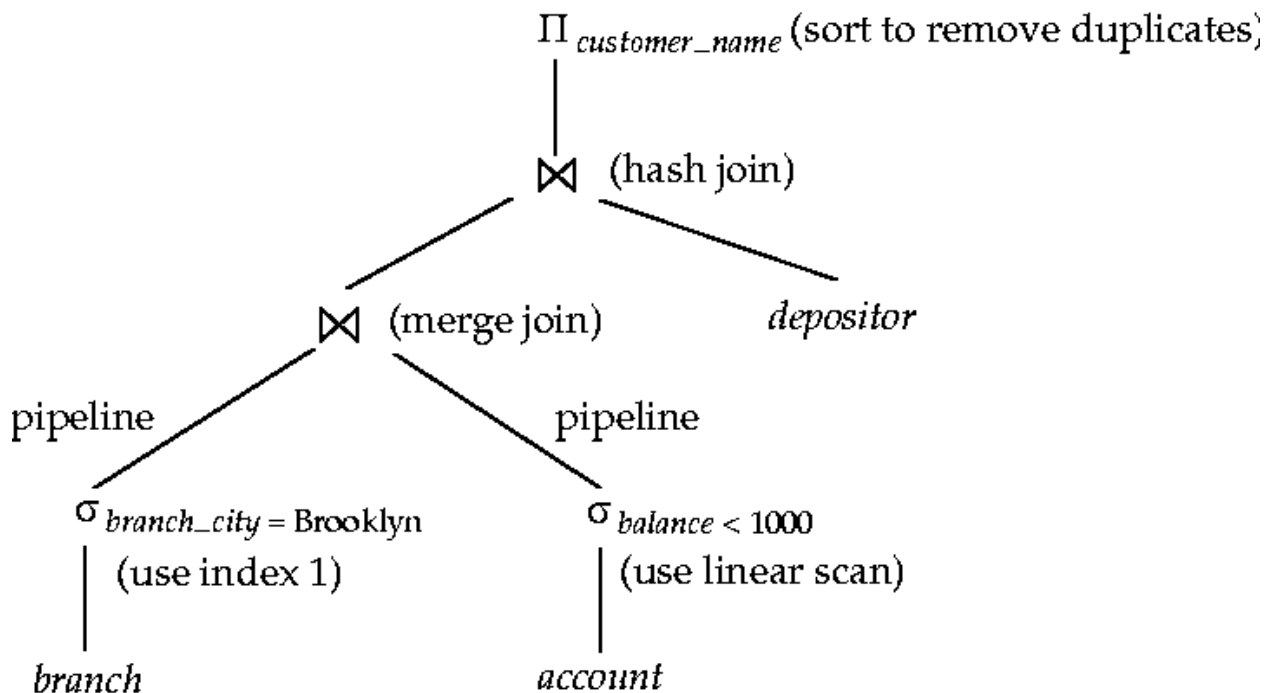
- The above approach is *very expensive* in *space* and *time*
- *Space requirements* reduced by *sharing common subexpressions*:
 - when E1 is *generated* from E2 by an *equivalence rule*:
 - usually *only* the *top level* of the two are *different*,
 - *subtrees* below are the *same* and *can be shared*
 - E.g. when applying *join associativity*
- *Time requirements* are reduced by *not generating all expressions*

Cost Estimation

- Cost of *each operator computes*
 - Need *statistics* of *input relations*
 - E.g. *number of tuples*, *sizes of tuples*
- *Inputs can be results of sub-expressions*
 - Need to *estimate statistics* of *expression results*
 - To do so, we *require additional statistics*
 - E.g. *number of distinct values* for an *attribute*

Evaluation Plan

- An *evaluation plan* defines exactly
 - *what algorithm* is used for *each operation*, and
 - *how the execution* of the operations is *coordinated*.



Choice of Evaluation Plans

- Must *consider* the interaction of *evaluation techniques*

- when choosing evaluation plans:
 - ▶ choosing the cheapest algorithm for *each* operation independently
 - may not *yield best overall algorithm*.
 - E.g., merge-join may be costlier than hash-join,
 - ▶ but may provide a *sorted output*
 - which *reduces* the cost for an *outer level aggregation*.
 - nested-loop join may *provide opportunity* for pipelining
 - Practical query optimizers incorporate elements of:
 - ▶ the following two broad approaches:
1. Search *all the plans* and choose the *best* plan in a *cost-based* fashion.
 2. Use *heuristics* to choose a plan.

Cost-Based Optimization

- Consider finding the best *join-order* for $r_1 \bowtie r_2 \dots r_n$.
- There are $(2(n-1))! / (n-1)!$ different *join orders* for above expression.
 - ▶ With $n = 7$, the number is 665280,
 - ▶ with $n = 10$, the number is greater than 17.6 billion!

BUT:

- No need to generate all the *join orders*.
- Using *dynamic programming*:
 - ▶ the least-cost *join order* for any subset of $\{r_1, r_2, \dots, r_n\}$ is:
 - *computed* only once and
 - *stored* for *future use*.

Dynamic Programming in Optimization

- To find best join tree for a set of n relations:
 - To find best plan for a set S of n relations,
 - ▶ consider all possible plans of the form: $S_1 \bowtie (S - S_1)$
 - where S_1 is any *non-empty* subset of S .
 - » i.e., $2^n - 1$ *alternatives*!
 - » e.g., 1023 for $\{r_1, r_2, \dots, r_{10}\}$
 - Recursively,
 - ▶ *Compute* costs for joining subsets of S to find the cost of each plan.
 - ▶ *Choose* the cheapest of the $2^n - 1$ *alternatives*.
 - When a plan for any subset is *computed*,
 - ▶ *store* it and *reuse* it when it is required again,
 - instead of recomputing it.

Heuristic Optimization

- Cost-based optimization is expensive,
 - ▶ **even with** dynamic programming.
- **Systems** may **use *heuristics*** to **reduce**
 - ▶ the number of choices that must be made in a cost-based fashion.
- **Heuristic** optimization **transforms the query-tree** by

- ▶ using a set of rules that typically (but *not in all cases*)
 - ▶ **improve** execution performance:
- Perform *selection early* (reduces the number of tuples)
- Perform *projection early* (reduces the number of attributes)
- Perform most restrictive selection and join operations
 - ▶ **before** other similar operations.
- **Some** systems use only heuristics,
 - ▶ **others** combine heuristics with *partial cost-based optimization*.

Structure of Query Optimizers

- The **System R/Starburst optimizer** considers **only** left-deep join orders.
 - This reduces optimization complexity and
 - generates **plans** amenable to pipelined evaluation.
 - **System R/Starburst** also
 - ▶ uses **heuristics** to **push** selections and projections **down** the query tree.
- Heuristic optimization used in some versions of **Oracle**:
 - Repeatedly pick “best” relation to join next
 - ▶ Starting from each of **n starting points**.
 - ▶ Pick best among these.
- For scans using secondary indices,
 - **some optimizers** take into account the **probability** that
 - ▶ the page containing the tuple is in the buffer.
- Intricacies of SQL complicate query optimization
 - E.g. **nested** subqueries
- Some query optimizers integrate
 - ▶ heuristic selection and
 - ▶ the **generation** of alternative *access plans*.
 - **System R** and **Starburst** use a *hierarchical procedure* based on
 - ▶ the *nested-block* concept of SQL:
 - heuristic rewriting followed by cost-based *join-order optimization*.
- **Even with** the use of *heuristics*,
 - *cost-based query optimization imposes a substantial overhead*.
- This **expense** is usually *more than offset* by:
 - **savings** at query-execution time,
 - ▶ *particularly* by *reducing* the number of *slow disk* accesses.

Optimizing Nested Subqueries

- SQL conceptually **treats** nested subqueries in the *where* clause
 - as functions that take *parameters* and
 - ▶ **return** a single value or set of values
 - *Parameters* are *variables from outer level* query that are
 - ▶ used in the nested subquery;
 - ▶ such variables are *called correlation variables*

- E.g. **select** *customer_name*
from *borrower*
where exists (**select** *
from *depositor*
where *depositor.customer_name* =
borrower.customer_name)
- Conceptually, nested subquery is **executed once for each tuple**
 - ▶ in the *cross-product* generated by the *outer level from* clause
 - Such evaluation is called **correlated evaluation**
 - Note: other conditions in **where** clause may be
 - ▶ **used** to compute a **join** (instead of a *cross-product*)
 - ▶ **before** executing the *nested subquery*

- Correlated evaluation **may be quite inefficient** since:
 - a *large number of calls* may be made to the *nested subquery*
 - there may be unnecessary **random I/O** as a result
- **SQL optimizers** attempt to
 - **transform** nested subqueries to **joins** where *possible*,
 - **enabling** use of *efficient join techniques*
 - ▶ E.g.: earlier nested query can be rewritten as:

```
select customer_name
from borrower, depositor
where depositor.customer_name = borrower.customer_name
      ▶ Note: above query doesn't correctly deal with duplicates,
      can be modified to do so as we will see
```

- **In general**, it is **not possible / straightforward**
 - to **move** the **entire** nested subquery into the *outer level query*
 - A temporary relation is **created** instead, and
 - ▶ used in **body** of *outer level query*

In general, SQL queries of the form below can be rewritten as shown

- Rewrite: **select** ...
from L_1
where P_1 and exists (**select** *
from L_2
where P_2)
- To: **create table** t_1 as
select distinct V
from L_2
where P_2^1
select ...
from L_1, t_1
where P_1 and P_2^2
 - P_2^1 contains *predicates* in P_2 that do *not involve* any *correlation variables*
 - P_2^2 reintroduces *predicates involving correlation variables*,
with relations renamed appropriately
 - V contains **all attributes** used in *predicates with correlation variables*

- In our example, the original nested query would be **transformed** to:

```
create table t1 as
  select distinct customer_name
    from depositor ;
  select customer_name
    from borrower, t1
  where t1.customer_name = borrower.customer_name
```

- The process of replacing a nested query **by** a *query with a join*
 - ▶ (possibly with a *temporary relation*)
 - is called **decorrelation**.
- **Decorrelation** is more **complicated** when
 - the nested subquery *uses* aggregation, or
 - when *the result* of the nested subquery is *used* to test for equality, or
 - when the condition *linking* the nested subquery *to the other query*
 - ▶ is **not exists**,
 - and *so on*.

Materialized Views

- A **materialized view** is
 - a **view** whose contents are computed and stored.
- Consider the view:

```
create view branch_total_loan(branch_name, total_loan) as
  select branch_name, sum(amount)
  from loan
 groupby branch_name
```

- Materializing the above **view** would be **very useful**
- if the **total loan amount** is **required** frequently
- **Saves the effort** of
 - ▶ *finding multiple tuples* and
 - ▶ *adding up* their amounts

Materialized View Maintenance

- The task of **keeping** a *materialized view* **up-to-date** with the *underlying data*
 - is known as *materialized view maintenance*
- Materialized views **can** be maintained by **recomputation** on *every update*
- A **better option** is:
 - ▶ to use **incremental view maintenance**
 - **Changes** to *database relations* are **used**
 - ▶ to **compute** *changes to materialized view*,
 - ▶ which is then **updated**
- View maintenance **can** be done by:
 - Manually **defining triggers** on insert, delete, and update of each relation in the view definition

- Manually **written code** to update the view whenever database relations are updated
- **OR:** Supported **directly by the database**

Incremental View Maintenance

- The **changes** (inserts and deletes) to a relation or expressions are
 - *referred to* as its **differential**
 - **Set of tuples** inserted to and deleted from **r** are
 - denoted **i_r** and **d_r**
- To *simplify* our description, we *only consider* inserts and deletes
 - We replace **updates** to a tuple by
 - *deletion* of the tuple followed by
 - *insertion* of the update tuple
- We describe how to **compute** the change:
 - to the **result of each relational operation**,
 - **given** changes to its **inputs**
- We then **outline** *how to handle* relational algebra *expressions*

TRANSACTIONS

- Transaction Concept
- Transaction State
- Concurrent Executions
- Serializability
- Recoverability
- Implementation of Isolation
- Transaction Definition in SQL
- Testing for Serializability

Transaction Concept

- A **transaction** is a *unit* of program execution that
 - **accesses** and possibly **updates** various data items.
- A transaction **must** see a **consistent database**.
 - During transaction **execution**:
 - the database may be temporarily *inconsistent*.
 - When the transaction **completes** successfully (is committed),
 - the database **must** be consistent.
 - After a transaction commits,
 - the **changes** it has made to the database **persist**,
 - even if there are *system failures*.
- **Multiple** transactions **can** execute in parallel.
 - **Two main issues** to deal with:
 - Failures of various kinds, such as :
 - hardware *failures* and system *crashes*

- ▶ Concurrent execution of *multiple* transactions

ACID Properties

- A **transaction** is a *unit* of program execution that:
 - ▶ *accesses* and possibly *updates* various data items.
- To **preserve** the **integrity** of data, the database system must **ensure**:
 - **Atomicity.** Either all operations of the transaction are:
 - ▶ properly reflected in the database or none are.
 - **Consistency.** Execution of a transaction in isolation:
 - ▶ **preserves** the **consistency** of the database.
 - **Isolation.** Although multiple transactions may execute concurrently,
 - ▶ **each** *transaction* must be **unaware** of other concurrent *transactions*.
 - ▶ Intermediate transaction results **must be**:
 - **hidden** from other concurrently executed transactions.
 - ▶ That is, for every pair of transactions T_i and T_j , it appears to T_i that **either** T_j , finished execution before T_i started, **or** T_j started execution after T_i finished.
 - **Durability.** After a transaction **completes** successfully,
 - ▶ the **changes** it has made to the database **persist**,
 - even if there are system *failures*.


Example of Fund Transfer

- Transaction to transfer \$50 from account **A** to account **B**:
 1. **read(A)**
 2. $A := A - 50$
 3. **write(A)**
 4. **read(B)**
 5. $B := B + 50$
 6. **write(B)**
- **Atomicity requirement** :
 - if the transaction **fails** after step 3 and before step 6,
 - ▶ the **system** should **ensure** that :
 - its **updates** are *not reflected* in the database,
 - else an *inconsistency* will result.
- **Consistency requirement** :
 - the **sum** of **A** and **B** is:
 - ▶ **unchanged** by the execution of the transaction.
- **Isolation requirement** —
 - if between steps 3 and 6,
 - ▶ another transaction is allowed to access the partially updated database,
 - it will see an inconsistent database
 - (the sum $A + B$ will be less than it should be).
 - Isolation can be **ensured** trivially by:
 - ▶ running transactions **serially**,
 - that is **one** after the **other**.

- However, executing multiple *transactions concurrently*
 - has significant **benefits**, as we will see later.
- **Durability requirement** :
 - once the user has been notified that the transaction has **completed** :
 - (i.e., the transfer of the \$50 has taken place),
 - the **updates** to the database by the transaction **must persist**
 - despite *failures*.

Transaction State

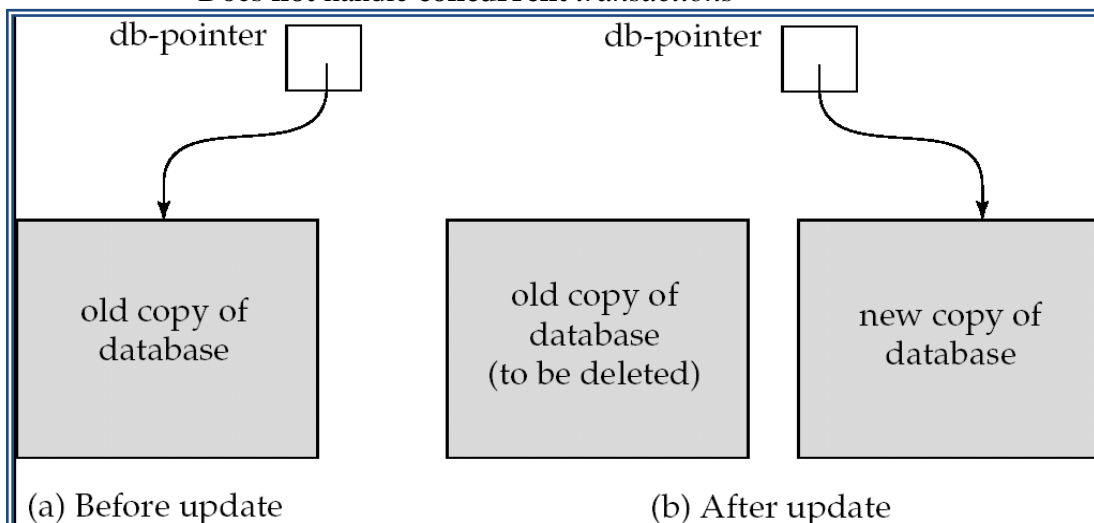
- **Active** – the initial state;
 - the transaction stays in this state while it is executing
- **Partially committed** – after the final statement has been executed.
- **Failed** -- after the discovery that *normal execution* can no longer proceed.
- **Aborted** – after the transaction has been rolled back and
 - the database restored to its state prior to the start of the *transaction*.
 - **Two options** after it has been aborted:
 - restart the transaction; can be done
 - only **if** no internal logical error
 - **kill** the transaction
- **Committed** – after **successful** completion.



Implementation of Atomicity and Durability

- The recovery-management component of a database system
 - implements the support for atomicity and durability.
- The *shadow-database* scheme:

- assume that only one *transaction* is active at a time.
 - a pointer called `db_pointer` always points to the current consistent copy of the database.
 - all updates are made on a *shadow copy* of the database, and `db_pointer` is made to point to the updated shadow copy only after the transaction reaches partial commit and all updated pages have been flushed to disk.
 - in case transaction fails, old consistent copy pointed to by `db_pointer` can be used, and the shadow copy can be deleted.
- Assumes **disks** do not fail
 - Useful for **text editors**, but
 - extremely **inefficient** for *large databases* (why?)
 - Does **not** handle **concurrent transactions**



- Assumes **disks** do not fail
- Useful for **text editors**, but
 - extremely **inefficient** for *large databases* (why?)
 - Does **not** handle **concurrent transactions**

Concurrent Executions

- Multiple transactions are allowed to run concurrently in the system:
 - **increased processor and disk utilization**,
 - ▶ leading to **better transaction throughput**:
 - **one transaction** can be using the CPU while
 - **another** is reading from or writing to the disk
 - **reduced average response time** for transactions:
 - ▶ **short transactions** need not wait behind **long ones**.
- **Concurrency control schemes** :
 - mechanisms to achieve isolation; that is,
 - ▶ to **control the interaction** among the concurrent *transactions*
 - in order to **prevent** them from destroying the **consistency** of the database
 - ▶ Will study in Chapter 16, after studying

- notion of correctness of *concurrent executions*.

Schedules

- **Schedule** – a sequences of instructions that **specify** the chronological order
 - in which *instructions* of concurrent transactions are *executed*
- a schedule for a *set of transactions* must consist of
 - *all instructions* of those *transactions*
- must **preserve** the **order** in which
 - the instructions **appear** in each individual transaction.
- A *transaction* that successfully **completes** its execution
 - will have a **commit** *instructions* as the last statement
 - (will be omitted if it is obvious)
- A *transaction* that **fails** to successfully complete its execution
 - will have an **abort** *instructions* as the last statement
 - (will be omitted if it is obvious)

Schedule 1

- Let :
 - **T₁** transfer **\$50** from A to B, and
 - **T₂** transfer **10%** of the balance from A to B.
- A **serial** schedule in which **T₁** is followed by **T₂**:

<i>T₁</i>	<i>T₂</i>
read(<i>A</i>) $A := A - 50$ write (<i>A</i>) read(<i>B</i>) $B := B + 50$ write(<i>B</i>)	read(<i>A</i>) $temp := A * 0.1$ $A := A - temp$ write(<i>A</i>) read(<i>B</i>) $B := B + temp$ write(<i>B</i>)

Schedule 2

T_1	T_2
$\text{read}(A)$ $A := A - 50$ $\text{write}(A)$ $\text{read}(B)$ $B := B + 50$ $\text{write}(B)$	$\text{read}(A)$ $\text{temp} := A * 0.1$ $A := A - \text{temp}$ $\text{write}(A)$ $\text{read}(B)$ $B := B + \text{temp}$ $\text{write}(B)$

Schedule 3

- Let T_1 and T_2 be the transactions defined previously:
 - The following schedule is **not** a **serial** schedule,
 - but it is *equivalent* to Schedule 1.

T_1	T_2
$\text{read}(A)$ $A := A - 50$ $\text{write}(A)$ $\text{read}(B)$ $B := B + 50$ $\text{write}(B)$	$\text{read}(A)$ $\text{temp} := A * 0.1$ $A := A - \text{temp}$ $\text{write}(A)$ $\text{read}(B)$ $B := B + \text{temp}$ $\text{write}(B)$

- In Schedules 1, 2 and 3:
 - the **sum**($A + B$) is *preserved*.

Schedule 4

T_1	T_2
<code>read(A)</code> $A := A - 50$ <code>write(A)</code> <code>read(B)</code> $B := B + 50$ <code>write(B)</code>	<code>read(A)</code> $temp := A * 0.1$ $A := A - temp$ <code>write(A)</code> <code>read(B)</code> $B := B + temp$ <code>write(B)</code>

Serializability

- **Basic Assumption:**
 - Each **transaction** preserves database **consistency**.
- Thus:
 - **Serial** execution of a **set of transactions** preserves database **consistency**.
- A (possibly concurrent) **schedule** is **serializable** :
 - if it is **equivalent** to a **serial schedule**.
- Different forms of **schedule equivalence** give rise to the notions of:
 1. **conflict** serializability
 2. **view** serializability
- We ignore operations other than **read** and **write** instructions, and we assume that:
 - transactions may perform:
 - ▶ arbitrary computations on data in local buffers,
 - ▶ in between **reads** and **writes**.
 - Our *simplified schedules* consist of :
 - ▶ only **read** and **write** instructions.

Conflicting Instructions

- Instructions l_i and l_j of transactions T_i and T_j respectively,
 - **conflict** if and only if there exists
 - ▶ some item Q accessed by **both** l_i and l_j , and
 - ▶ at least one of these instructions **wrote** Q .
 - 1. $l_i = \text{read}(Q)$, $l_j = \text{read}(Q)$. l_i and l_j don't conflict.
 - 2. $l_i = \text{read}(Q)$, $l_j = \text{write}(Q)$. They conflict.
 - 3. $l_i = \text{write}(Q)$, $l_j = \text{read}(Q)$. They conflict
 - 4. $l_i = \text{write}(Q)$, $l_j = \text{write}(Q)$. They conflict
- Intuitively, a **conflict** between l_i and l_j **forces**
 - a (logical) **temporal order** between them.

- If l_i and l_j are consecutive in a schedule and they do not conflict,
 - their **results** would remain the **same**
 - ▶ even if they had been interchanged in the *schedule*.

Conflict Serializability

- If a *schedule* S can be **transformed** into a *schedule* S' by
 - a *series of swaps* of **non-conflicting** instructions,
 - ▶ we say that S and S' are **conflict equivalent**.
- We say that a *schedule* S is **conflict serializable**
 - if it is *conflict equivalent* to a **serial** schedule
- Schedule 3 can be transformed into Schedule 6,
 - a *serial schedule* where T_2 follows T_1 ,
 - ▶ by *series of swaps* of **non-conflicting** instructions.
 - Therefore:
 - ▶ Schedule 3 is **conflict serializable**.

T_1	T_2
read(A) write(A)	read(A) write(A)
read(B) write(B)	
	read(B) write(B)

Schedule 3

T_1	T_2
read(A) write(A) read(B) write(B)	read(A) write(A)
	read(B) write(B)

Schedule 6

- Example of a **schedule** that is **not conflict serializable**:

T_3	T_4
read(Q)	write(Q)
write(Q)	

- We are **unable** to **swap instructions** in the above schedule to obtain:
 - either the serial schedule $\langle T_3, T_4 \rangle$,
 - or the serial schedule $\langle T_4, T_3 \rangle$.

View Serializability

- Let S and S' be two schedules with the **same** set of transactions.
 - S and S' are **view equivalent**, if the following **three** conditions are met:
 - For each data item Q , if transaction T_i **reads** the initial value of Q in schedule S ,
 - then transaction T_i **must**, in schedule S' , also **read** the initial value of Q .
 - For each data item Q if transaction T_i executes **read(Q)** in schedule S , and that **value** was produced by transaction T_j (if any),
 - then transaction T_i must in schedule S' also **read** the value of Q that was produced by transaction T_j .
 - For each data item Q , the transaction (if any) that **performs** the final **write(Q)** operation in schedule S
 - must perform** the final **write(Q)** operation in schedule S' .

As can be seen, **view** equivalence is also based purely on **reads** and **writes** alone.

- A schedule S is **view serializable**,
 - if it is view equivalent to a **serial schedule**.
- Every** conflict serializable schedule is **also** view serializable.
- Below is a schedule which is view-serializable but **not** conflict serializable.

T_3	T_4	T_6
read(Q)	write(Q)	
write(Q)		
		write(Q)

- What **serial schedule** is above equivalent to?
- Every view serializable schedule that is **not** conflict serializable:
 - has **blind writes**.

Other Notions of Serializability

- The schedule below produces **same** outcome as the **serial** schedule $\langle T_1, T_5 \rangle$,
 - yet!, is **not** conflict equivalent or view equivalent to it!


T_1	T_5
read(A) $A := A - 50$ write(A)	read(B) $B := B - 10$ write(B)
read(B) $B := B + 50$ write(B)	
	read(A) $A := A + 10$ write(A)

- Determining** such equivalence requires:

- *analysis of operations* other than **read** and **write**.

Testing for Serializability


- Consider some schedule of a set of transactions T_1, T_2, \dots, T_n
- **Precedence graph** :
 - ▶ a **direct graph** where the vertices **are** the *transactions* (names).
- We draw an arc from T_i to T_j
 - if the two transaction conflict, and
 - T_i accessed the data item on which the conflict arose **earlier**.
- We may label the arc by the **item** that was **accessed**.
- **Ex. 1:**



Example Schedule (Schedule A) + *Precedence Graph*

T_1	T_2	T_3	T_4	T_5
read(Y) read(Z)	read(X)			read(V) read(W) read(W)
	read(Y) write(Y)	write(Z)		
read(U)			read(Y) write(Y) read(Z) write(Z)	
read(U) write(U)				

Test for Conflict Serializability



- A schedule is *conflict serializable*
 - ▶ if and only if its *precedence graph* is **acyclic**.
- **Cycle-detection** algorithms exist which:
 - take **order n^2** time,
 - ▶ where n is the *number of vertices* in the graph.
 - Better algorithms take **order $n + e$**
 - ▶ where e is the *number of edges*.
- If precedence graph is **acyclic**, the *serializability order* can be obtained by a **topological sorting** of the graph.
 - This is a linear order *consistent* with:
 - ▶ the partial order of the graph.
 - For example, a *serializability order* for Schedule A would be:

$$T_5 \rightarrow T_1 \rightarrow T_3 \rightarrow T_2 \rightarrow T_4$$
 - ▶ Are there others?

Test for View Serializability

- The precedence graph test for conflict serializability
 - **cannot** be used *directly* to test for **view** serializability.
 - Extension to test for *view serializability* has:
 - ▶ **cost exponential** in the size of the precedence graph.
- The problem of **checking** if a schedule is **view** serializable:
 - falls in the class of:
 - ▶ **NP-complete** problems.
 - Thus existence of an efficient algorithm is:
 - ▶ **extremely unlikely**.
- However practical algorithms that:
 - just check some **sufficient conditions** for view serializability
 - ▶ **can still be used**.

Recoverable Schedules

Need to address the effect of *transaction failures* on concurrently running transactions:

- **Recoverable schedule:**
 - if a transaction T_j **reads** a data item *previously written* by a transaction T_i ,
 - then the **commit** operation of T_i appears **before** the **commit** operation of T_j .
- The following schedule (Schedule 11) is **not** recoverable
 - ▶ if T_9 **commits** immediately after the read

T_8	T_9
read(A) write(A) read(B)	read(A)

- If T_8 should **abort**, T_9 would:
 - **have read** (and possibly shown to the user) an *inconsistent database state*.
- Hence, database must ensure that schedules are recoverable.

Cascading Rollbacks

- **Cascading rollback:**
 - a **single transaction failure leads** to a series of transaction rollbacks.
 - Consider the following schedule where:
 - ▶ **none** of the transactions has yet **committed**
 (so the schedule is **recoverable**)

T_{10}	T_{11}	T_{12}
read(A) read(B) write(A)	read(A) write(A)	read(A)

- If T_{10} *fails*, T_{11} and T_{12} must also be *rolled back*.
- Can **lead** to the *undoing* of a *significant amount of work*

Cascadeless Schedules

- **Cascadeless schedules:**
 - cascading rollbacks cannot occur;
 - ▶ for each pair of transactions T_i and T_j such that
 - T_j **reads** a data item previously **written** by T_i ,
 - the **commit** operation of T_i appears:
 - » before the **read** operation of T_j .
- Every **cascadeless** schedule:
 - ▶ is also **recoverable**
- It is *desirable* to **restrict** the schedules to:

- ▶ those that are *cascadeless*

Concurrency Control

- A database must **provide** a mechanism that
 - will ensure that **all** possible schedules are:
 - ▶ either conflict or view serializable, and
 - ▶ are recoverable and
 - *preferably* cascadeless
- A policy in which only one transaction can execute at a time :
 - ▶ generates serial schedules,
 - ▶ but provides a *poor degree of concurrency*
 - Are serial schedules recoverable/cascadeless?
- Testing a schedule for serializability:
 - *after* it has executed
 - ▶ is a little too late!
- **Goal:**
 - to develop concurrency control protocols that:
 - ▶ will *assure serializability*.

Concurrency Control vs Serializability Tests

- Concurrency control protocols:
 - allow concurrent schedules,
 - ▶ but **ensure** that the schedules are conflict / view serializable, and
 - ▶ are recoverable and cascadeless.
- Concurrency control protocols generally
 - do not examine the precedence graph as it is being created
 - Instead a protocol **imposes a discipline** that
 - ▶ avoids nonserializable schedules.
 - We study such protocols in Chapter 16.
- **Different** concurrency control protocols:
 - ▶ provide **different** tradeoffs between
 - the *amount of concurrency* they allow and
 - the *amount of overhead* that they incur.
- Tests for serializability **help** us understand **why** a concurrency control protocol is **correct**.

Weak Levels of Consistency

- Some applications are willing to live with weak levels of consistency,
 - **allowing** schedules that are not serializable
 - *E.g.* a *read-only transaction* that wants to
 - ▶ get an approximate total balance of all accounts
 - *E.g.* database statistics computed
 - ▶ for query optimization can be approximate (why?)
 - Such transactions :
 - ▶ need not be serializable with respect to other transactions
- Tradeoff :
 - ▶ accuracy for

- performance

Levels of Consistency in SQL

- **Serializable** : default
- **Repeatable read** :
 - **only** committed records to be read,
 - **repeated reads** of same record must return same value.
 - However, a transaction **may not** be serializable –
 - it may **find** some records **inserted** by a transaction
 - but **not find others!**
- **Read committed** :
 - **only** committed records can be read,
 - but **successive reads** of record
 - may return different (but committed) values.
- **Read uncommitted** :
 - even uncommitted records may be read.
- Lower degrees of consistency useful for:
 - gathering *approximate information* about the database

Transaction Definition in SQL

- Data manipulation language must include a construct for:
 - specifying the **set of actions** that comprise a transaction.
- In SQL, a transaction begins implicitly.
- A transaction in SQL ends by:
 - **Commit work** commits current transaction and begins a new one.
 - **Rollback work** causes current transaction to abort.
- Levels of consistency specified by **SQL-92**:
 - **Serializable** — default
 - **Repeatable read**
 - **Read committed**
 - **Read uncommitted**

CONCURRENCY CONTROL

- **Lock-Based** Protocols
- **Timestamp-Based** Protocols
- **Validation-Based** Protocols
- Multiple **Granularity**
- **Multiversion** Schemes
- **Deadlock** Handling
- Insert and Delete **Operations**
- Concurrency in **Index** Structures

Lock-Based Protocols

- A **lock** is a *mechanism* to *control concurrent access* to a data item
- Data items can be locked in *two modes* :
 1. **exclusive** (X) *mode*. Data item can be both read as well as written. X-lock is requested using **lock-X** instruction.
 2. **shared** (S) *mode*. Data item can only be read. S-lock is requested using **lock-S** instruction.
- Lock requests are made to *concurrency-control manager*.
 - Transaction **can proceed** only after *request* is granted.
- **Lock-compatibility** matrix

	S	X
S	true	false
X	false	false

- A **transaction** may be granted a lock on an item
 - if the *requested lock* is **compatible** with
 - ▶ locks already held on the item by **other transactions**
- Any number of transactions can hold **shared** locks on an item,
 - **but** if any transaction holds an **exclusive** on the item
 - ▶ no other transaction may hold any lock on the item.
- If a lock cannot be granted,
 - the requesting transaction is made to **wait** till:
 - ▶ **all** incompatible locks held by other transactions have been released.
 - The lock is then granted.
- Example of a transaction performing locking:
 T_2 : **lock-S**(A);
 read (A);
 unlock(A);
 lock-S(B);
 read (B);
 unlock(B);
 display(A+B)
- Locking as above is **not sufficient** to guarantee serializability:
 - if A and B get updated in-between the read of A and B,
 - ▶ the displayed sum would be wrong.

- A **locking protocol** is a set of rules *followed by*
 - **all transactions** while requesting and releasing locks.
- Locking protocols **restrict** the set of possible **schedules**.

Pitfalls of Lock-Based Protocols

- Consider the **partial** schedule

T_3	T_4
lock-X(B) read(B) $B := B - 50$ write(B) lock-X(A)	 lock-S(A) read(A) lock-S(B)

- Neither T_3 nor T_4 can make progress:
 - executing **lock-S(B)** causes T_4 to wait for T_3 to release its lock on B , while
 - executing **lock-X(A)** causes T_3 to wait for T_4 to release its lock on A .
- Such a situation is *called* a **deadlock**.
 - To **handle** a deadlock one of T_3 or T_4 **must** be rolled back
 - ▶ and its locks released.
- The potential for deadlock exists in:
 - ▶ **most locking protocols**.
 - *Deadlocks are a necessary evil.*
- **Starvation** is also *possible*
 - ▶ **if concurrency control manager is *badly designed*.**

For example:

- A transaction may be waiting for an **X-lock** on an item,
 - ▶ while a *sequence* of other *transactions*
 - request and are granted an **S-lock** on the same item.
- The **same transaction** is
 - ▶ repeatedly rolled back due to deadlocks.
- *Concurrency control manager can be designed to:*
 - ▶ **prevent** starvation.

The Two-Phase Locking Protocol

- This is a protocol which ensures conflict-serializable schedules.
- Phase 1: **Growing** Phase
 - transaction may obtain locks
 - transaction *may not release* locks
- Phase 2: **Shrinking** Phase
 - transaction may release locks
 - transaction *may not obtain* locks
- The protocol **assures** serializability.
 - It can be proved that the transactions can be
 - ▶ serialized *in the order of* their **lock points**
 - ▶ i.e. the point where a transaction acquired its final lock.
- Two-phase locking *does not* ensure freedom from deadlocks
- Cascading roll-back is possible under two-phase locking.
 - To avoid this, follow a modified protocol :
 - ▶ called **strict two-phase locking**.
 - Here a transaction must hold **all** its *exclusive locks* :
 - till it commits/aborts.
- **Rigorous two-phase locking** is even stricter:
 - here **all locks** are **held**:
 - ▶ till commit/abort.
 - In this protocol transactions can be serialized
 - ▶ *in the order* in which they commit.
- There can be **conflict** serializable schedules that:
 - cannot be obtained **if** *two-phase locking* is used!
- However, in the **absence** of *extra information* (e.g., ordering of access to data),
 - *two-phase locking* is needed for:
 - ▶ **conflict** serializability in the following sense:
 - Given a transaction T_i that does not follow two-phase locking,
 - we can find a transaction T_j that uses two-phase locking,
 - and a schedule for T_i and T_j that
 - n **is not** conflict serializable.

Lock Conversions

- *Two-phase locking* with *lock conversions*:
 - First Phase:

- can acquire a lock-S on item
- can acquire a lock-X on item
- can convert a lock-S to a lock-X (upgrade)
- Second Phase:
 - can release a lock-S
 - can release a lock-X
 - can convert a lock-X to a lock-S (downgrade)
- This protocol assures *serializability*.
- But still **relies on the programmer**
 - ▶ to insert the *various* locking instructions.

Automatic Acquisition of Locks

- A transaction T_i issues the standard read/write instruction,
 - without *explicit locking calls*.
- The operation **read**(D) is processed as:


```

      if  $T_i$  has a lock on  $D$ 
      then
        read( $D$ )
      else begin
        if necessary wait until no other
        transaction has a lock-X on  $D$ 
        grant  $T_i$  a lock-S on  $D$ ;
        read( $D$ )
      end
      
```
- **write**(D) is processed as:

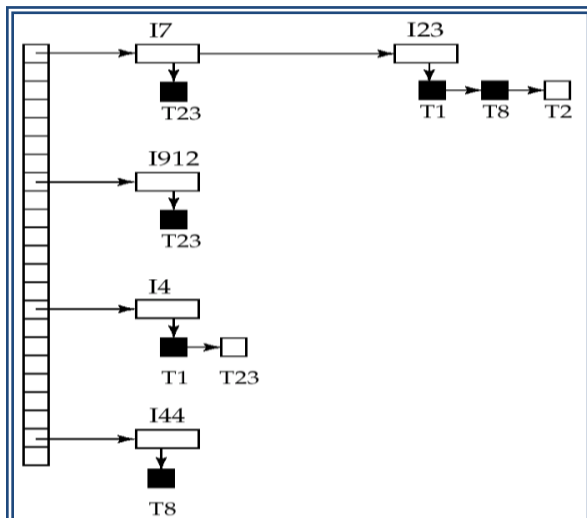

```

      if  $T_i$  has a lock-X on  $D$ 
      then
        write( $D$ )
      else begin
        if necessary wait until no other trans. has any lock on  $D$ ,
        if  $T_i$  has a lock-S on  $D$ 
        then
          upgrade lock on  $D$  to lock-X
        else
          grant  $T_i$  a lock-X on  $D$ 
        write( $D$ )
      end;
      
```
- All locks are released after commit or abort

Implementation of Locking

- A **lock manager** can be implemented as:
 - a separate process to which:
 - transactions send lock and unlock requests
- The lock manager **replies** to a lock request by
 - sending a lock grant messages
 - or a message asking the transaction to roll back, in case of a deadlock.
- The requesting transaction waits:
 - until its request is answered
- The lock manager maintains a data-structure called a **lock table** to
 - record granted locks and pending requests
- The lock table is usually *implemented* as:
 - an in-memory *hash table* indexed on:
 - the **name** of the **data item** being locked

Lock Table




- **Black rectangles** indicate granted locks,
 - **white ones** indicate waiting requests
- Lock table also records:
 - the **type of lock** granted or requested
- New request is added to the end of:
 - the **queue** of requests for the data item, and
 - granted **if** it is **compatible** with all earlier locks

- Unlock requests result in
 - the request being deleted, and
 - later requests are checked
 - ▶ to see **if** they can now be granted
- If transaction ***aborts***,
- **all** waiting or granted **requests** of the transaction are **deleted**
- lock manager may keep a *list of locks* held by each transaction,
- to implement this *efficiently*

Graph-Based Protocols

- **Graph-based** protocols are:
 - an ***alternative*** to *two-phase locking*
- **Impose** a partial ordering \rightarrow on the set $\mathbf{D} = \{d_1, d_2, \dots, d_h\}$ of ***all*** data ***items***.
 - If $d_i \rightarrow d_j$ then:
 - ▶ any transaction accessing both d_i and d_j
 - ▶ must access d_i *before* accessing d_j .
 - Implies that the set \mathbf{D} may now be viewed as
 - ▶ a directed acyclic graph,
 - ▶ called a ***database graph***.
- The ***tree-protocol*** is:
 - ▶ a simple kind of graph protocol.

Tree Protocol



- **Only exclusive** locks are allowed.
- The **first** lock by T_i may be on any data item.
 - Subsequently, a ***data Q*** can be locked by T_i
 - ▶ ***only if*** the **parent of Q** is currently locked by T_i .
- ***Data items*** may be unlocked at **any time**.
- The tree protocol ensures:

- **conflict** serializability as well as freedom from *deadlock*.
- Unlocking may occur earlier in the tree-locking protocol
 - ▶ than in the *two-phase locking* protocol.
 - shorter waiting times, and increase in *concurrency*
 - protocol is deadlock-free, no rollbacks are required
- Drawbacks
 - Protocol does not guarantee *recoverability* or *cascade freedom*
 - ▶ Need to introduce **commit dependencies**
 - n to ensure recoverability
 - Transactions may have to lock data **items** that they do not access.
 - ▶ increased locking overhead, and additional waiting time
 - ▶ potential decrease in concurrency
- **Schedules** not possible under *two-phase locking* are:
 - possible under tree protocol, and vice versa.


Timestamp-Based Protocols

- Each transaction is issued a timestamp when it enters the system.
 - If an old transaction T_i has time-stamp $TS(T_i)$,
 - ▶ a new transaction T_j is assigned time-stamp $TS(T_j)$
 - such that $TS(T_i) < TS(T_j)$.
- The protocol manages concurrent execution such that
 - the time-stamps determine the serializability order.
- In order to assure such behavior,
 - the protocol maintains for each data Q **two timestamp** values:
 - ▶ **W-timestamp**(Q) is the largest time-stamp of any transaction that executed **write**(Q) successfully.
 - ▶ **R-timestamp**(Q) is the largest time-stamp of any transaction that executed **read**(Q) successfully.
- The **timestamp ordering** protocol ensures that:
 - any conflicting **read** and **write** operations are:
 - ▶ executed in timestamp order.
- Suppose a transaction T_i issues a **read**(Q):
 - If $TS(T_i) < \mathbf{W}\text{-timestamp}(Q)$, then T_i needs to **read** a value of Q that was already overwritten.
 - ▶ Hence, the **read** operation is **rejected**, and T_i is rolled back.
 - If $TS(T_i) \geq \mathbf{W}\text{-timestamp}(Q)$, then the **read** operation is **executed**, and $\mathbf{R}\text{-timestamp}(Q)$ is set to the maximum of $\mathbf{R}\text{-timestamp}(Q)$ and $TS(T_i)$.
- Suppose that transaction T_i issues **write**(Q):

- If $TS(T_i) < \mathbf{R}\text{-timestamp}(Q)$, then the value of Q that T_i is producing was needed previously, and the system assumed that that value would never be produced.
 - ▶ Hence, the **write** operation is **rejected**, and T_i is rolled back.
- If $TS(T_i) < \mathbf{W}\text{-timestamp}(Q)$, then T_i is attempting to write an obsolete value of Q .
 - ▶ Hence, this **write** operation is **rejected**, and T_i is **rolled back**.
- Otherwise, the **write** operation is **executed**, and $\mathbf{W}\text{-timestamp}(Q)$ is set to $TS(T_i)$.

Example Use of the Protocol

A partial schedule for several data items for transactions with timestamps 1, 2, 3, 4, 5



Correctness of Timestamp-Ordering Protocol

- The timestamp-ordering protocol *guarantees serializability*
 - since **all** the arcs in the precedence graph are of the form:



Thus, there will be **no cycles** in the precedence graph

- Timestamp protocol **ensures** *freedom from deadlock* as
 - **no** transaction ever **waits**.
- But the schedule may **not** be *cascade-free*, and
 - may **not** even be *recoverable*.

Recoverability and Cascade Freedom

- Problem with timestamp-ordering protocol:
 - Suppose T_i aborts, but T_j has read a data item *written by* T_i
 - Then T_j must abort:
 - ▶ *if* T_j had been *allowed to commit earlier*,
 - the schedule is not recoverable!
 - ▶ *otherwise*, any T_k that has read a data item *written by* T_j must abort too!
 - This can lead to cascading rollbacks!
- Solution 1:
 - A transaction is structured such that:
 - ▶ **all** its **writes** are performed **at the end** of its processing
 - **All** writes of a transaction form an **atomic** action:
 - ▶ no transaction may **execute** while a transaction is being written
 - A **transaction** that aborts is restarted with a **new timestamp**
- Solution 2: Limited form of locking: wait for data to be **committed** before reading it
- Solution 3: Use **commit dependencies** to ensure recoverability

Thomas' Write Rule

- *Modified version* of the timestamp-ordering protocol in which:
 - obsolete **write** operations *may be ignored* under certain circumstances.
- When T_i attempts to write data item Q ,
 - if $TS(T_i) < W\text{-timestamp}(Q)$,
 - ▶ then T_i is *attempting to write an obsolete value* of $\{Q\}$.
 - Rather than *rolling back* T_i :
 - ▶ as the *timestamp ordering* protocol would have done,
 - ▶ this **{write}** operation can be ignored.
- *Otherwise*, this protocol is the same as the timestamp ordering protocol.

- Thomas' Write Rule **allows** *greater potential concurrency*:
 - **Allows** some *view-serializable* schedules that:
 - ▶ **are not** *conflict-serializable*.

Validation-Based Protocol

- Execution of transaction T_i is done in **three phases**:
 1. **Read and execution** phase: Transaction T_i writes:
 - only to *temporary local variables*
 2. **Validation** phase: Transaction T_i performs a "**validation** test":
 - to determine if local variables can be written *without violating serializability*.
 3. **Write** phase: If T_i is **validated**, the updates are applied to the database
 - otherwise, T_i is **rolled back**!
- The **three phases** of concurrently executing transactions can be **interleaved**,
 - but each transaction must go through the three phases in that order.
 - Assume for simplicity that:
 - ▶ the validation and write phase **occur together**, *atomically* and *serially*
 - ▶ I.e., only one transaction **executes** validation/write **at a time**.
- Also called as **optimistic concurrency control** since:
 - transaction **executes fully** in the hope that all will go well during validation
- Each transaction T_i has **3** timestamps:
 - Start(T_i) : the time when T_i started its execution
 - Validation(T_i): the time when T_i entered its validation phase
 - Finish(T_i) : the time when T_i finished its write phase
- Serializability order is **determined** by:
 - ▶ timestamp given at validation time,
 - ▶ to **increase** concurrency.
 - Thus TS(T_i) is given the value of Validation(T_i).
- This protocol is useful and
 - ▶ gives **greater** degree of concurrency
 - ▶ **if** probability of conflicts is **low**.

because:

- ▶ the serializability order is **not** pre-decided, and
- ▶ relatively **few** transactions will have to be rolled back.

Validation Test for Transaction T_j

- If for all T_i with $TS(T_i) < TS(T_j)$ either one of the following condition holds:
 - **finish(T_i) < start(T_j)** or
 - **start(T_j) < finish(T_i) < validation(T_j) and**
 - ▶ the set of data items written by T_i does not intersect with
 - ▶ the set of data items read by T_j .
 then validation succeeds and T_j can be **committed**.
- **Otherwise**, validation **fails** and T_j is **aborted**.

Justification:

- Either the first condition is satisfied, and
 - there is **no** overlapped execution,
- Or, the second condition is satisfied and
 - the writes of T_j do **not** affect reads of T_i since
 - ▶ they *occur after* T_i has finished its reads.
 - the writes of T_i do **not** affect reads of T_j since
 - ▶ T_j does *not read* any item written by T_i .

Schedule Produced by Validation

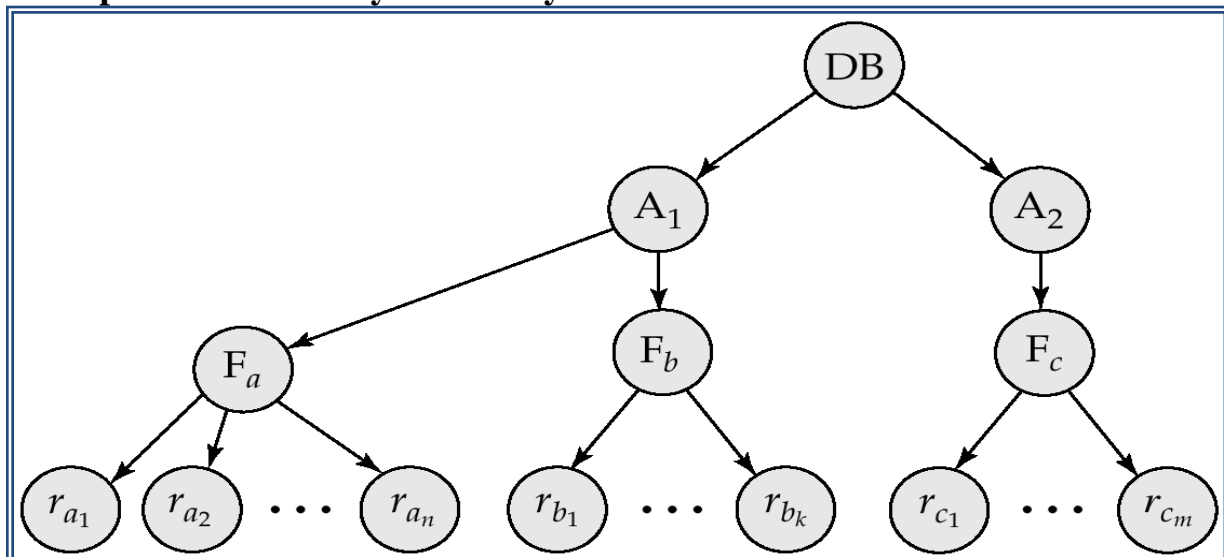
- Example of schedule produced using validation

T_{14}	T_{15}
read(B) read(A) (<i>validate</i>) display ($A+B$)	read(B) $B := B - 50$ read(A) $A := A + 50$ (<i>validate</i>) write (B) write (A)

Multiple Granularity

- Allow **data items** to be of **various sizes** and
 - **define** a hierarchy of data granularities, where :
 - the small granularities are **nested** within larger ones
- Can be represented graphically as a **tree**
 - but **don't confuse** with *tree-locking* protocol
- When a transaction **locks** a node in the tree *explicitly*,
 - it **implicitly locks** all the node's descendents in the same mode.
- Granularity of locking (level in tree where locking is done):
 - **fine** granularity (lower in tree):
 - high concurrency, high locking overhead
 - **coarse** granularity (higher in tree):
 - low locking overhead, low concurrency

Example of Granularity Hierarchy



The levels, starting from the coarsest (top) level are

- *database*
- *area*
- *file*
- *record*

Intention Lock Modes

- In addition to **S** and **X lock modes**, there are three additional lock modes with multiple granularity:

- **intention-shared (IS)**: indicates explicit locking at a lower level of the tree but only with shared locks.
- **intention-exclusive (IX)**: indicates explicit locking at a lower level with exclusive or shared locks
- **shared and intention-exclusive (SIX)**:
 - ▶ the subtree rooted by that node is locked explicitly in shared mode and
 - ▶ explicit locking is being done at a lower level with exclusive-mode locks.
- **intention locks allow** a higher level node to **be locked** in **S** or **X** mode
 - without having to check *all descendent* nodes.

Compatibility Matrix with Intention Lock Modes

- The compatibility matrix for all lock modes is:

	I	I	S	SIX	X
I	✓	✓	✓	✓	×
I	✓	✓	×	×	×
S	✓	×	✓	×	×
SIX	✓	×	×	×	×
X	×	×	×	×	×

Multiple Granularity Locking Scheme

- Transaction T_i can lock a node Q , using the following **rules**:
 1. The lock compatibility matrix **must be** observed.
 2. The root of the tree **must be** locked **first**, and may be locked in *any mode*.
 3. A node Q can be **locked** by T_i in **S** or **IS** mode **only if**:

- the parent of Q is *currently locked* by T_i in either **IX** or **IS** mode.
- 4. A node Q can be **locked** by T_i in **X**, **SIX**, or **IX** mode:
 - only if the parent of Q is *currently locked* by T_i in either **IX** or **SIX** mode.
- 5. T_i can **lock** a node **only if**:
 - it has **not** *previously unlocked* any node (i.e., T_i is **two-phase**).
- 6. T_i can **unlock** a node Q **only if**:
 - **none** of the children of Q are *currently locked* by T_i .
- Observe that locks are **acquired** in root-to-leaf order,
 - whereas they are **released** in leaf-to-root order.

Multiversion Schemes

- **Multiversion** schemes **keep** old versions of data item to increase **concurrency**.
 - Multiversion *Timestamp* Ordering
 - Multiversion *Two-Phase* Locking
- Each successful **write** results in:
 - ▶ the **creation** of a new version of the data item written.
- Use timestamps to label versions.
- When a **read**(Q) operation is issued,
 - **select** an appropriate version of Q *based on* the timestamp of the transaction, and return the value of the selected version.
- **reads** **never have to wait** as an appropriate version is returned immediately.

Multiversion *Timestamp* Ordering

- Each data item Q has a sequence of versions $\langle Q_1, Q_2, \dots, Q_m \rangle$. Each version Q_k contains three data fields:
 - **Content** -- the value of version Q_k .
 - **W-timestamp**(Q_k) -- timestamp of the transaction that created (wrote) version Q_k
 - **R-timestamp**(Q_k) -- largest timestamp of a transaction that successfully read version Q_k
- when a transaction T_i creates a new version Q_k of Q ,
 - Q_k 's W-timestamp and R-timestamp are initialized to **TS**(T_i).
- R-timestamp of Q_k is updated whenever a transaction T_j reads Q_k , and:
 - $\text{TS}(T_j) > \text{R-timestamp}(Q_k)$.
- Suppose that transaction T_i issues a **read**(Q) or **write**(Q) operation.

- Let Q_k denote the version of Q whose write timestamp is the largest write timestamp less than or equal to $TS(T_i)$.
- If transaction T_i issues a **read**(Q), then the value returned is the content of version Q_k .
- If transaction T_i issues a **write**(Q)
 - ▶ if $TS(T_i) < \text{R-timestamp}(Q_k)$, then transaction T_i is **rolled back**.
 - ▶ if $TS(T_i) = \text{W-timestamp}(Q_k)$, the contents of Q_k are overwritten
 - ▶ else a new version of Q is created.
- Observe that
 - Reads always succeed
 - A write by T_i is rejected if some other transaction T_j that (in the serialization order defined by the timestamp values) should read T_i 's write,
 - ▶ has already read a version created by a transaction older than T_i .
- Protocol guarantees **serializability**

Multiversion *Two-Phase* Locking

- Differentiates between read-only *transactions* and update *transactions*
- *Update transactions* **acquire** read and write **locks**, and:
 - **hold all locks** up to the end of the transaction.
 - ▶ That is, update transactions follow **rigorous two-phase locking**.
 - Each successful **write** results in:
 - ▶ the creation of a new version of the data item written.
 - each version of a data item has a **single** timestamp:
 - ▶ whose value is obtained from a counter **ts-counter** that is:
 - incremented during **commit** processing.
- *Read-only transactions* are assigned a timestamp by :
 - reading the current value of **ts-counter** before they **start** execution;
 - they **follow** the *multiversion timestamp-ordering* protocol:
 - ▶ for performing reads.
- When an update *transaction* wants to read a data item:
 - it obtains a shared lock on it, and reads the *latest version*.
- When it wants to write an item:
 - it obtains **X-lock** on;
 - it then creates a new version of the item and
 - ▶ **sets** this version's timestamp to ∞ .
- When update transaction T_i completes, commit processing occurs:
 - T_i **sets** timestamp on the versions it has created to **ts-counter** + 1

- T_i increments **ts-counter** by 1
- Read-only transactions that **start** after T_i increments **ts-counter**:
 - **will see** the values **updated** by T_i .
- Read-only transactions that **start** before T_i increments the **ts-counter**:
 - will see the value before the **updates** by T_i .
- Only *serializable schedules* are produced.

Deadlock Handling

- Consider the following two transactions:

T_1 : write (X) T_2 : write(Y)
 write(Y) write(X)

- Schedule with **deadlock**

T_1	T_2
X-lock on X write (X) wait for X-lock on Y	X-lock on Y write (X) wait for X-lock on X

Deadlock Handling

- System is **deadlocked** if there is a set of transactions such that:
 - every transaction in the set **is waiting** for another transaction in the set.
- **Deadlock prevention** protocols **ensure** that:
 - the system will **never** enter into a deadlock state.
- Some **prevention** strategies :
 - **Require** that each transaction **locks all** its data items:
 - ▶ before it begins **execution (predeclaration)**.
 - **Impose partial ordering** of **all data items** and require that:
 - ▶ a transaction can lock data items:
 - only in the order *specified* by the partial **order**
 - (graph-based protocol).


More Deadlock *Prevention* Strategies

- Following schemes use transaction timestamps:
 - for the sake of deadlock **prevention** alone.
- **wait-die** scheme — *non-preemptive*
 - **older** transaction may **wait** for younger one to release data item.
 - Younger transactions **never** wait for **older** ones;
 - ▶ they are **rolled back** instead.
 - a transaction may die several times before **acquiring** needed data item
- **wound-wait** scheme — *preemptive*
 - **older** transaction **wounds** (forces rollback) of younger transaction:
 - ▶ instead of waiting for it.
 - Younger transactions may **wait** for **older** ones.
 - may be **fewer rollbacks** than **wait-die** scheme.
- Both in **wait-die** and in **wound-wait** schemes,
 - a rolled back transactions is **restarted** with its **original** timestamp.
 - **Older** transactions thus have **precedence** over newer ones,
 - and starvation is hence **avoided**.
- **Timeout-Based** Schemes :
 - a transaction waits for a **lock**:
 - ▶ only for a specified amount of **time**.
 - ▶ After that, the wait **times out** and:
 - the transaction is rolled back.
 - thus deadlocks are **not** possible
 - **simple** to implement;
 - but starvation is **possible**.
 - Also difficult to determine *good value* of the **timeout** interval.


Deadlock *Detection*

- **Deadlocks** can be described as a *wait-for graph*,
 - which consists of a pair $G = (V, E)$,
 - ▶ V is a set of vertices (all the transactions in the system)
 - ▶ E is a set of edges; each element is an ordered pair $T_i \rightarrow T_j$.
- If $T_i \rightarrow T_j$ is in E , then there is a directed edge from T_i to T_j ,
 - implying that T_i is waiting for T_j to release a data item.
- When T_i requests a data item currently being held by T_j ,
 - then the edge $T_i \rightarrow T_j$ is inserted in the wait-for graph.
 - This edge is removed only when:
 - ▶ T_j is no longer holding a *data item* needed by T_i .

- The system is in a **deadlock** state:
 - *if and only if* the wait-for graph **has a cycle**.
- **Must invoke** a **deadlock-detection** algorithm periodically:
 - to look for **cycles**.



Wait-for graph **without** a cycle



Wait-for graph **with** a cycle

Deadlock Recovery

- When deadlock is detected :
 - Some transaction will **have to be rolled back** (made a victim):
 - ▶ to **break** deadlock.
 - ▶ **Select** that transaction as victim that:
 - will **incur minimum** cost.
 - Rollback -- determine how far to roll back transaction
 - ▶ **Total** rollback:
 - **Abort** the transaction and then **restart** it.
 - ▶ More effective to roll back transaction **only**:
 - **as far as necessary to break** deadlock.
 - Starvation happens if **same** transaction is:
 - ▶ always chosen as **victim**.
 - ▶ **Include** the *number of rollbacks* in:
 - the **cost factor** to **avoid** starvation

Insert and Delete: *phantom* phenomenon

- If **two-phase** locking is used :
 - A **delete** operation may be performed only if the transaction deleting the tuple has an **X-lock** on the tuple to be deleted.
 - A transaction that **inserts** a new tuple into the database is given an **X-lock** on the tuple
- **Insertions and deletions can lead to the phantom phenomenon:**

- T_i **scans** a relation (e.g., **find all** accounts in Perryridge) and
- T_j **inserts** a tuple in the relation (e.g., insert a new account at Perryridge)
 - ▶ may conflict in spite of **not accessing** any tuple in common.
- If only **tuple locks** are used:
 - ▶ **non-serializable** schedules **can result**:
 - ▶ the **scan** transaction **may not see** the new account,
 - ▶ (yet may be serialized before the insert transaction).
- The transaction **scanning** the relation is reading *information* that:
 - ▶ indicates *what tuples* the relation contains,
 - while a transaction **inserting** a tuple updates the *same information*.
 - This *information should* be **locked**.
- One **solution**:
 - **Associate** a *data item* with the relation,
 - ▶ to **represent** the *information* about what tuples the relation contains.
 - Transactions **scanning** the relation:
 - ▶ **acquire** a *shared lock* in the *data item*,
 - Transactions **inserting** or **deleting** a tuple:
 - ▶ **acquire** an *exclusive lock* on the *data item*.
 - ▶ (Note: locks on the *data item* do not conflict with locks on individual tuples.)
- Above protocol provides very **low** concurrency for insertions/deletions.
- **Index locking** protocols provide higher **concurrency**
 - While **preventing** the phantom phenomenon,
 - by **requiring** *locks* on certain *index buckets*.

Index Locking Protocol

- Every relation **must have at least** one index.
- **Access** to a relation **must be** made:
 - **only** through **one** of the indices on the relation.
- A transaction T_i that performs a lookup:
 - **must lock all** the *index buckets* that it accesses, in S-mode.
- A transaction T_i **may not** insert a tuple t_i into a relation r
 - without updating **all** indices to r .
- T_i must perform a lookup on **every** index to find:
 - **all** *index buckets* that could have possibly contained a pointer to tuple t_i ,

- ▶ had it existed already, and
- **obtain** locks in **X-mode** on *all these index buckets*.
- T_i **must** also **obtain** locks in **X-mode** on *all index buckets* that it **modifies**.
- The rules of the **two-phase** locking protocol must be observed
 - **Guarantees** that *phantom phenomenon won't occur!*

Weak Levels of Consistency

- **Degree-two consistency:**
 - differs from **two-phase** locking in that:
 - ▶ **S-locks** *may be released* at any time, and
 - ▶ **locks** *may be acquired* at any time
 - **X-locks** *must be held till* end of transaction
 - Serializability is **not guaranteed**,
 - ▶ *programmer* must ensure that **no erroneous** database state will **occur!**
- **Cursor stability (CS):**
 - *Special case of degree-two consistency*
 - For **reads**, each tuple is:
 - ▶ locked,
 - ▶ read, and
 - ▶ [lock is immediately released]
 - **X-locks** are *held till* end of transaction

Weak Levels of Consistency in SQL

- SQL allows **non-serializable** executions:
 - **Serializable:** is the default
 - **Repeatable read:**
 - ▶ **allows only committed records** to be read, and
 - ▶ **repeating** a read should return the same value
 - (so read locks should be retained)
 - ▶ However, the *phantom* phenomenon **need not be prevented**
 - T_1 may see some **records inserted** by T_2 ,
 - but may not see **others inserted** by T_2
 - **Read committed:**
 - ▶ same as **degree-two** consistency,
 - ▶ but *most systems* implement it as cursor-stability
 - **Read uncommitted:**
 - ▶ **allows even uncommitted data** to be **read**

Concurrency in *Index Structures*

- **Indices** are **unlike** other database items in that:
 - ▶ their **only** job is to **help** in **accessing data**.
- **Index-structures** are typically **accessed very often**,
 - ▶ *much more than* **other** database **items**.
- **Treating index-structures like other database items** leads to:
 - **low concurrency**.
 - **Two-phase locking** on an **index** may result in:
 - ▶ transactions **executing** practically one-at-a-time!
- It is **acceptable** to **have nonserializable** concurrent **access** to an index:
 - as long as the **accuracy** of the index is **maintained**.
- In particular, the exact values read in an **internal node** of a **B⁺-tree** are **irrelevant** so long as **we land up** in the **correct** leaf node.
- There are *index concurrency protocols* where:
 - **locks** on **internal nodes** are **released** early,
 - ▶ and **not** in a two-phase fashion.

Example of *index concurrency protocol*:

- Use **crabbing** instead of **two-phase** locking on the nodes of the **B⁺-tree**, as follows.
- During search/insertion/deletion:
 - First **lock** the **root** node in shared mode.
 - After **locking all required children** of a **node** in shared mode,
 - ▶ release the **lock** on **the node**.
 - During insertion/deletion,
 - ▶ **upgrade** leaf node **locks** to exclusive mode.
 - When splitting or coalescing requires changes to a **parent**,
 - ▶ **lock** the **parent** in exclusive mode.
- Above protocol can cause excessive **deadlocks**.
 - Better protocols are available;
 - ▶ E.g the **B-link tree** protocol

RECOVERY SYSTEM

- Failure Classification
- Storage Structure
- Recovery and Atomicity
- Log-Based Recovery
- Shadow Paging
- Recovery With Concurrent Transactions

- Buffer Management
- Failure with Loss of Nonvolatile Storage
- Advanced Recovery Techniques
- ARIES Recovery Algorithm
- Remote Backup Systems

Failure Classification

- **Transaction failure :**
 - **Logical errors:** transaction cannot complete due to some internal error condition
 - **System errors:** the database system must terminate an active transaction due to an error condition (e.g., deadlock)
- **System crash:** a *power* failure or other *hardware* or *software* failure causes the system to crash.
 - **Fail-stop assumption:** non-volatile *storage contents* are assumed to *not be corrupted* by system crash
 - ▶ Database systems have *numerous integrity checks* to *prevent corruption* of disk data
- **Disk failure:** a head crash or similar disk failure destroys all or part of disk storage
 - **Destruction** is assumed to be **detectable**: disk drives use checksums to detect failures

Recovery Algorithms

- Recovery algorithms are *techniques* to **ensure** database **consistency** and transaction **atomicity** and **durability** despite failures
- Recovery algorithms have *two parts*:
 - Actions taken during normal transaction processing to **ensure** enough information **exists** to recover from failures
 - Actions taken after a failure to **recover** the database contents to a **state** that **ensures** atomicity, consistency and durability

Storage Structure

- **Volatile storage:**
 - does not survive system crashes
 - Ex: main memory, cache memory
- **Nonvolatile storage:**
 - survives system crashes

- Ex: disk, tape, flash memory,
non-volatile (battery backed up) RAM
- **Stable storage:**
 - a *mythical* form of storage that survives **all failures**
 - *approximated by* maintaining multiple copies on distinct nonvolatile media

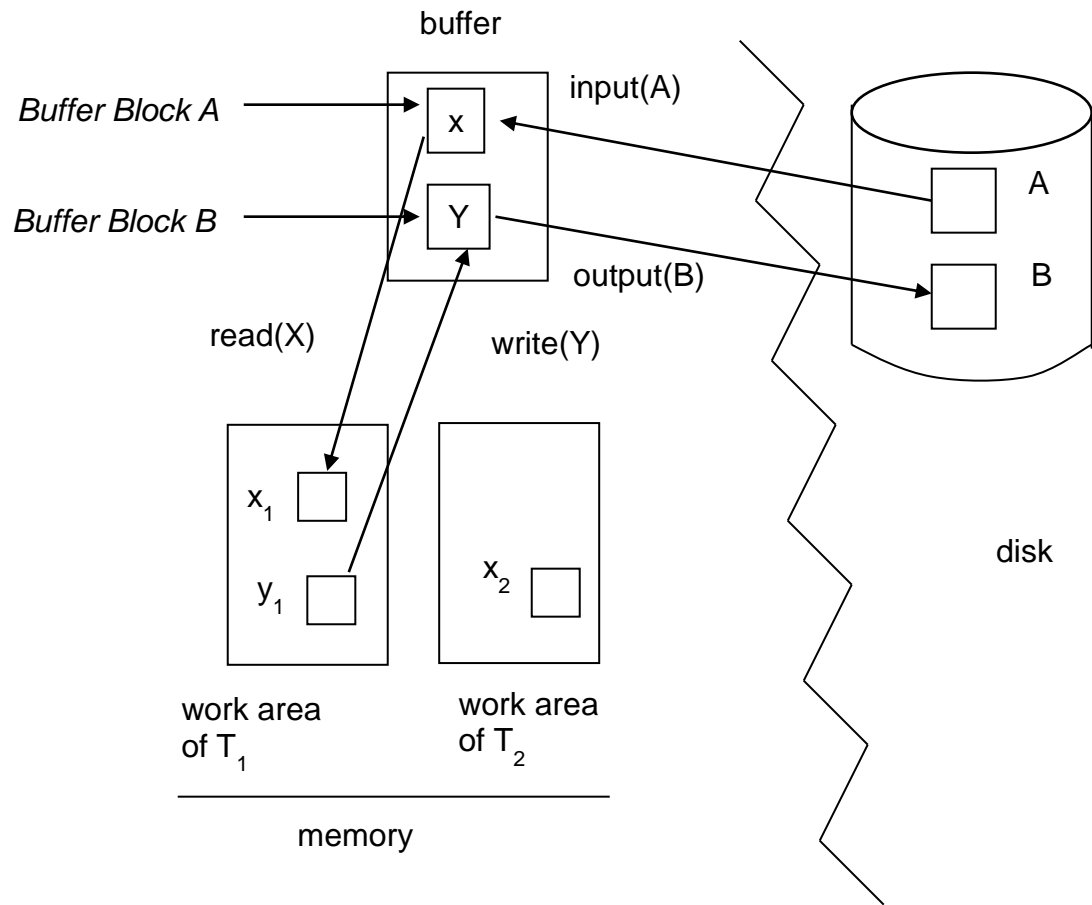
Stable-Storage Implementation

- Maintain **multiple** copies of each block on **separate** disks
 - copies can be at **remote** sites to *protect against disasters* such as **fire** or **flooding**.
- Failure during data transfer can still result in **inconsistent copies**:
 - Block transfer can result in:
 - ▶ Successful completion
 - ▶ Partial **failure**: destination block has incorrect information
 - ▶ Total **failure**: destination block was never updated
- Protecting storage media from **failure** during data transfer (one solution):
 - Execute output operation as follows (assuming **two copies** of each block):
 - ▶ **Write** the information onto the first physical block.
 - ▶ **When** the first write successfully completes, **write** the same information onto the second physical block.
 - ▶ The output is completed only after the second write successfully completes.
- Copies of a block **may differ** due to failure during output operation.
To **recover** from failure:
 1. First **find** inconsistent blocks:
 1. *Expensive solution*: **Compare** the two copies of every disk block.
 2. *Better solution*:
 - Record in-progress disk writes on non-volatile storage (Non-volatile RAM or special area of disk).
 - Use this information during recovery to find blocks that may be inconsistent, and only compare copies of these.
 - Used in hardware RAID systems
 2. **If** either copy of an inconsistent block is **detected** to have an error (bad checksum), **overwrite** it by the other copy. **If** both have **no error**, but are different, overwrite the second block by the first block.

Data Access

- **Physical blocks** are those blocks residing on the disk.
- **Buffer blocks** are the blocks residing temporarily in main memory.
- Block **movements** between *disk* and *main* memory are initiated through the following two operations:
 - **input**(B) transfers the *physical block* B to main memory.
 - **output**(B) transfers the *buffer block* B to the disk, and replaces the appropriate physical block there.
- Each transaction T_i has its **private** *work-area* in which *local copies* of **all** data items accessed and updated by it are kept.
 - T_i 's local copy of a data item X is called x_i .
- We assume, for *simplicity*, that each data item **fits** in, and is stored inside, a single block.
- Transaction *transfers* data items between *system* buffer blocks and its *private work-area* using the following operations :
 - **read**(X) assigns the value of data item X to the *local variable* x_i .
 - **write**(X) assigns the value of local variable x_i to data item $\{X\}$ in the buffer block.
 - both these *commands* may *necessitate* the issue of an **input**(B_X) instruction before the assignment, if the block B_X in which X resides is not already in memory.
- Transactions
 - Perform **read**(X) while accessing X for the first time;
 - All subsequent accesses are to the *local* copy.
 - After last access, transaction executes **write**(X).
- **output**(B_X) need **not** immediately follow **write**(X).
 - System can perform the **output** operation *when it deems fit*.

Example of Data Access



Recovery and Atomicity

- **Modifying** the database **without** ensuring that the **transaction** will **commit**
 - ▶ may **leave** the database in an inconsistent state.
- Consider transaction T_i that transfers \$50 from account *A* to account *B*; goal is:
 - **either** to perform **all** database modifications made by T_i
 - **or none** at all.
- **Several output** operations may be required for T_i (to output *A* and *B*).
 - A failure may occur **after one** of these *modifications* have been made
 - but before all of them are made.
- To **ensure** atomicity despite failures,
 - we first **output** information **describing** the modifications to stable storage without *modifying* the *database* itself.

- We study two approaches:
 - **log-based recovery**, and
 - **shadow-paging**
- We assume (initially) that *transactions run serially*, that is, one after the other.

Log-Based Recovery

- A **log** is kept on stable storage.
 - The log is a sequence of **log records**, and
 - ▶ maintains a record of update activities on the database.
- When transaction T_i starts, it registers itself by writing a **<T_i start>** log record
- *Before* T_i executes **write(X)**, a log record **<T_i, X, V₁, V₂>** is written, where V_1 is the value of X before the write, and V_2 is the value to be written to X.
 - Log record notes that T_i has performed a write on data item X_j :
 - ▶ X_j had value V_1 before the write, and will have value V_2 after the write.
- When T_i finishes its *last statement*, the log record **<T_i commit>** is written.
- We *assume* for now that log records are written *directly* to *stable storage*
 - (that is, they are **not buffered**)
- Two approaches using logs:
 - **Deferred** database modification
 - **Immediate** database modification

Deferred Database Modification

- The **deferred** database modification scheme records all modifications to the log,
 - but defers **all the writes** to *after partial commit*.
- Assume that *transactions execute serially*
- Transaction starts by writing **<T_i start>** record to log.
- A **write(X)** operation results in a log record **<T_i, X, V>** being written, where V is the new value for X
 - Note: **old value** is not needed for this scheme
- The **write** is **not performed** on X at this time, but is deferred.
- When T_i partially commits, **<T_i commit>** is written to the log
- Finally, the log records are **read and used** to:
 - actually **execute** the previously deferred **writes**.
- During **recovery** after a crash, a transaction needs to be **redone**:

- *if and only if both* $\langle T_i \text{ start} \rangle$ and $\langle T_i \text{ commit} \rangle$ **are** there in the log.
- Redoing a transaction T_i (**redo T_i**) sets the value of **all data** items:
 - updated by the transaction to the new values.
- Crashes can occur while:
 - the transaction is executing the original updates,
 - **or:** while **recovery action** is being taken
- Ex: transactions T_0 and T_1 (T_0 executes before T_1):

T_0 : read (A)	T_1 : read (C)
A:- A - 50	C:- C - 100
Write (A)	write (C)
read (B)	
B:- B + 50	
write (B)	
- Below we show the log as it appears at three instances of time:

$\langle T_0 \text{ start} \rangle$ $\langle T_0, A, 950 \rangle$ $\langle T_0, B, 2050 \rangle$	$\langle T_0 \text{ start} \rangle$ $\langle T_0, A, 950 \rangle$ $\langle T_0, B, 2050 \rangle$ $\langle T_0 \text{ commit} \rangle$ $\langle T_1 \text{ start} \rangle$ $\langle T_1, C, 600 \rangle$	$\langle T_0 \text{ start} \rangle$ $\langle T_0, A, 950 \rangle$ $\langle T_0, B, 2050 \rangle$ $\langle T_0 \text{ commit} \rangle$ $\langle T_1 \text{ start} \rangle$ $\langle T_1, C, 600 \rangle$ $\langle T_1 \text{ commit} \rangle$
(a)	(b)	(c)

- If log on stable storage at **time of crash** is as in case:
 - (a) No redo actions need to be taken
 - (b) **redo(T_0)** must be performed since $\langle T_0 \text{ commit} \rangle$ is present
 - (c) **redo(T_0)** must be performed followed by **redo(T_1)** since $\langle T_0 \text{ commit} \rangle$ and $\langle T_i \text{ commit} \rangle$ are present

Immediate Database Modification

- The **immediate database modification** scheme allows:
 - database **updates** of an **uncommitted transaction** to be made:
 - ▶ *as the writes are issued*
 - since, undoing may be needed,
 - ▶ update **logs** must have **both** old value and new value
- Update log record **must be written before** database item is written
 - We assume that the **log** record is *output directly* to stable storage
 - ▶ Can be extended to **postpone** log record output,

- Prior to execution of an **output**(B) operation for a *data block* B ,
 - ▶ **all** log records corresponding to items B **must be flushed** to stable storage
- Output of updated blocks can take place at:
 - *any time* before or after transaction **commit**
- Order in which blocks are output *can be different* from:
 - the order in which they are written.

Log	Write	Output
$\langle T_0 \text{ start} \rangle$		
$\langle T_0, A, 1000, 950 \rangle$		
$T_0, B, 2000, 2050$		
	$A = 950$	
	$B = 2050$	
$\langle T_0 \text{ commit} \rangle$		
$\langle T_1 \text{ start} \rangle$		
$\langle T_1, C, 700, 600 \rangle$		
	$C = 600$	
		B_B, B_C
$\langle T_1 \text{ commit} \rangle$		
		B_A

- Note: B_X denotes block containing X .
- **Recovery** procedure has *two operations* instead of one:
 - **undo**(T_i) restores the value of all data items updated by T_i to their old values, *going backwards* from the last log record for T_i
 - **redo**(T_i) sets the value of all data items updated by T_i to the new values, *going forward* from the first log record for T_i
- **Both** operations must be **idempotent**
 - That is, even if the operation is executed *multiple times* the effect is the same as if it is executed once
 - ▶ Needed since operations may get *re-executed* during recovery
- When recovering after failure:

- Transaction T_i needs to be undone if the log contains the record $\langle T_i \text{ start} \rangle$,
 - ▶ but does not contain the record $\langle T_i \text{ commit} \rangle$.
- Transaction T_i needs to be redone if the log contains:
 - ▶ both the record $\langle T_i \text{ start} \rangle$ and the record $\langle T_i \text{ commit} \rangle$.
- **Undo** operations are **performed first**, then **redo** operations.

Immediate DB Modification Recovery

Below we show the log as it appears at three instances of time:

$\langle T_0 \text{ start} \rangle$	$\langle T_0 \text{ start} \rangle$	$\langle T_0 \text{ start} \rangle$
$\langle T_0, A, 1000, 950 \rangle$	$\langle T_0, A, 1000, 950 \rangle$	$\langle T_0, A, 1000, 950 \rangle$
$\langle T_0, B, 2000, 2050 \rangle$	$\langle T_0, B, 2000, 2050 \rangle$	$\langle T_0, B, 2000, 2050 \rangle$
	$\langle T_0 \text{ commit} \rangle$	$\langle T_0 \text{ commit} \rangle$
	$\langle T_1 \text{ start} \rangle$	$\langle T_1 \text{ start} \rangle$
	$\langle T_1, C, 700, 600 \rangle$	$\langle T_1, C, 700, 600 \rangle$
		$\langle T_1 \text{ commit} \rangle$
(a)	(b)	(c)

Recovery actions in each case above are:


- (a) **undo** (T_0): B is restored to 2000 and A to 1000.
- (b) **undo** (T_1) and **redo** (T_0): C is restored to 700, and
then A and B are set to 950 and 2050 respectively.
- (c) **redo** (T_0) and **redo** (T_1): A and B are set to 950 and 2050 respectively.
Then C is set to 600

Checkpoints

- **Problems** in recovery procedure *as discussed earlier* :
 1. searching the entire log is **time-consuming**
 2. we might *unnecessarily redo* transactions which have *already output* their updates to the database.
- **Streamline recovery** procedure by periodically performing **checkpointing**
 1. Output **all log records** currently residing in main memory onto stable storage.
 2. Output **all** modified buffer blocks to the disk.

3. Write a log record **< checkpoint >** onto stable storage.
- During recovery we need to **consider** only the most recent transaction T_i that started before the checkpoint, and transactions that started after T_i .
 1. Scan **backwards** from end of log to find the most recent **< checkpoint >** record
 2. Continue scanning **backwards** till a record **< T_i start >** is found.
 3. Need only **consider** the part of log **following** above **start** record.
Earlier part of log can be ignored during recovery, and can be erased whenever desired.
 4. For **all transactions** (starting from T_i or later) with no **< T_i commit >**, execute **undo(T_i)**. (Done only in case of immediate modification.)
 5. Scanning **forward** in the log, for **all transactions** starting from T_i or later with a **< T_i commit >**, execute **redo(T_i)**.

Example of Checkpoints



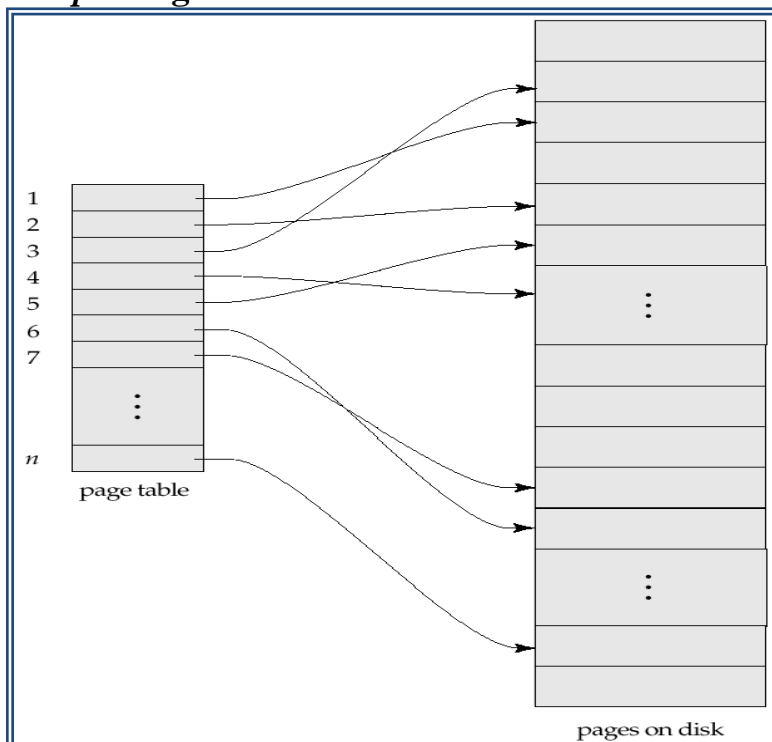
- T_1 can be ignored (*updates already output to disk due to checkpoint*)
- T_2 and T_3 redone.
- T_4 undone

Shadow Paging

- **Shadow paging** is an **alternative** to *log-based recovery*;
 - this scheme is useful **if** *transactions execute serially*
- Idea: maintain **two** page tables during the lifetime of a transaction:
 - the **current** page table, and the **shadow** page table
- Store the **shadow** page table in nonvolatile storage,
 - such that state of the database prior to transaction execution may be recovered.

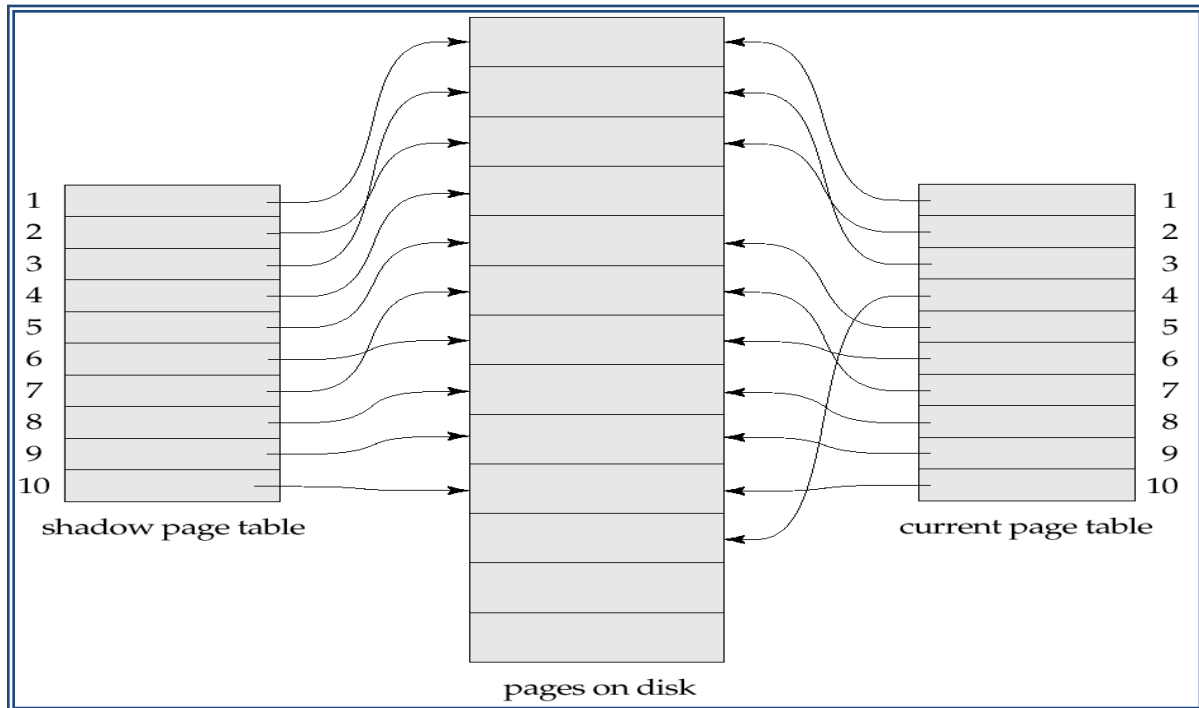
- *Shadow* page table is **never modified** during execution
- To start with, **both** the page tables are **identical**.
 - **Only current** page table is used for:
 - ▶ data item *accesses* during execution of the transaction.
- Whenever **any page** is about to be written for the first time
 - A **copy** of this page is **made** onto an unused page.
 - The **current** page table is then **made** to point to the copy
 - The update is **performed** on the copy

Sample Page Table



Example of *Shadow* Paging

Shadow and **current** page tables after *write* to page **4**



- To **commit** a transaction :

1. Flush **all** *modified* pages in main memory to **disk**
2. Output **current** page table to disk
3. Make the **current** page table the new **shadow** page table, as follows:
 - keep a **pointer** to the **shadow** page table at a fixed (known) location on **disk**.
 - to make the **current** page table the new **shadow** page table,
 - ▶ **simply update** the **pointer** to point to:
 - **current** page table on **disk**
- **Once pointer** to **shadow** page table has been *written*,
 - ▶ transaction is **committed**.
- **No recovery** is **needed** after a crash:
 - ▶ new transactions can **start** right away, using the **shadow** page table.
- **Pages** *not pointed* to from **current/shadow** page table:
 - ▶ should be **freed** (garbage collected).
- Advantages of *shadow-paging* over *log-based* schemes
 - **no** overhead of **writing log records**
 - recovery is **trivial**

- Disadvantages :
 - Copying the entire page table is very **expensive**
 - ▶ Can be reduced by using a page table structured like a B⁺-tree
 - No need to copy entire tree, only need to copy paths in the tree that lead to updated leaf nodes
 - Commit **overhead** is high even with above extension
 - ▶ Need to **flush** every updated page, and page table
 - Data gets **fragmented** (related pages get separated on disk)
 - After every transaction completion, the database pages containing old versions of modified data need to be **garbage** collected
 - **Hard** to extend algorithm to **allow** transactions to run **concurrently**
 - ▶ Easier to extend log based schemes

Recovery With *Concurrent* Transactions

- We **modify** the *log-based* recovery schemes to:
 - ▶ allow **multiple** transactions to execute **concurrently**.
 - **All** transactions share a **single** disk buffer and a **single** log
 - A buffer block **can** have data items updated by **one** or **more** transactions
- We assume *concurrency control* using **strict two-phase locking**;
 - i.e. the **updates** of uncommitted transactions:
 - ▶ **should not be visible** to other transactions
 - ▶ Otherwise how to perform undo if T1 updates A, then T2 updates A and commits, and finally T1 has to abort?
- Logging is done as described *earlier*.
 - Log records of different transactions *may be interspersed* in the log.
- The checkpointing technique and *actions* taken on *recovery*:
 - ▶ have to be changed!
 - since, *several transactions may be active*,
 - ▶ **when** a checkpoint is performed.
- 1. **Checkpoints** are performed as before,
 - ▶ except that the *checkpoint log record* is now of the form:
 - < **checkpoint** *L* >
 - where *L* is the list of transactions active at the time of the checkpoint
 - We assume **no** updates are in progress ,
 - ▶ **while** the checkpoint is carried out
- When the system **recovers** from a crash, it first does the following:

1. **Initialize** *undo-list* and *redo-list* to empty
 2. **Scan** the log **backwards** from the end,
 - stopping when the first **<checkpoint L>** record is found.
- For each record found **during** the *backward scan*:
- if the record is **<T_i commit>**, **add** *T_i* to *redo-list*
 - if the record is **<T_i start>**, then **if** *T_i* is not in *redo-list*, **add** *T_i* to *undo-list*
3. For **every** *T_i* in *L*, if *T_i* is not in *redo-list*, **add** *T_i* to *undo-list*
- At this point *undo-list* consists of *incomplete transactions* which must be **undone**, and *redo-list* consists of *finished transactions* that must be **redone**.
 - Recovery now continues as follows:
 4. **Scan** again the log **backwards** from the end:
 - During the scan, perform **undo** for each log record that belongs to a transaction in *undo-list*.
 - Stop the scan when **<T_i start>** records have been *encountered* for **all** *T_i* in *undo-list*
 5. **Locate** the most recent **<checkpoint L>** record.
 6. **Scan** the log **forwards** from the **<checkpoint L>** record:
 - During the scan, perform **redo** for each log record that belongs to a transaction on *redo-list*
 - Stop the scan at the *end of the log*.

Example of Recovery

- Go over the steps of the recovery algorithm on the following log:

```

<T0 start>
<T0, A, 0, 10>
<T0 commit>
<T1 start>          /* Scan in Step 4 stops here */
<T1, B, 0, 10>
<T2 start>
<T2, C, 0, 10>
<T2, C, 10, 20>
<checkpoint {T1, T2}>
<T3 start>
<T3, A, 10, 20>
<T3, D, 0, 10>
<T3 commit>

```

Log Record Buffering

- Log record *buffering*:
 - Log records are buffered in **main memory**,
 - ▶ instead of being *output directly* to stable storage.
 - Log records are *output* to stable storage:
 - ▶ when a block of log records in the *buffer* is **full**,
 - ▶ or a **log force** operation is executed.
- Log force is performed to **commit** a transaction by:
 - *forcing* all its log records (including the commit record) to stable storage.
- **Several** log records can thus be output using a *single output* operation,
 - **reducing** the *I/O cost*.
- The **rules** below **must be followed if** log records are buffered:
 - Log records are *output* to stable storage:
 - ▶ in the **order** in which they are created.
 - Transaction T_i enters the **commit** state:
 - ▶ **only** when the log record $\langle T_i \text{ commit} \rangle$ has been *output* to stable storage.
 - Before a **block** of data in main memory is *output* to the database,
 - ▶ **all** log records pertaining to data in that **block**:
 - n **must** have been *output* to stable storage.
 - ▶ This rule is called the **write-ahead** logging or **WAL** rule:
 - n Strictly speaking **WAL** only requires *undo information* to be *output*

Database Buffering

- Database maintains an **in-memory buffer** of data blocks
 - When a *new block* is needed, if **buffer** is **full**:
 - ▶ an existing *block* needs to be removed from buffer
 - If the *block* chosen for removal has been **updated**,
 - ▶ it **must** be output to disk
- As a result of the **write-ahead** logging rule,
 - if a *block* with *uncommitted updates* is *output* to disk,
 - ▶ log records with *undo information* for the updates are:
 - *output to the log* on stable storage **first**.
- **No updates** should be in progress on a **block**:
 - when it is **output** to disk.
 - Can be ensured *as follows*.

- Before **writing** a *data item*,
 - transaction acquires **exclusive lock** on **block** containing the data item
 - Lock can be **released** once the *write is completed*.
 - ▶ Such *locks* held for *short duration* are called **latches**.
- Before a **block** is *output to disk*,
 - the system acquires an **exclusive latch** on the **block**
 - **Ensures no update** can be *in progress* on the **block**
- Database **buffer** can be implemented:
 - either **in** an area of **real** main-memory **reserved for the database**,
 - or **in virtual** memory
- Implementing **buffer** in reserved main-memory has **drawbacks**:
 - Memory is *partitioned before-hand* between database buffer and applications,
 - ▶ *limiting flexibility*.
 - **Needs may change**, and
 - ▶ although operating system **knows best**:
 - **how** memory *should be divided up* at **any time**,
 - **it cannot change** the partitioning of memory.
- Database **buffers** are:
 - **generally implemented in virtual** memory
 - ▶ in spite of some **drawbacks** (*as follows*):
- When operating system **needs** to **evict** a page that has been modified,
 - ▶ to make **space** for another page,
 - the page is written to **swap** space on **disk**.
- When database **decides** to **write** buffer page to **disk**,
 - buffer page *may be* in **swap** space, and
 - ▶ *may have* to be **read** from **swap** space on **disk** and
 - ▶ **output** to the database on disk,
 - ▶ resulting in **extra I/O!**
 - Known as ***dual paging** problem*.
 - **Ideally** when *swapping out* a **database** buffer page,
 - ▶ operating system *should pass control to* database,
 - which in turn **outputs** page to database (space)
 - instead of to **swap** space
 - ▶ (making sure to **output** log records **first**)
 - ***Dual paging*** can thus **be avoided**,
 - ▶ **but** common operating systems
 - do **not support** such *functionality*.

Failure with *Loss of Nonvolatile Storage*

- So far we assumed:
 - ▶ **no loss** of **non-volatile** storage
- *Technique* similar to checkpointing used
 - ▶ to deal with **loss** of **non-volatile** storage.
- Periodically **dump** the entire content of the database to **stable** storage
- No transaction may be **active** *during the dump* procedure;
 - ▶ a **procedure** similar to checkpointing **must** take place.
 - Output **all** log records:
 - ▶ currently residing in main memory onto stable storage.
 - Output **all** buffer blocks onto the **disk** (*i.e., database*).
 - Copy the **contents** of the database to stable storage (*i.e., archival dump*).
 - Output a record **<dump>** to **log** on stable storage.
- To **recover** from *disk failure*
 - restore **database** from most recent **dump**.
 - Consult the **log** and **redo all** transactions that committed after the **dump**
- Can be **extended** to allow **transactions** to be **active** during **dump**;
 - known as **fuzzy dump** or **online dump**
 - Will study *fuzzy checkpointing* later.

Advanced Recovery Algorithm

Advanced Recovery Techniques

- Support **high-concurrency** *locking techniques*,
 - ▶ such as those **used** for **B⁺-tree** *concurrency control*
- Operations like **B⁺-tree** insertions and deletions **release** locks **early**.
 - They **cannot** be **undone** *by restoring old values* (**physical undo**),
 - ▶ since once a lock is released,
 - other transactions may have updated the **B⁺-tree**.
 - Instead, insertions (resp. deletions) are **undone** by:
 - ▶ executing a deletion (resp. insertion) operation (known as **logical undo**).
- For such operations, **undo** log records should contain:
 - ▶ the *undo operation* to be executed
 - called **logical undo** logging, in contrast to *physical undo logging*.
- **Redo** information is logged **physically**
 - ▶ (that is, new value for each write) even for such operations

- Logical **redo** is very **complicated** !
 - since database state on **disk** may **not** be “*operation consistent*”
- Operation logging is done as follows:
 1. When operation starts, log $\langle T_i, O_j, \text{operation-begin} \rangle$.
 - Here O_j is a **unique** identifier of the operation instance.
 2. While operation is executing,
 - normal log records with physical redo and physical undo *information* are logged.
 3. When operation completes, $\langle T_i, O_j, \text{operation-end}, U \rangle$ is logged,
 - where U contains *information* needed to perform a **logical** undo.
- If crash/rollback occurs **before** operation completes:
 - the **operation-end** log record is **not** found, and
 - the **physical** undo information is used to undo operation.
- If crash/rollback occurs **after** the operation completes:
 - the **operation-end** log record **is** found, and in this case
 - **logical** undo is performed using U ;
 - the **physical** undo *information* for the operation is *ignored*.
- **Redo** of operation (after **crash**):
 - still uses **physical** redo *information*.

Rollback of transaction T_i is done as follows:

- Scan the log backwards
 1. If a log record $\langle T_i, X, V_1, V_2 \rangle$ is found,
 - perform the undo and
 - **log** a special **redo-only** log record $\langle T_i, X, V_1 \rangle$.
 2. If a $\langle T_i, O_j, \text{operation-end}, U \rangle$ record is found
 - Rollback the operation **logically** using the undo information U .
 - Updates performed during roll back **are logged**
 - » just like during *normal operation* execution.
 - At the end of the operation rollback,
 - » instead of logging an **operation-end** record,
 - » **generate** a record $\langle T_i, O_j, \text{operation-abort} \rangle$.
 - **Skip** all preceding log records for T_i until:
 - the record $\langle T_i, O_j, \text{operation-begin} \rangle$ is found
 3. If a **redo-only** record is found **ignore it**

4. If a $\langle T_i, O_j, \text{operation-abort} \rangle$ record is found:
 - skip all preceding log records for T_i until :
 - the record $\langle T_i, O_j, \text{operation-begin} \rangle$ is found.
 5. Stop the scan when the record $\langle T_i, \text{start} \rangle$ is found
 6. Add a $\langle T_i, \text{abort} \rangle$ record to the log
- Some points to note:
 - Cases 3 and 4 above can occur **only if**:
 - the database crashes while a transaction is being rolled back.
 - Skipping of log records as in case 4 is **important** :
 - to **prevent multiple** rollback of the *same operation*.

The following actions are taken when **recovering** from **system** crash:

1. Scan log forward from last $\langle \text{checkpoint } L \rangle$ record
 1. Repeat history by physically redoing :
 - all updates of all transactions,
 2. Create an undo-list during the scan as follows:
 - *undo-list* is set to L initially
 - Whenever $\langle T_i \text{ start} \rangle$ is found T_i is added to *undo-list*
 - Whenever $\langle T_i \text{ commit} \rangle$ or $\langle T_i \text{ abort} \rangle$ is found, T_i is deleted from *undo-list*
- This **brings database** to state as of crash,
 - ▶ with committed as well as uncommitted transactions having been **redone**.
- Now *undo-list* contains transactions that are **incomplete**, that is,
 - ▶ have neither committed nor been **fully** rolled back.
2. Scan log backwards, performing undo on log records of transactions found in *undo-list*.
 - Transactions are rolled back
 - ▶ as described earlier.
 - When $\langle T_i \text{ start} \rangle$ is found for a transaction T_i in *undo-list*,
 - ▶ write a $\langle T_i \text{ abort} \rangle$ log record.
 - Stop scan when $\langle T_i \text{ start} \rangle$ records have been found for **all** T_i in *undo-list*
- This **undoes** the effects of incomplete *transactions*
 - ▶ (those with neither **commit** nor **abort** log records).
- Recovery is now **complete**.
- Checkpointing is done as follows:
 1. Output all log records in memory to **stable storage**
 2. Output all modified buffer blocks to **disk**

3. Output a < **checkpoint** *L* > record to log on **stable storage** .
 4. Transactions are **not** allowed to perform **any** actions
 - while **checkpointing** is in progress.
- **Fuzzy checkpointing allows transactions to progress**
 - while the most time consuming **parts** of checkpointing are in progress
 - **Fuzzy checkpointing** is done as follows:
 1. Temporarily **stop all** updates by transactions
 2. Write a <**checkpoint** *L*> log record and **force** log to *stable storage*
 3. Note list *M* of modified buffer blocks
 4. Now permit transactions to **proceed** with their actions
 5. Output to disk **all** modified buffer blocks in list *M*
 - blocks should not be updated **until** being **output**
 - Follow **WAL**: **all** log records pertaining to a block must be **output** before the block is output
 6. Store a pointer to the **checkpoint** record
 - in a fixed position **last_checkpoint** on **disk**
 - When **recovering** using a **fuzzy** checkpoint,
 - **start scan** from the **checkpoint** record pointed to by **last_checkpoint**
 - Log records **before last_checkpoint**
 - ▶ have their updates reflected in database on disk, and
 - ▶ **need not** be redone.
 - Incomplete checkpoints,
 - ▶ where **system** had crashed while performing checkpoint,
 - are handled **safely**

PARALLEL DATABASES

- Introduction
- I/O Parallelism
- Interquery Parallelism
- Intraquery Parallelism
- Intraoperation Parallelism
- Interoperation Parallelism
- Design of Parallel Systems

Introduction

- Parallel machines are becoming quite common and affordable
 - Prices of microprocessors, memory and disks have dropped sharply
 - Recent desktop computers feature multiple processors and this trend is projected to accelerate
- Databases are growing increasingly large
 - large volumes of transaction data are collected and stored for later analysis.
 - multimedia objects like images are increasingly stored in databases
- Large-scale parallel database systems increasingly used for:
 - storing large volumes of data
 - processing time-consuming decision-support queries
 - providing high throughput for transaction processing

Parallelism in Databases

- Data can be partitioned across multiple disks for parallel I/O.
- Individual relational operations (e.g., sort, join, aggregation) can be executed in parallel
 - data can be partitioned and each processor can work independently on its own partition.
- Queries are expressed in high level language (SQL, translated to relational algebra)
 - makes parallelization easier.
- Different queries can be run in parallel with each other. Concurrency control takes care of conflicts.
- Thus, databases naturally lend themselves to parallelism.

I/O Parallelism

- Reduce the time required to retrieve relations from disk by partitioning
- the relations on multiple disks.

- Horizontal partitioning – tuples of a relation are divided among many disks such that each tuple resides on one disk.
- Partitioning techniques (number of disks = n):

Round-robin:

Send the i^{th} tuple inserted in the relation to disk $i \bmod n$.

Hash partitioning:

- Choose one or more attributes as the partitioning attributes.
- Choose hash function h with range $0 \dots n - 1$
- Let i denote result of hash function h applied to the partitioning attribute value of a tuple. Send tuple to disk i .
- **Range partitioning:**
 - Choose an attribute as the partitioning attribute.
 - A partitioning vector $[v_0, v_1, \dots, v_{n-2}]$ is chosen.
 - Let v be the partitioning attribute value of a tuple. Tuples such that $v_i \leq v_{i+1}$ go to disk $i + 1$. Tuples with $v < v_0$ go to disk 0 and tuples with $v \geq v_{n-2}$ go to disk $n-1$.

E.g., with a partitioning vector $[5, 11]$, a tuple with partitioning attribute value of 2 will go to disk 0, a tuple with value 8 will go to disk 1, while a tuple with value 20 will go to disk 2.

Comparison of Partitioning Techniques

- Evaluate how well partitioning techniques support the following types of data access:
 1. Scanning the entire relation.
 2. Locating a tuple associatively – **point queries**.
 - E.g., $r.A = 25$.
 3. Locating all tuples such that the value of a given attribute lies within a specified range – **range queries**.
 - E.g., $10 \leq r.A < 25$.

Round robin:

- Advantages
 - Best suited for sequential scan of entire relation on each query.
 - All disks have almost an equal number of tuples; retrieval work is thus well balanced between disks.
- Disadvantages
 - Range queries are difficult to process

- No clustering -- tuples are scattered across all disks

Hash partitioning:

- Good for sequential access
 - **Assuming** hash function is good, and partitioning attributes form a key, tuples will be equally distributed between disks
 - Retrieval work is then well balanced between disks.
- Good for point queries on partitioning attribute
 - Can lookup single disk, leaving others available for answering other queries.
 - Index on partitioning attribute can be local to disk, making lookup and update more efficient
- No clustering, so difficult to answer range queries

Range partitioning:

- Provides data clustering by partitioning attribute value.
- Good for sequential access
- Good for point queries on partitioning attribute: only one disk needs to be accessed.
- For range queries on partitioning attribute, one to a few disks may need to be accessed
 - Remaining disks are available for other queries.
 - Good if result tuples are from one to a few blocks.
 - If many blocks are to be fetched, they are still fetched from one to a few disks, and potential parallelism in disk access is wasted
 - ▶ Example of execution skew.

Partitioning a Relation across Disks

- If a relation contains only a few tuples which will fit into a single disk block, then assign the relation to a single disk.
- **Large relations** are preferably partitioned **across all** the available **disks**.
- If a relation consists of m disk blocks and there are n disks available in the system, then the relation should be allocated $\min(m,n)$ disks.

Handling of Skew

- The distribution of tuples to disks may be **skewed** — that is, some disks have many tuples, while others may have fewer tuples.
- **Types of skew:**
 - 1) **Attribute-value skew.**

- ▶ Some values appear in the partitioning attributes of **many tuples**; all the tuples with the same value for the partitioning attribute end up in the same partition.
- ▶ Can **occur** with range-partitioning and hash-partitioning.

2) **Partition skew.**


- ▶ With range-partitioning, **badly** chosen partition **vector** may assign too many tuples to some partitions and too few to others.
- ▶ Less likely with hash-partitioning if a **good** hash-function is chosen.

Handling Skew in Range-Partitioning

- To create a **balanced** partitioning vector (assuming partitioning attribute forms a key of the relation):
 - **Sort** the relation on the partitioning attribute.
 - Construct the **partition vector** by scanning the relation in sorted order as follows.
 - ▶ After **every** $1/n^{th}$ of the relation has been read, the value of the partitioning attribute of the next tuple is **added to** the partition **vector**.
 - n denotes the **number** of **partitions** to be constructed.
 - Duplicate entries or imbalances can result **if duplicates** are present in partitioning attributes.
- Alternative technique based on **histograms** used in practice

Handling Skew using Histograms

- **Balanced** partitioning vector can be constructed from **histogram** in a relatively straightforward fashion
 - **Assume** uniform distribution within each range of the histogram
- Histogram can be constructed by scanning relation, or sampling (blocks containing) tuples of the relation



Handling Skew Using Virtual Processor Partitioning

- Skew in range partitioning can be handled **elegantly** using **virtual processor partitioning**:
 - create a large number of partitions (say **10 to 20** times the number of processors)
 - Assign virtual processors to partitions either in round-robin fashion or based on estimated cost of processing each virtual partition
- **Basic idea**:
 - If any normal partition would have been skewed, it is very likely the **skew** is **spread** over a number of virtual partitions
 - Skewed virtual partitions get **spread** across a number of **processors**, so work gets distributed evenly!

Interquery Parallelism

- **Queries**/transactions execute in **parallel** with one another.
- Increases transaction **throughput**; used primarily to **scale up** a transaction processing system to support a **larger number** of **transactions** per second.
- Easiest form of parallelism to support, particularly in a **shared-memory** parallel database, because even sequential database systems support concurrent processing.
- More **complicated** to implement **on** shared-disk or shared-nothing architectures:

- **Locking** and logging must be coordinated by passing messages between processors.
- Data in a **local buffer** may have been updated at another processor.
- **Cache-coherency** has to be **maintained** — reads and writes of data in buffer must find **latest version of data**.

Cache Coherency Protocol

- Example of a cache coherency **protocol** for shared disk systems:
 - Before **reading/writing** to a page, the **page** must be **locked** in shared/exclusive mode.
 - On locking a page, the page must be read from disk
 - Before **unlocking** a page, the **page** must be **written** to disk if it was modified.
- More **complex protocols** with fewer disk reads/writes **exist**.
- Cache coherency **protocols** for shared-nothing systems are similar. Each database page is assigned a **home processor**. Requests to fetch the page or write it to disk are sent to the home processor.

Intraquery Parallelism

- Execution of a **single query** in **parallel** on multiple processors/disks; important for **speeding up** long-running queries.
- **Two** complementary **forms** of intraquery parallelism :
 - **Intraoperation Parallelism** – parallelize the execution of **each** individual operation in the query.
 - **Interoperation Parallelism** – execute the **different** operations in a query expression in parallel.
- the **first** form **scales better** with increasing parallelism because the number of tuples processed by each operation **is** typically **more than** the number of operations in a query

Parallel Processing of Relational Operations

- Our discussion of parallel algorithms **assumes**:
 - *read-only* queries
 - shared-nothing architecture
 - ***n*** processors, ***P*₀, ..., *P*_{*n*-1}**, and ***n*** disks ***D*₀, ..., *D*_{*n*-1}**, where disk ***D*_{*i*}** is associated with processor ***P*_{*i*}**.
- If a processor has multiple disks they can simply simulate a **single disk** ***D*_{*i*}**.

- Shared-nothing architectures can be efficiently simulated on shared-memory and shared-disk systems.
 - Algorithms for shared-nothing systems can thus be run on shared-memory and shared-disk systems.
 - However, some optimizations may be possible.

Parallel Sort

Range-Partitioning Sort

- Choose processors P_0, \dots, P_m , where $m \leq n - 1$ to do sorting.
- Create range-partition vector with m entries, on the sorting attributes
- Redistribute the relation using range partitioning
 - all tuples that lie in the i^{th} range are sent to processor P_i
 - P_i stores the tuples it received temporarily on disk D_i .
 - This step requires **I/O** and **communication overhead**.
- Each processor P_i sorts its partition of the relation locally.
- Each processor executes same operation (sort) **in parallel** with other processors, without any interaction with the others (**data parallelism**).
- Final **merge** operation is **trivial**: range-partitioning ensures that, for $1 \leq j \leq m$, the key values in processor P_i are all less than the key values in P_j .

Parallel External Sort-Merge

- Assume the relation has already been partitioned among disks D_0, \dots, D_{n-1} (in whatever manner).
- Each processor P_i **locally** sorts the data on disk D_i .
- The sorted runs on each processor are then merged to get the final sorted output.
- **Parallelize the merging** of sorted runs as follows:
 - The sorted partitions at each processor P_i are **range-partitioned** across the processors P_0, \dots, P_{m-1} .
 - Each processor P_i performs a merge on the streams as they are received, to get a single sorted run.
 - The sorted runs on processors P_0, \dots, P_{m-1} are concatenated to get the final result.


Parallel Join

- The join operation requires **pairs of tuples** to be tested to see if they satisfy the join condition, and if they do, the pair is added to the join output.
- Parallel join algorithms attempt to **split the pairs** to be tested **over several processors**. Each processor then computes part of the join locally.

- In a final step, the **results** from each processor can be **collected** together to produce the final result.

Partitioned Join

- For equi-joins and natural joins, it is possible to *partition* the two input relations across the processors, and compute the join locally at each processor.
- Let r and s be the input relations, and we want to compute $r \bowtie_{r.A = s.B} s$.
- r and s each are partitioned into n partitions, denoted r_0, r_1, \dots, r_{n-1} and s_0, s_1, \dots, s_{n-1} .
- Can use either *range partitioning* or *hash partitioning*.
- r and s must be partitioned on their join attributes $r.A$ and $s.B$, using the same range-partitioning vector or hash function.
- Partitions r_i and s_i are sent to processor P_i ,
- Each processor P_i locally computes $r_i \bowtie_{r.A = s.B} s_i$. Any of the standard join methods can be used.

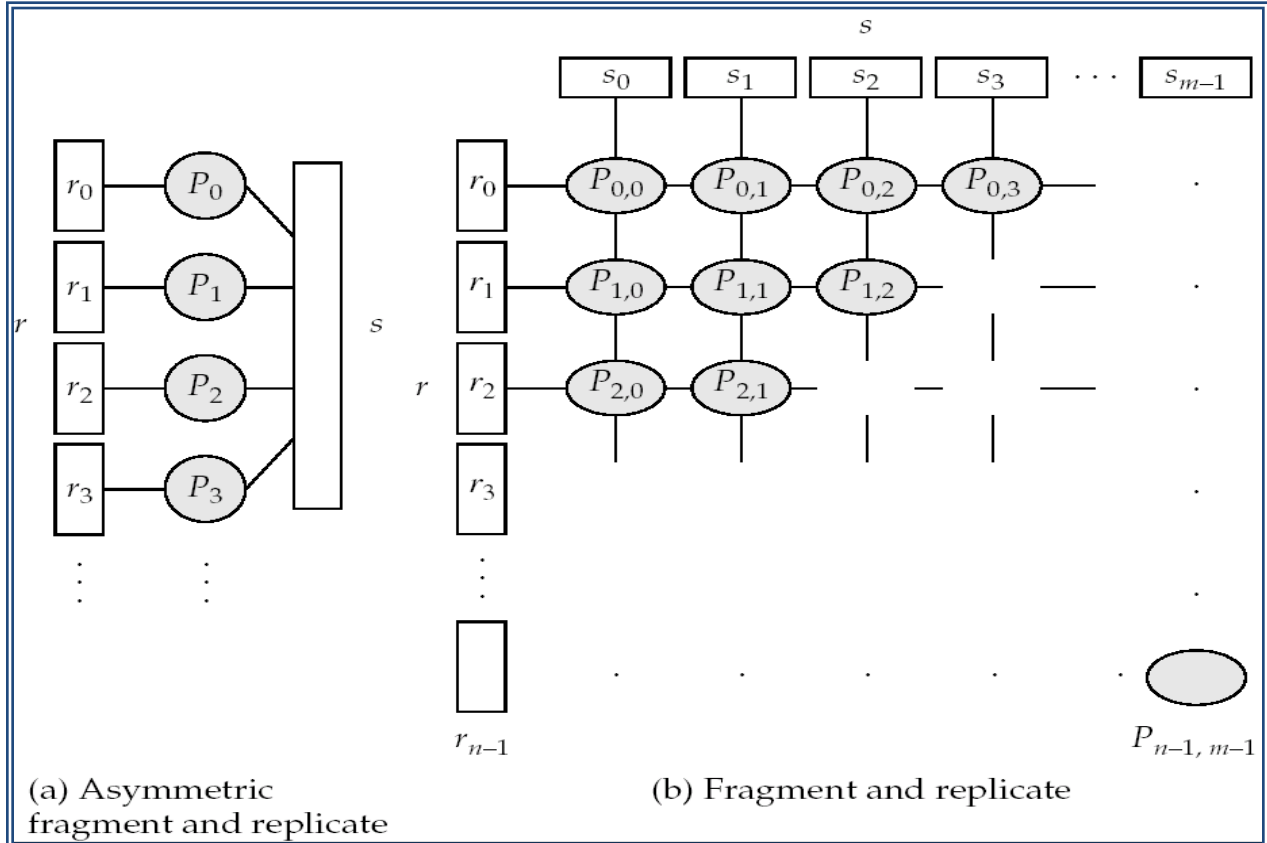


Fragment-and-Replicate Join

- Partitioning **not possible** for some join conditions
 - e.g., non-equijoin conditions, such as $r.A > s.B$.
- For joins where partitioning is not applicable, parallelization **can be** accomplished by **fragment and replicate** technique
- Special case – **asymmetric fragment-and-replicate**:

- One of the relations, say r , is partitioned; any partitioning technique can be used.
- The other relation, s , is replicated across all the processors.
- Processor P_i then locally computes the join of r_i with all of s using any join technique.

Depiction of Fragment-and-Replicate Joins



General case: reduces the sizes of the relations at each processor.

- r is partitioned into n partitions, r_0, r_1, \dots, r_{n-1} ; s is partitioned into m partitions, s_0, s_1, \dots, s_{m-1} .
- Any partitioning technique may be used.
- There must be at least $m * n$ processors.
- Label the processors as
- $P_{0,0}, P_{0,1}, \dots, P_{0,m-1}, P_{1,0}, \dots, P_{n-1,m-1}$.
- $P_{i,j}$ computes the join of r_i with s_j . In order to do so, r_i is replicated to $P_{i,0}, P_{i,1}, \dots, P_{i,m-1}$, while s_j is replicated to $P_{0,j}, P_{1,j}, \dots, P_{n-1,j}$.
- Any join technique can be used at each processor $P_{i,j}$.

- Both versions of fragment-and-replicate **work** with any join condition, since **every** tuple in r can be **tested** with **every** tuple in s .
- Usually has a **higher cost** than partitioning, since one of the relations (for asymmetric fragment-and-replicate) or both **relations** (for general fragment-and-replicate) have to be **replicated**.
- Sometimes **asymmetric** fragment-and-replicate is **preferable** even though partitioning could be used.
 - E.g., say s is **small** and r is **large**, and already partitioned. It may be **cheaper** to **replicate** s across all processors, rather than repartition r and s on the join attributes.

Partitioned Parallel Hash-Join

Parallelizing partitioned hash join:

- Assume s is **smaller** than r and therefore s is chosen as the build relation.
- A hash function h_1 takes the join attribute value of each tuple in s and maps this tuple to one of the n processors.
- Each processor P_i reads the tuples of s that are on its disk D_i , and **sends** each tuple to the appropriate processor based on **hash** function h_1 . Let s_i denote the tuples of relation s that are sent to processor P_i .
- As **tuples** of relation s are **received** at the destination processors, they are **partitioned** further using another hash function, h_2 , which is used to **compute** the hash-join **locally**.
- Once the tuples of s have been **distributed**, the larger relation r is redistributed across the m processors using the hash function h_1
- Let r_i denote the tuples of relation r that are sent to processor P_i .
- As the r tuples are received at the destination processors, they are repartitioned using the function h_2
- (just as the probe relation is partitioned in the sequential hash-join algorithm).
- Each processor P_i **executes** the build and probe **phases** of the hash-join algorithm on the local partitions r_i and s_i of r and s to produce a partition of the final result of the hash-join.
- Note: Hash-join **optimizations** can be applied to the **parallel** case
- e.g., the **hybrid** hash-join algorithm can be used to **cache** some of the **incoming tuples** in memory and avoid the cost of writing them and reading them back in.

Parallel Nested-Loop Join

- Assume that
 - relation s is much smaller than relation r and that r is stored by partitioning.
 - there is an index on a join attribute of relation r at each of the partitions of relation r .
- Use **asymmetric** fragment-and-replicate, with relation s being replicated, and using the existing partitioning of relation r .
- Each processor P_j where a partition of relation s is stored reads the tuples of relation s stored in D_j , and replicates the tuples to every other processor P_i .
 - At the end of this phase, relation s is replicated at all sites that store tuples of relation r .
- Each processor P_i performs an indexed nested-loop join of relation s with the i^{th} partition of relation r .

Other Relational Operations

Selection $\sigma_{\theta}(r)$:

- If θ is of the form $a_i = v$, where a_i is an attribute and v a value.
 - If r is **partitioned** on a_i the selection is performed at a single processor.
- If θ is of the form $l \leq a_i \leq u$ (i.e., θ is a range selection) and the relation has been **range-partitioned** on a_i
 - Selection is performed at each **processor** whose **partition overlaps** with the **specified range** of values.
- In all other cases: the **selection** is **performed** in parallel at **all** the **processors**.
- **Duplicate elimination:**
 - Perform by using **either of the parallel sort techniques**
 - **eliminate** duplicates **as soon as** they are **found** during sorting.
 - Can also **partition** the tuples (using either **range-** or **hash-** partitioning) and perform duplicate **elimination locally** at each processor.
- **Projection:**
 - Projection **without duplicate elimination** can be performed as tuples are read in from disk in parallel.
 - If **duplicate elimination** is **required**, any of the above duplicate elimination techniques can be used.

Grouping/Aggregation

- **Partition** the **relation** on the **grouping attributes** and then **compute** the aggregate values **locally** at each processor.
- Can **reduce cost of transferring tuples** during partitioning **by partly computing aggregate values** before partitioning.
- Consider the **sum** aggregation operation:
 - Perform aggregation operation at each processor **P_i** on those tuples stored on disk **D_i**
 - ▶ **results in** tuples with **partial sums** at each processor.
 - Result of the **local aggregation** is partitioned on the **grouping attributes**, and the aggregation **performed again** at each processor **P_i** to get the **final result**.
- **Fewer tuples need to be sent** to other processors during partitioning.

Cost of Parallel Evaluation of Operations

- If there is **no skew** in the partitioning, **and** there is **no overhead** due to the parallel evaluation, expected **speed-up will be 1/n**
- If **skew** and **overheads** are also to be taken into account, the **time taken** by a parallel operation can be estimated as
$$T_{\text{part}} + T_{\text{asm}} + \max (T_0, T_1, \dots, T_{n-1})$$
 - T_{part} is the time for partitioning the relations
 - T_{asm} is the time for assembling the results
 - T_i is the **time taken** for the operation at processor **P_i**
 - ▶ this needs to be estimated taking into account the **skew**, and the **time wasted in contentions**.

Interoperator Parallelism

- **Pipelined parallelism**
 - Consider a join of four relations
 - ▶ $r_1 \bowtie r_2 \bowtie r_3 \bowtie r_4$
 - Set up a pipeline that computes the three joins in parallel
 - ▶ Let **P1** be assigned the computation of
$$\text{temp1} = r_1 \bowtie r_2$$
 - ▶ And **P2** be assigned the computation of $\text{temp2} = \text{temp1} \bowtie r_3$
 - ▶ And **P3** be assigned the computation of $\text{temp2} \bowtie r_4$
 - **Each** of these **operations** can **execute in parallel**, sending **result** tuples it computes to the next operation even as it is computing further results

- ▶ Provided a **pipelineable join** evaluation **algorithm** (e.g. indexed **nested loops join**) is used

Factors Limiting Utility of Pipeline Parallelism

- **Pipeline parallelism is useful** since it **avoids writing intermediate** results to disk
- Useful with **small number** of processors, **but does not scale up well** with more processors. One reason is that pipeline chains do not attain **sufficient length**.
- **Cannot pipeline operators** which do not produce output **until all inputs** have been accessed (e.g. aggregate and **sort**)
- **Little speedup** is obtained for the frequent cases of **skew** in which **one operator's** execution **cost** is much **higher** than the **others**

Independent Parallelism

- **Independent parallelism**
 - **Consider a join** of four relations

$r_1 \bowtie r_2 \bowtie r_3 \bowtie r_4$

- ▶ Let **P₁** be assigned the computation of

$$\text{temp}_1 = r_1 \bowtie r_2$$
- ▶ And **P₂** be assigned the computation of $\text{temp}_2 = r_3 \bowtie r_4$
- ▶ And **P₃** be assigned the computation of $\text{temp}_1 \bowtie \text{temp}_2$
- ▶ **P₁** and **P₂** can work independently **in parallel**
- ▶ **P₃** has to wait for **input** from **P₁** and **P₂**
 - **Can pipeline** output of **P₁** and **P₂** to **P₃**, combining independent parallelism and **pipelined parallelism**
- **Does not provide a high degree of parallelism**
 - ▶ useful with a lower degree of parallelism.
 - ▶ **less useful in a highly parallel system,**

Query Optimization

- Query **optimization in parallel databases** is significantly **more complex** than query optimization in sequential databases.
- **Cost models** are more **complicated**, since we must take into account partitioning **costs** and issues such as **skew** and resource **contention**.
- When **scheduling execution tree** in parallel system, must **decide**:
 - How to **parallelize** each **operation** and how many **processors** to use for it.

- What operations to pipeline, **what operations** to execute independently in parallel, and what operations to execute sequentially, one after the other.
- **Determining** the amount of resources **to allocate** for each operation is a problem.
 - E.g., allocating more processors than optimal **can result** in high **communication overhead**.
- **Long pipelines** should be avoided as the final operation may wait a lot for inputs, while **holding** precious **resources**
- The **number of** parallel **evaluation plans** from which to choose from is **much larger** than the number of sequential evaluation plans.
 - Therefore **heuristics** are **needed** while optimization
- Two alternative **heuristics** for choosing parallel plans:
 - **No pipelining** and inter-operation pipelining; just parallelize **every operation** across all processors.
 - ▶ Finding best plan is now much easier --- use standard optimization technique, but with **new cost model**
 - ▶ **Volcano** parallel database popularize the **exchange-operator** model
 - exchange operator is introduced into **query plans** to partition and distribute tuples
 - each operation works independently on local data on each processor, in **parallel** with **other copies of the operation**
 - **First** choose most efficient sequential plan and **then** choose how best to parallelize the operations in that plan.
 - ▶ Can explore pipelined parallelism as an option
- Choosing a **good** physical organization (**partitioning** technique) is important to speed up queries.

Design of Parallel Systems

Some **issues** in the **design** of parallel systems:

- **Parallel loading** of **data** from external sources is needed in order to handle large volumes of incoming data.
- **Resilience to failure** of some processors or disks.
 - **Probability** of some disk or processor **failing** is **higher** in a parallel system.
 - **Operation** (perhaps with degraded performance) should be **possible** in spite of failure.

- Redundancy achieved by **storing extra copy** of every data item at another processor.
- **On-line reorganization of data and schema changes must be supported.**
 - For example, **index** construction on **terabyte databases** can take hours or days even on a parallel system.
 - ▶ Need to **allow other processing** (insertions/deletions/updates) to be performed on relation even as index is being constructed.
 - Basic idea: index construction tracks changes and ``**catches up**'' on changes at the end.
- Also **need** support for **on-line** repartitioning and **schema** changes (executed **concurrently** with **other processing**).